

# Data Modelling Methods-IV

CS771: Introduction to Machine Learning  
Purushottam Kar



# Mid Semester Examination

- September 21<sup>st</sup>, 2017 (Thursday) 1300–1500 hrs
- Venue L18, 19, L20 (all OROS)
- Syllabus: till whatever we covered on Wednesday + maybe one question from today's lecture
- Open notes (handwritten only)
- No printed/photocopied material
- No laptops, i-pads, mobile phones (switched off)
- Please bring a notepad with you for rough work
- Please bring a pencil/eraser with you – we will not provide these
- Answers will have to be written on the question paper itself

# Outline of today's discussion

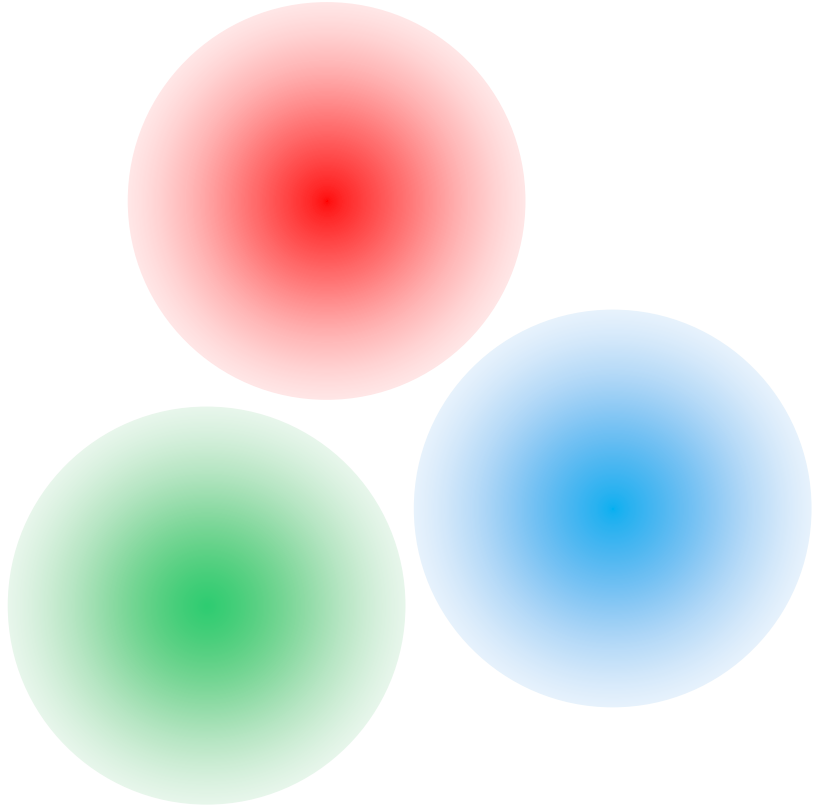
## DIMENSIONALITY REDUCTION TECHNIQUES

- Study an appropriate generative model for low-dim. Data
- Study the MLE for zero-noise condition (PCA)
- Study the MLE for noisy conditions (PPCA)
- See an efficient “power method” to solve the MLE

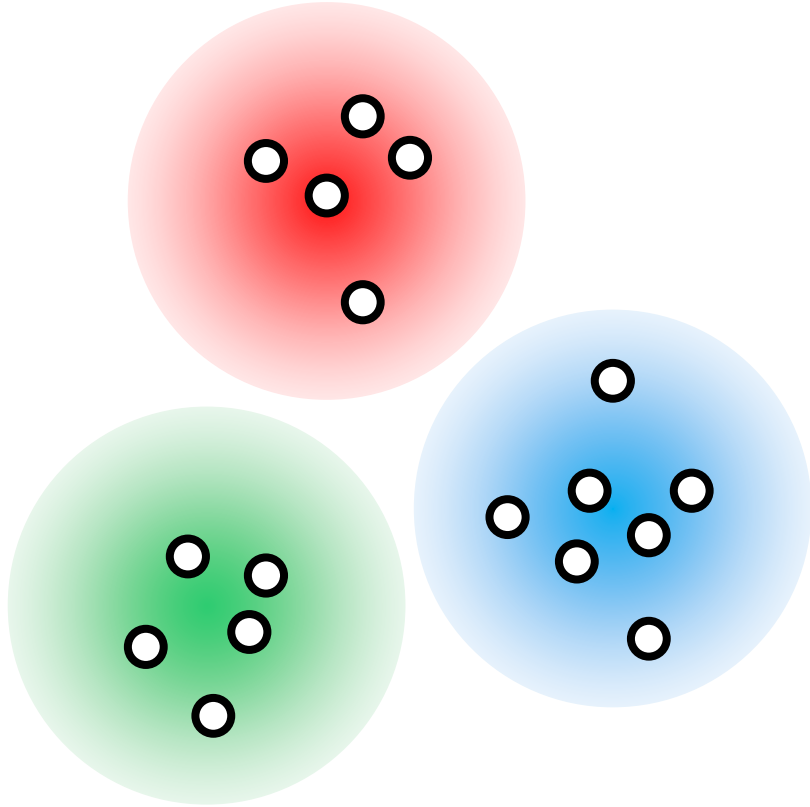
## AFTER MID-SEMS

- See a “soft”-assignment approach to solving the MLE
- See how the “hard” assignment rule can be used to solve PCA
- The One EM to Rule them All, Kernels, Deep learning, RecSys

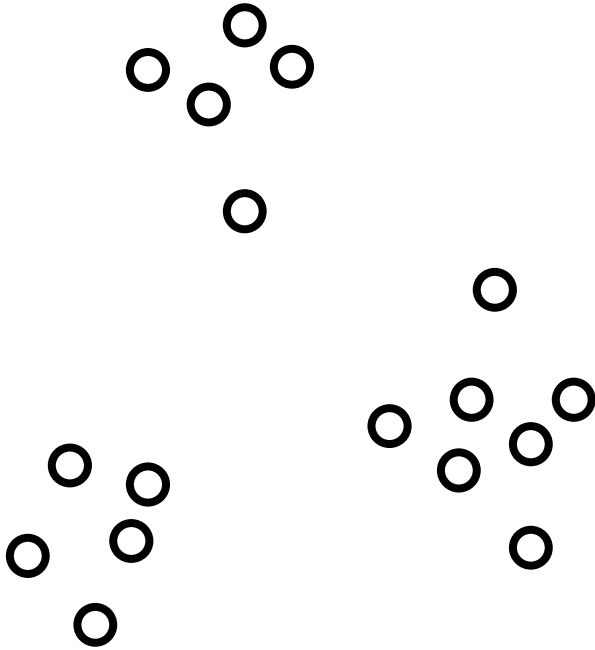
# Recap



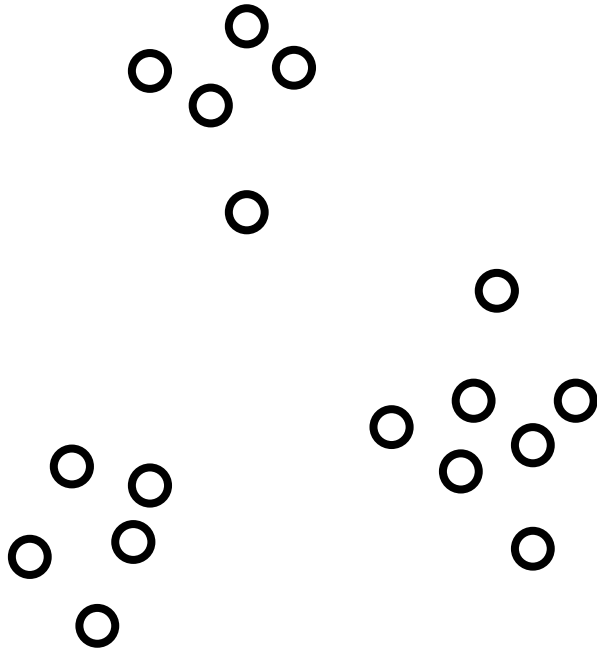
# Recap



# Recap

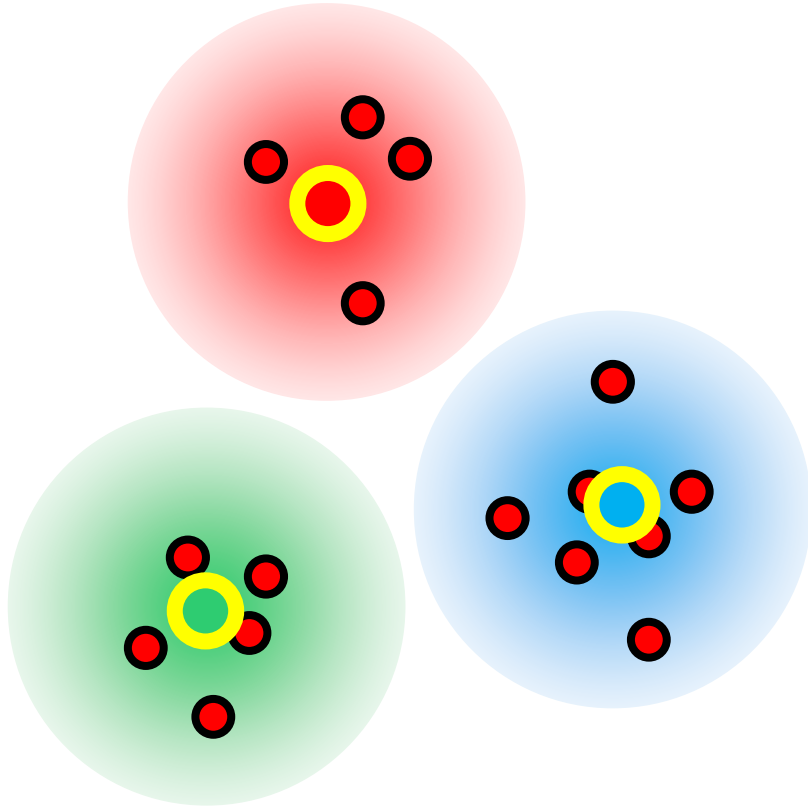


# Recap



Hard assignment – k-means  
Soft assignment – soft k-means

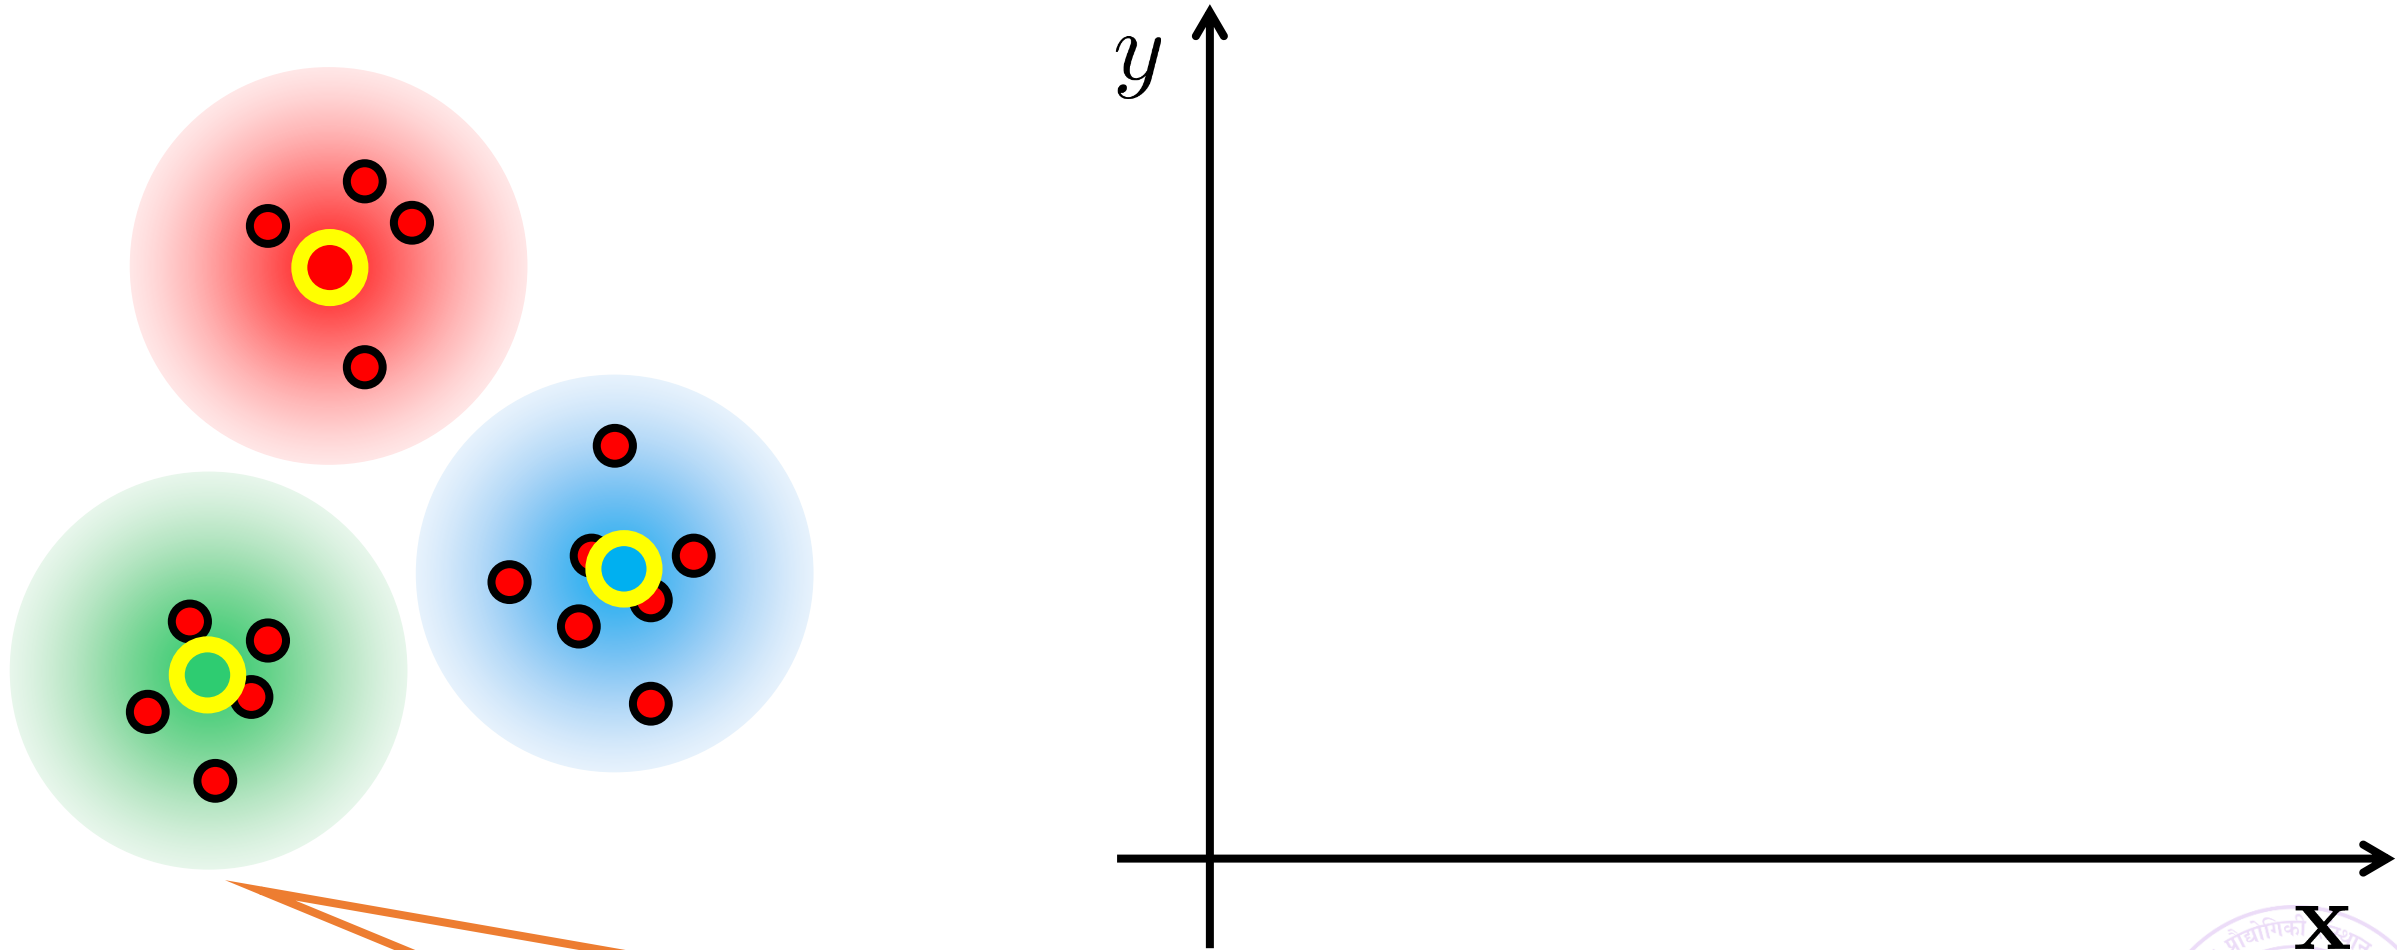
# Recap



Hard assignment – k-means  
Soft assignment – soft k-means

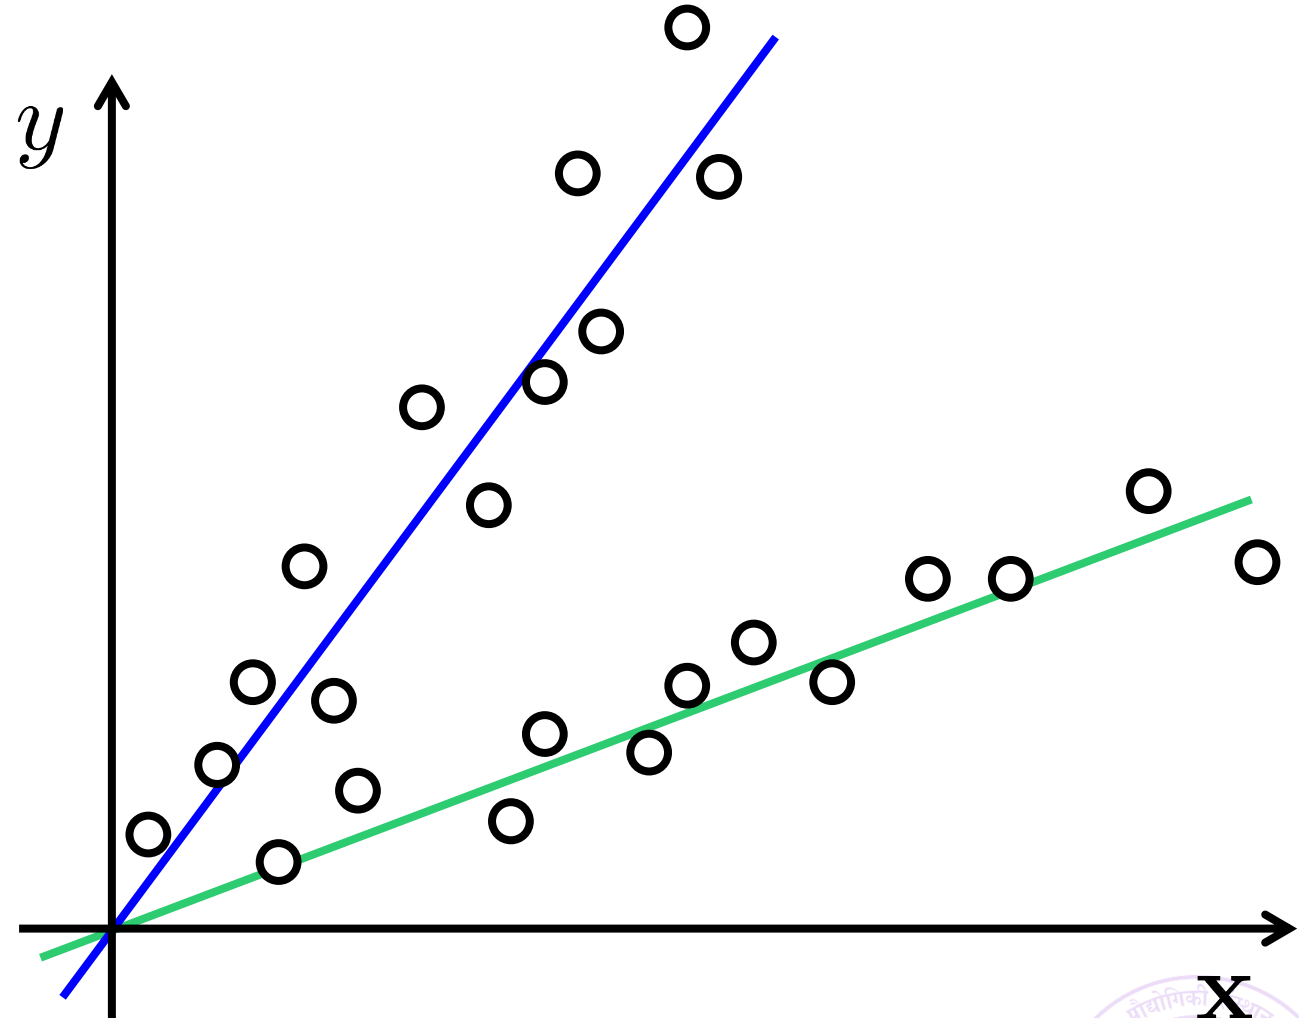
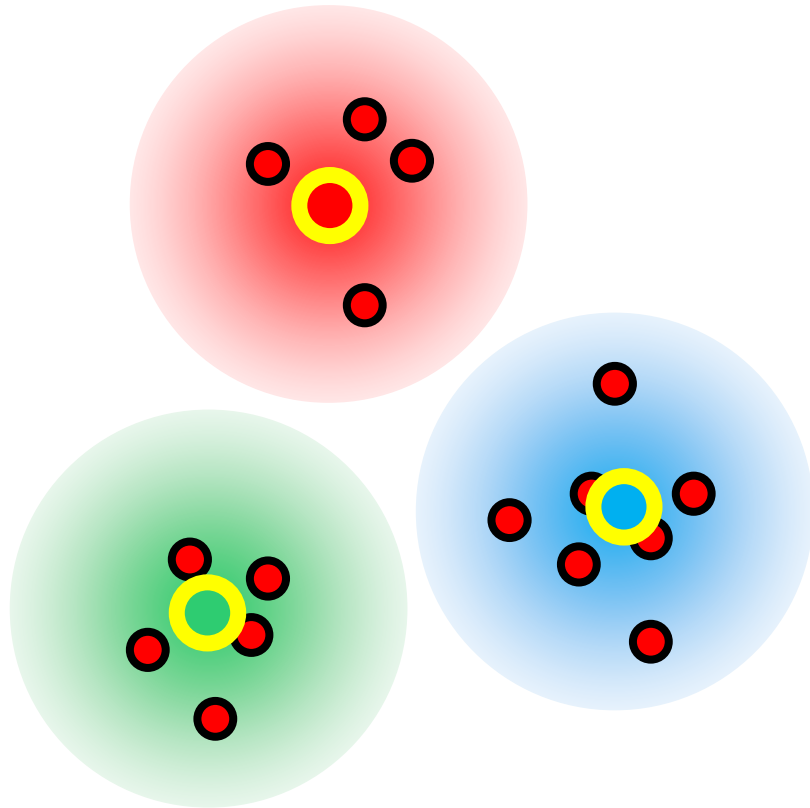


# Recap



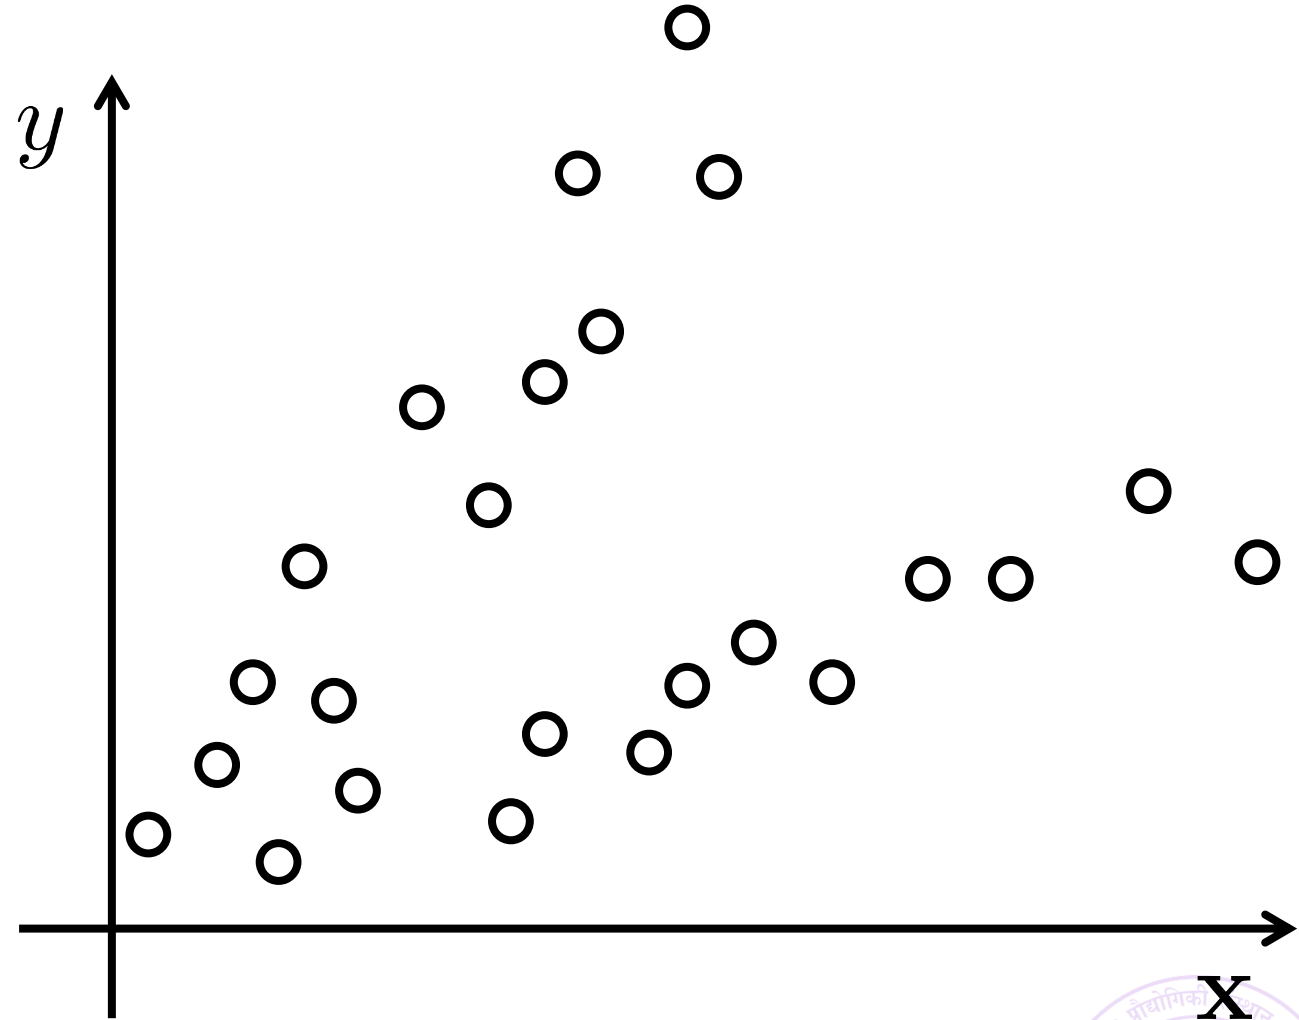
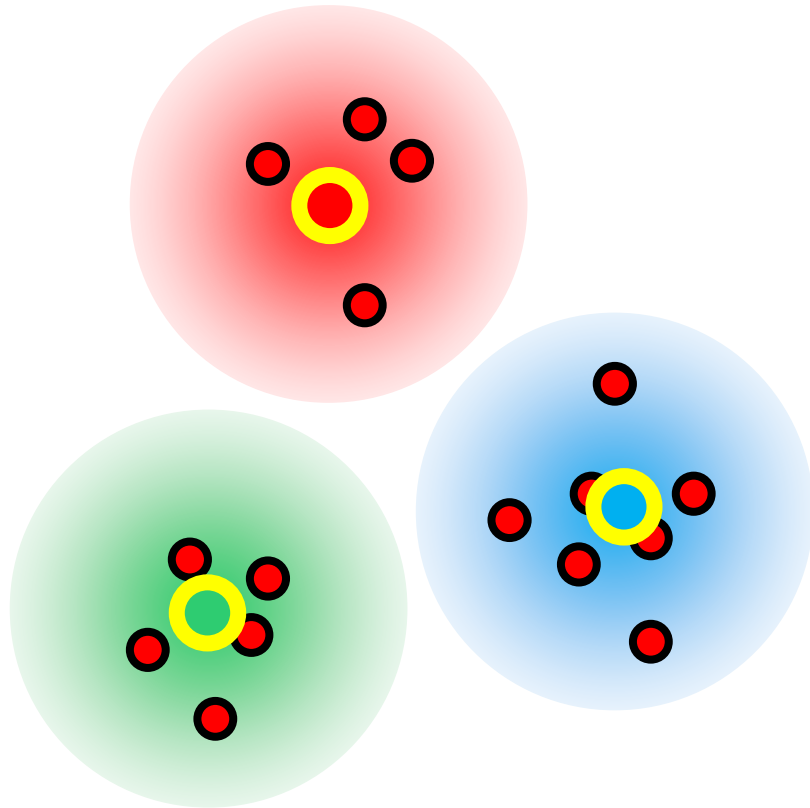
Hard assignment – k-means  
Soft assignment – soft k-means

# Recap



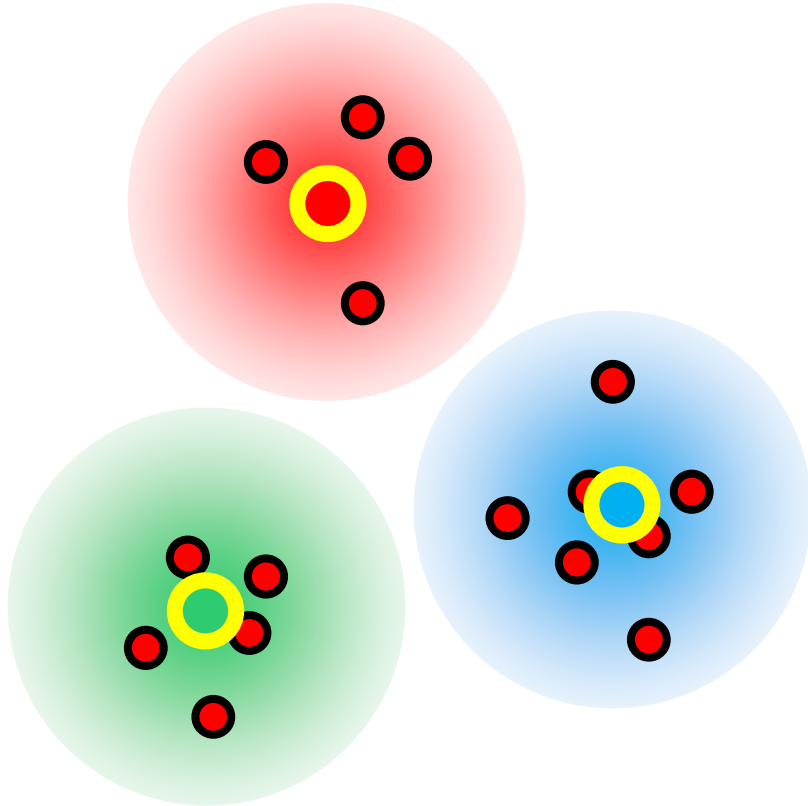
Hard assignment – k-means  
Soft assignment – soft k-means

# Recap

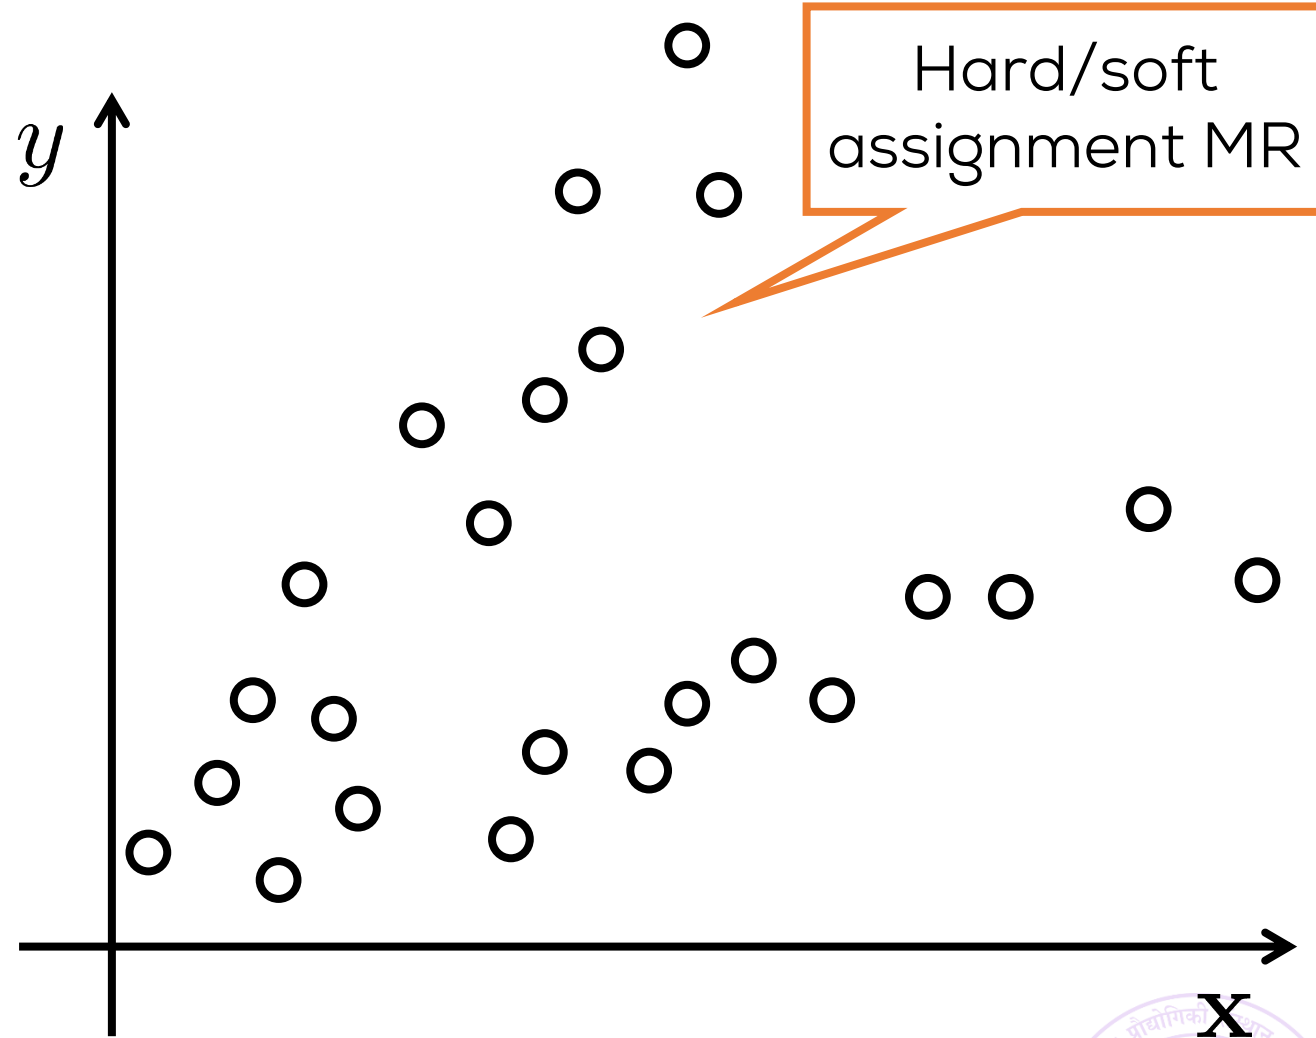


Hard assignment – k-means  
Soft assignment – soft k-means

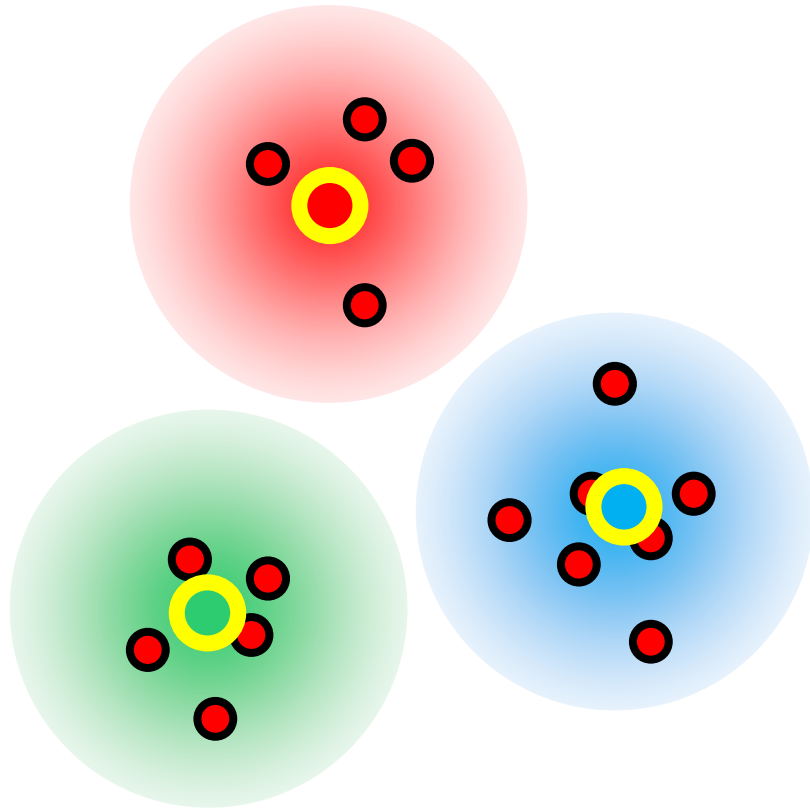
# Recap



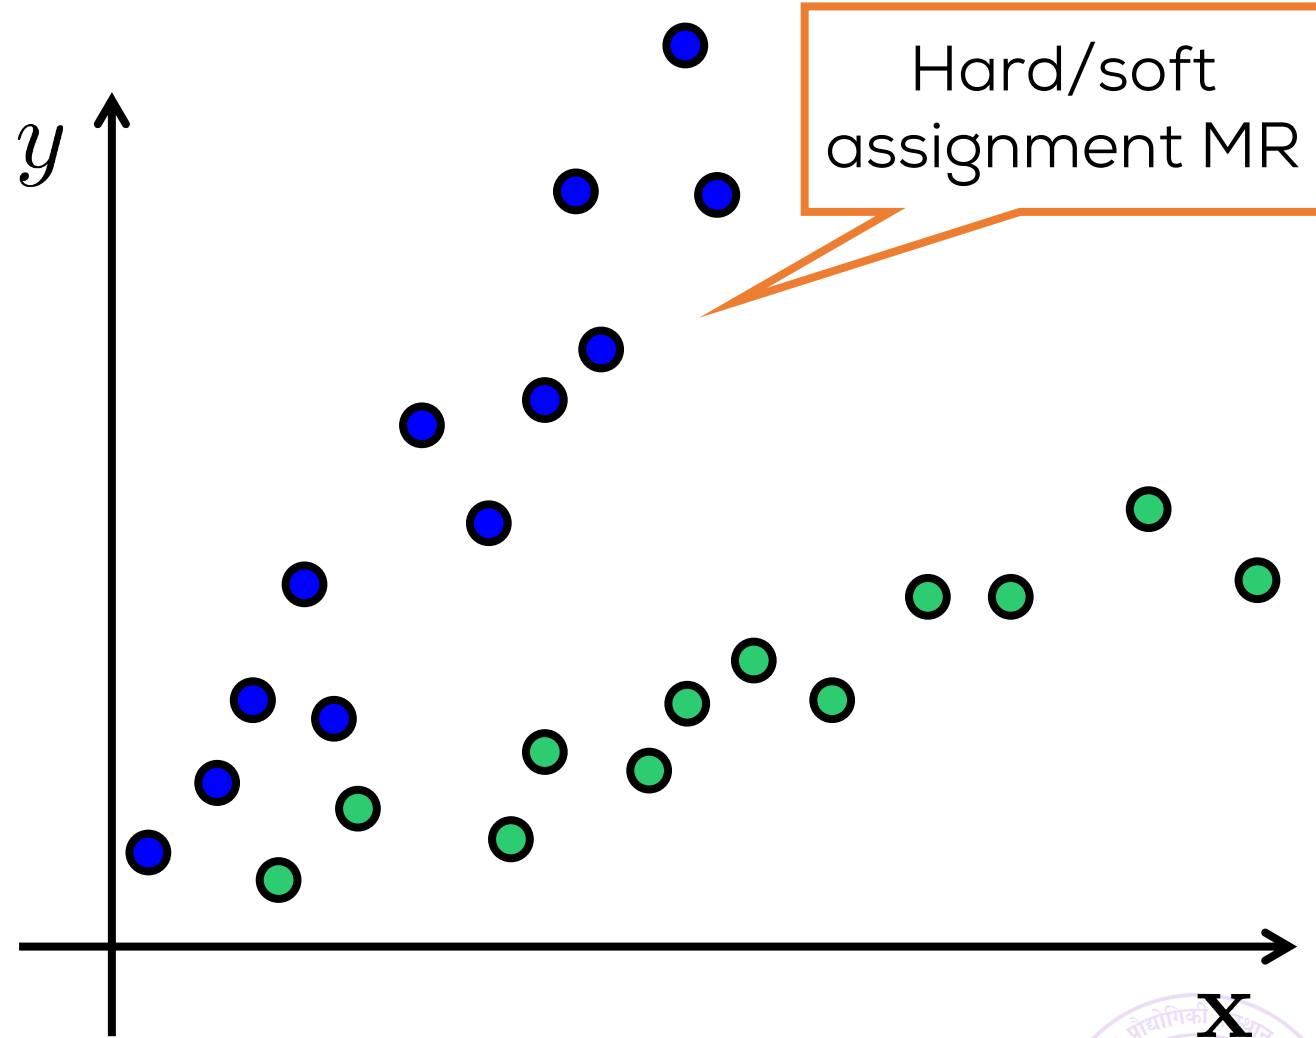
Hard assignment – k-means  
Soft assignment – soft k-means



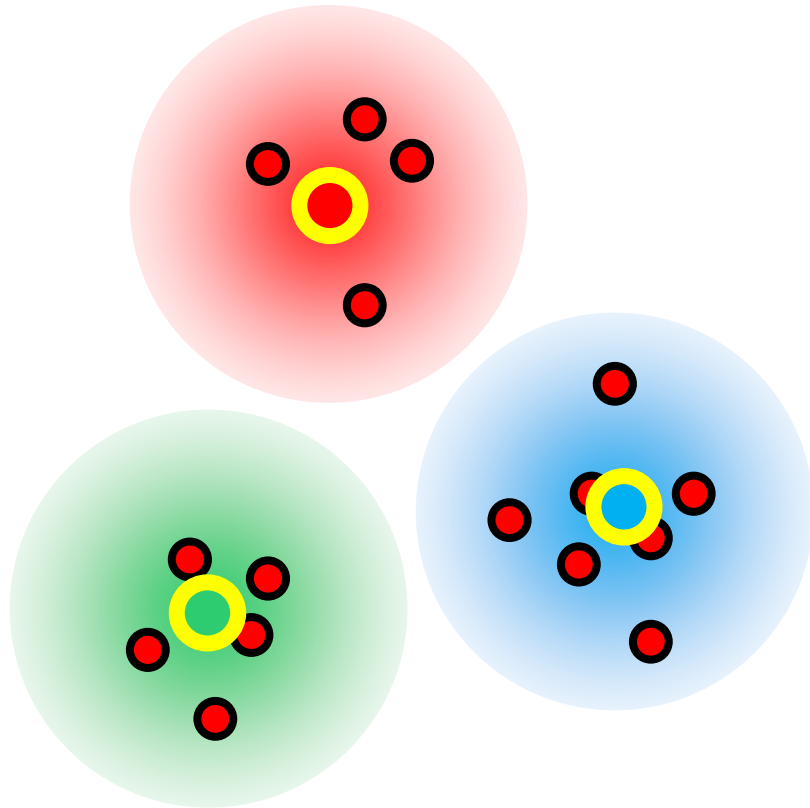
# Recap



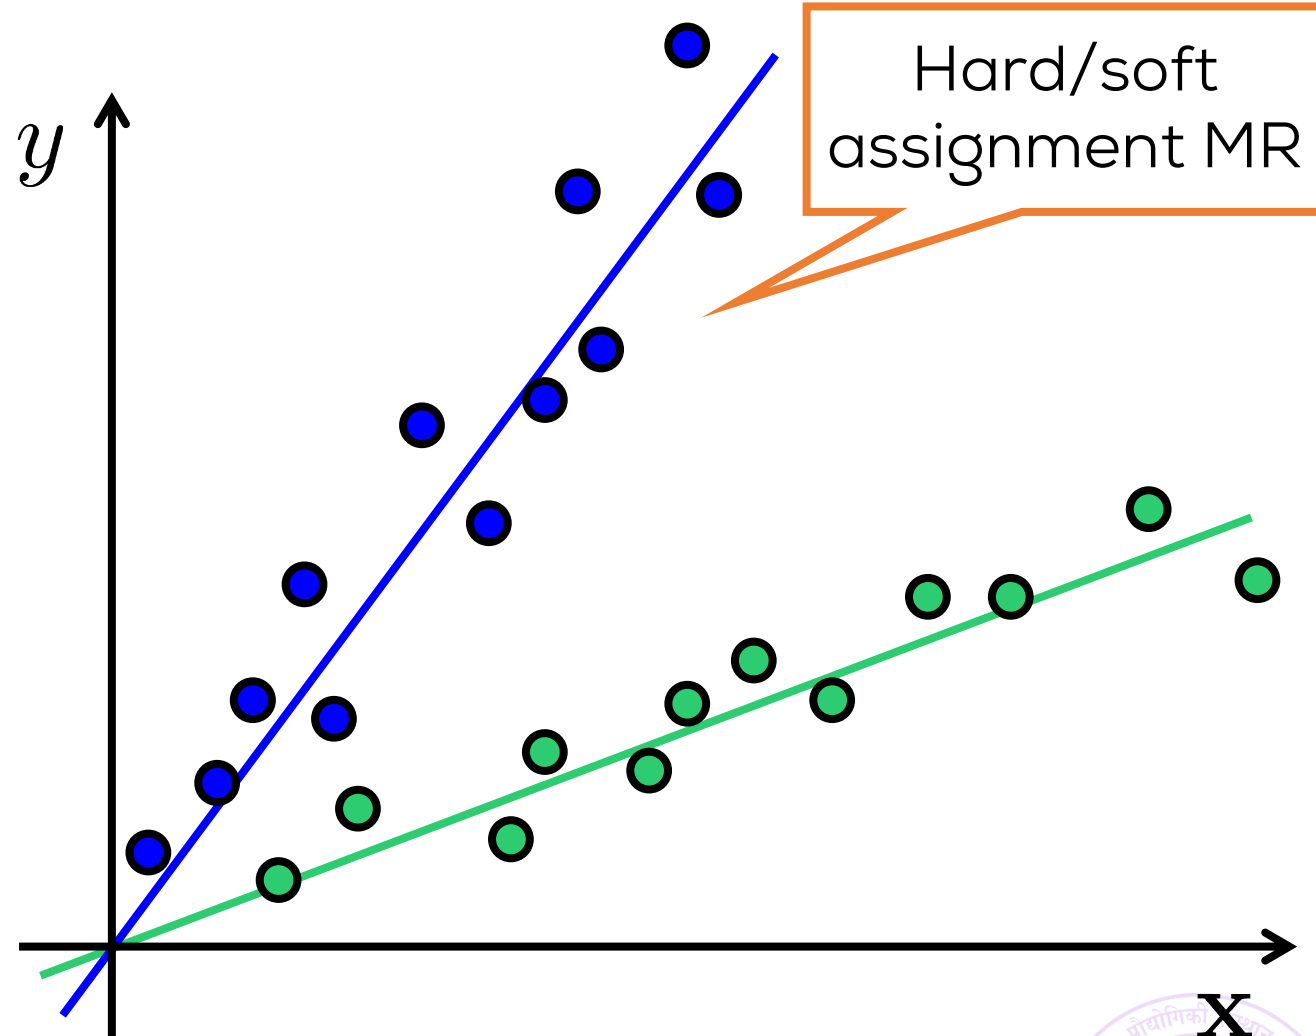
Hard assignment – k-means  
Soft assignment – soft k-means



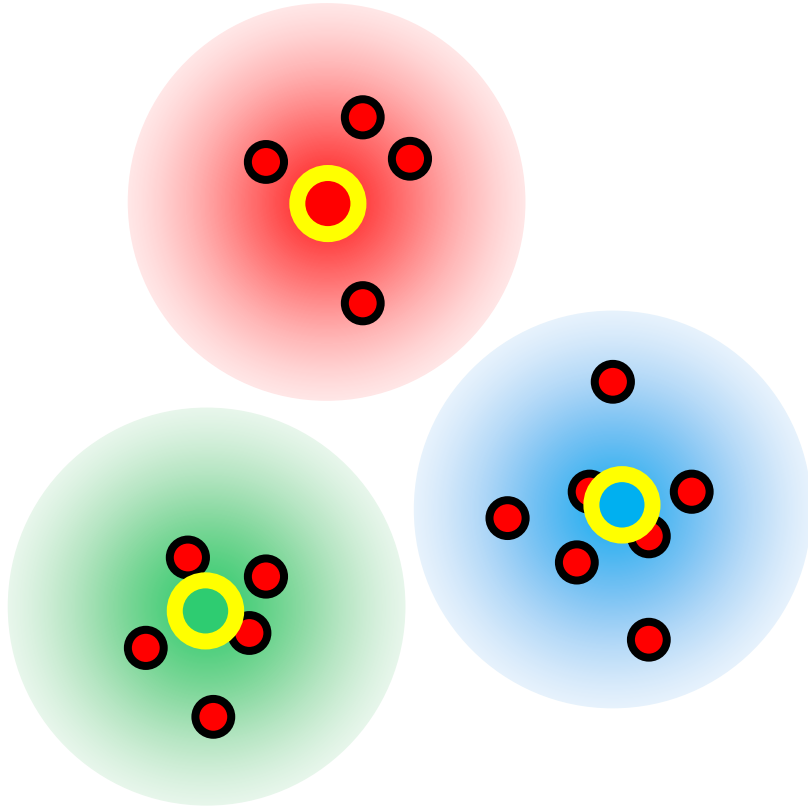
# Recap



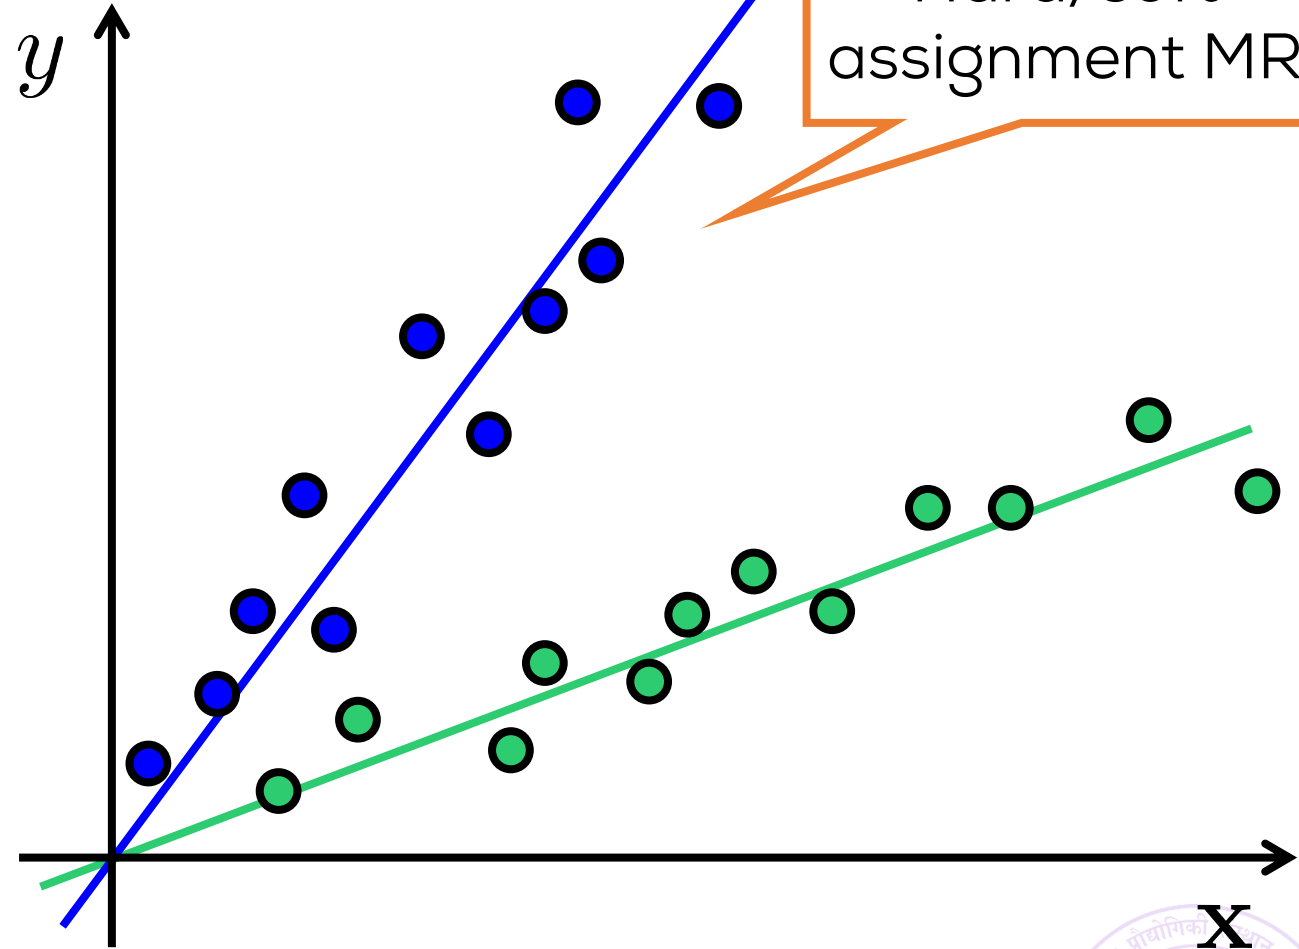
Hard assignment – k-means  
Soft assignment – soft k-means



# Recap



Hard assignment – k-means  
Soft assignment – soft k-means



Discovering hidden  
structure in data

# Low-dimensional Structure in Data

Sept 15, 2017

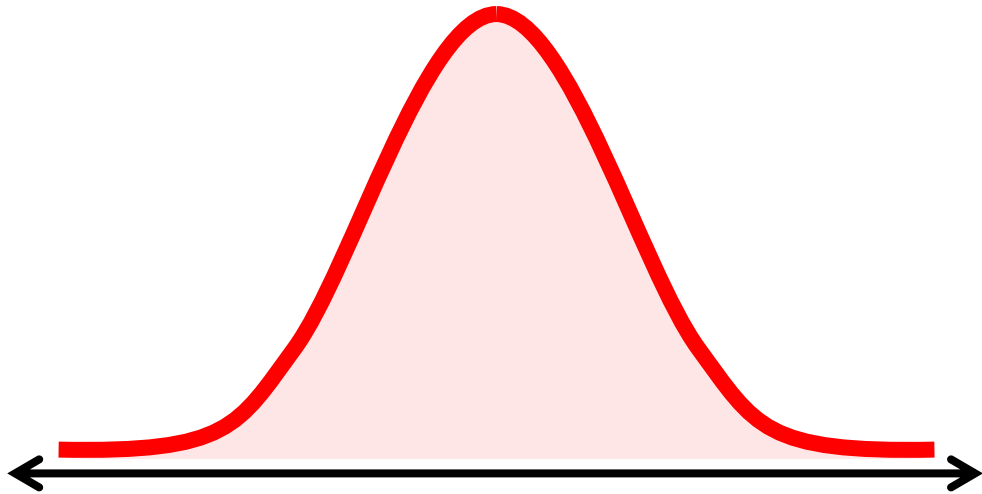




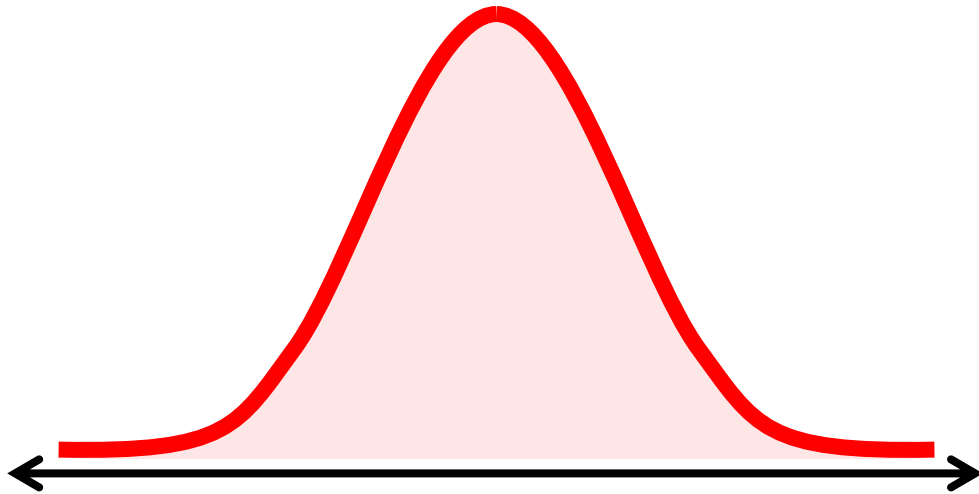
# Low-dimensional Structure in Data



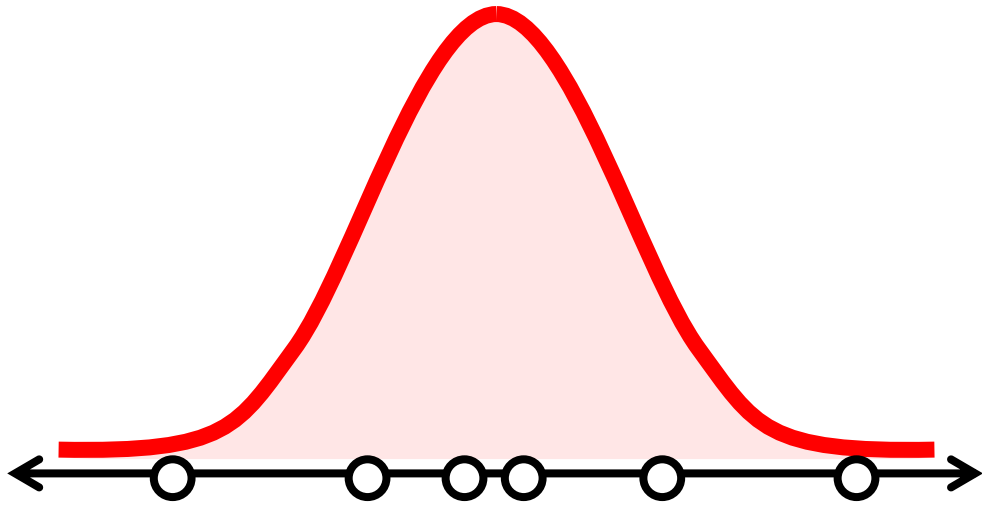
# Low-dimensional Structure in Data



# Low-dimensional Structure in Data



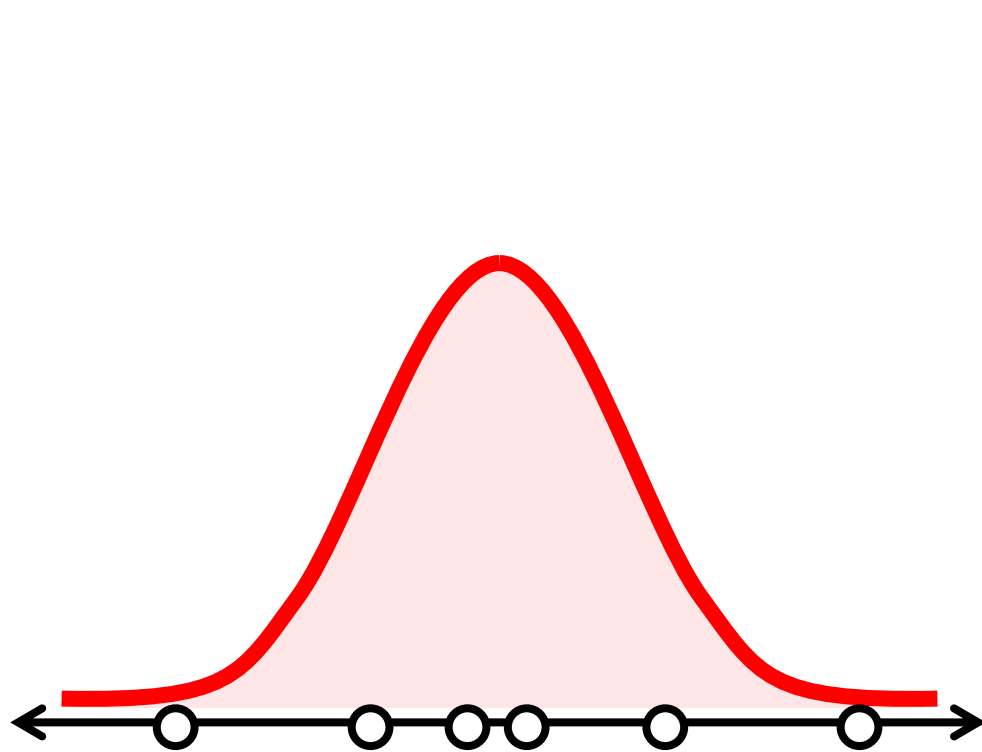
# Low-dimensional Structure in Data



$$\mathbf{z}^i \sim \mathcal{N}(\mathbf{0}, I_k)$$



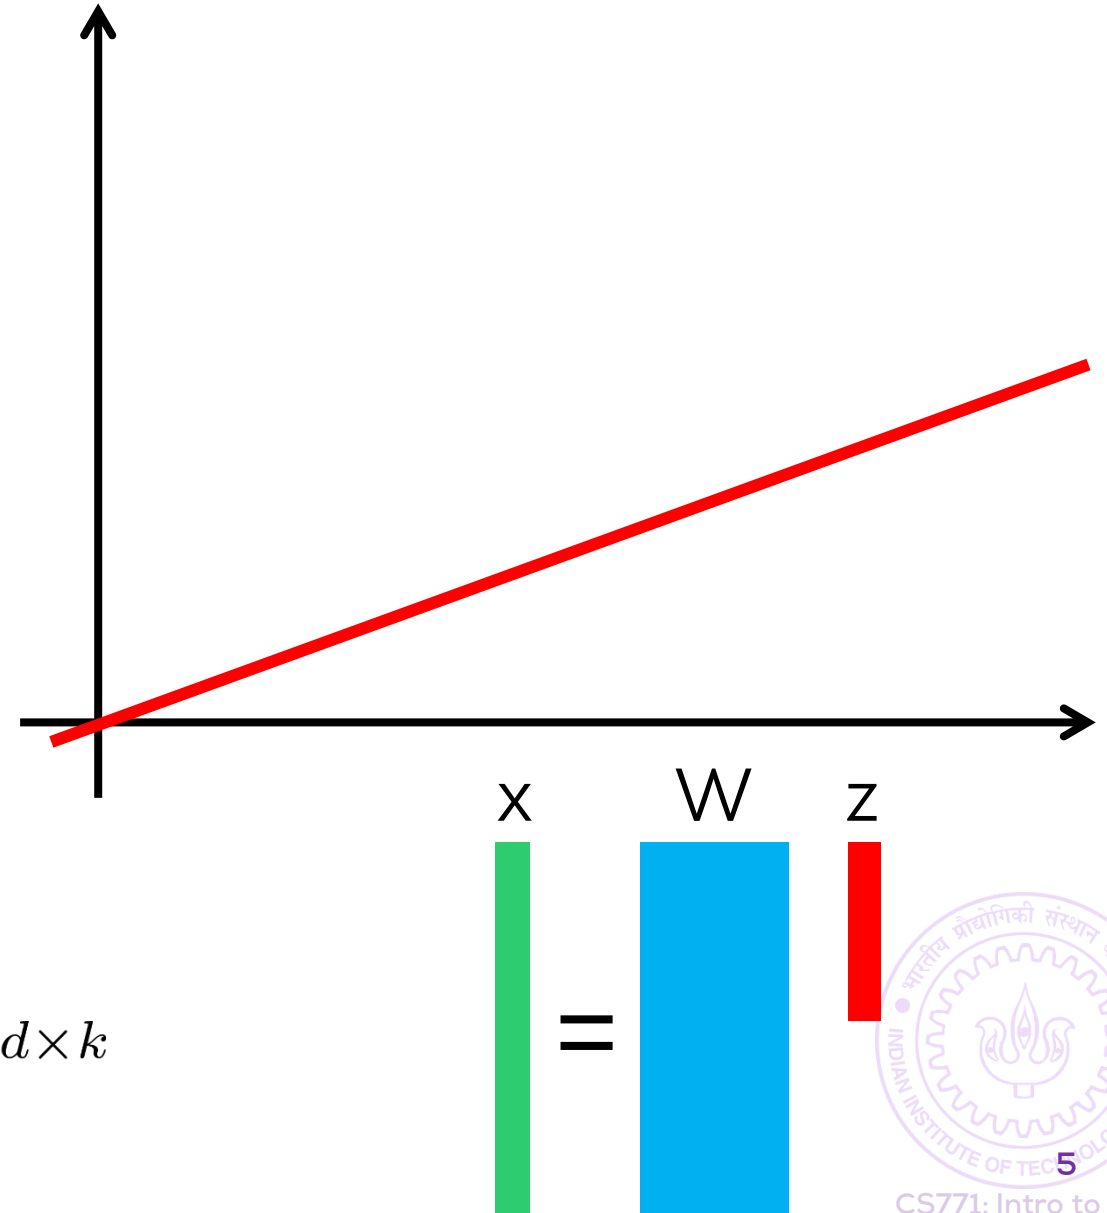
# Low-dimensional Structure in Data



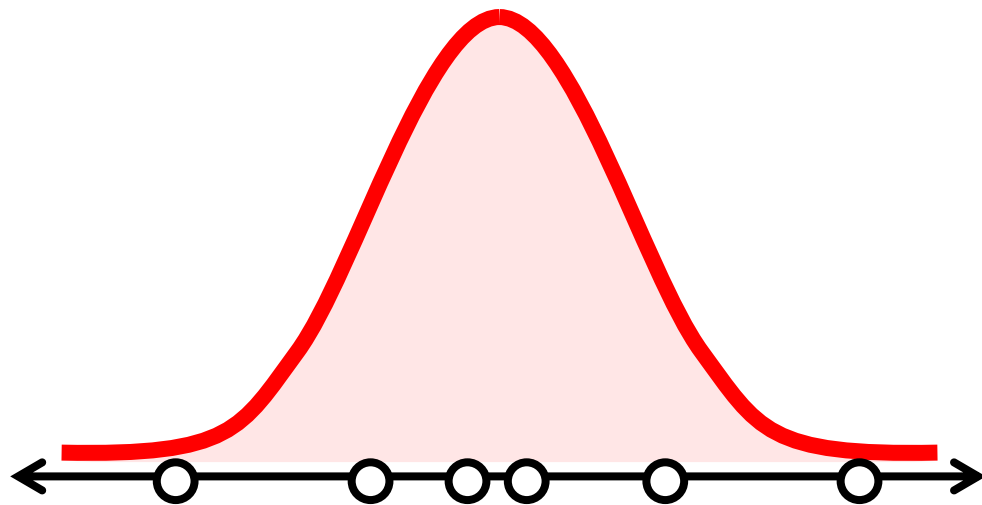
$$\mathbf{z}^i \sim \mathcal{N}(\mathbf{0}, I_k)$$

$$\mathbf{x}^i = W \mathbf{z}^i$$

$$W \in \mathbb{R}^{d \times k}$$



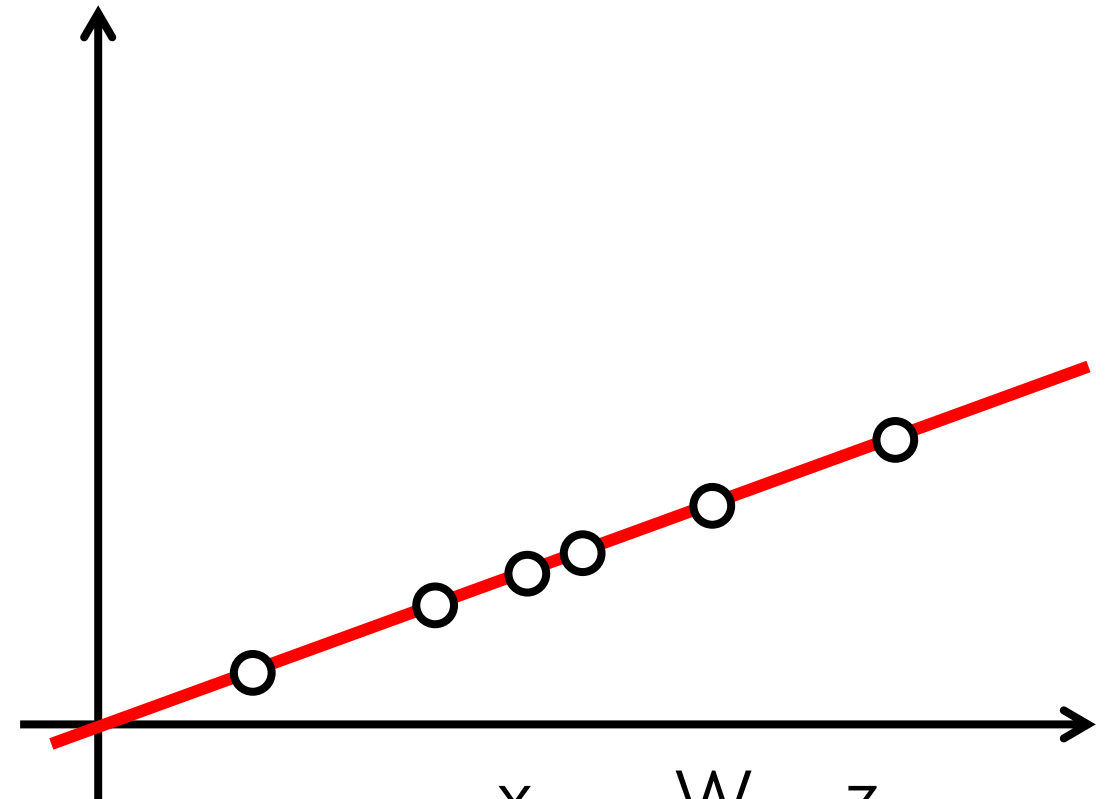
# Low-dimensional Structure in Data



$$\mathbf{z}^i \sim \mathcal{N}(\mathbf{0}, I_k)$$

$$\mathbf{x}^i = W \mathbf{z}^i$$

$$W \in \mathbb{R}^{d \times k}$$



$\mathbf{x}$



$=$

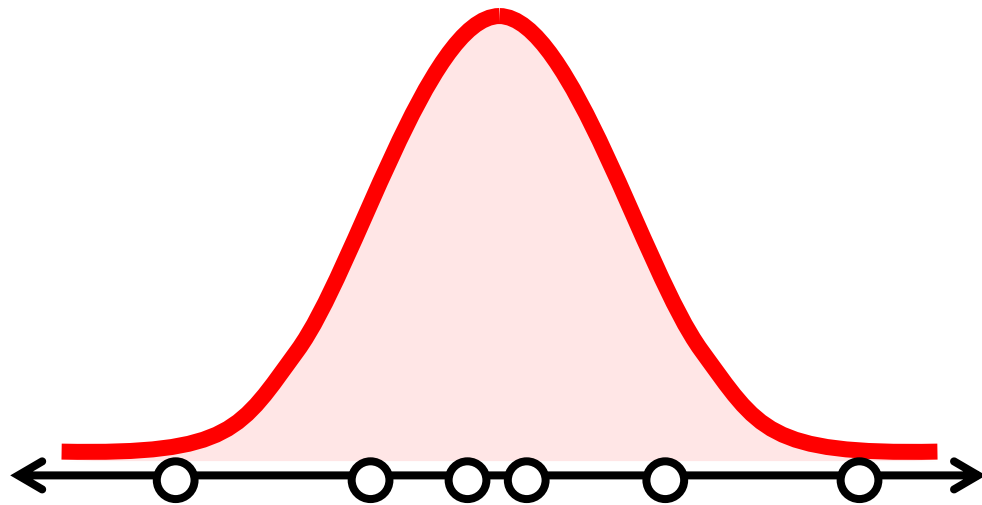
$W$



$\mathbf{z}$



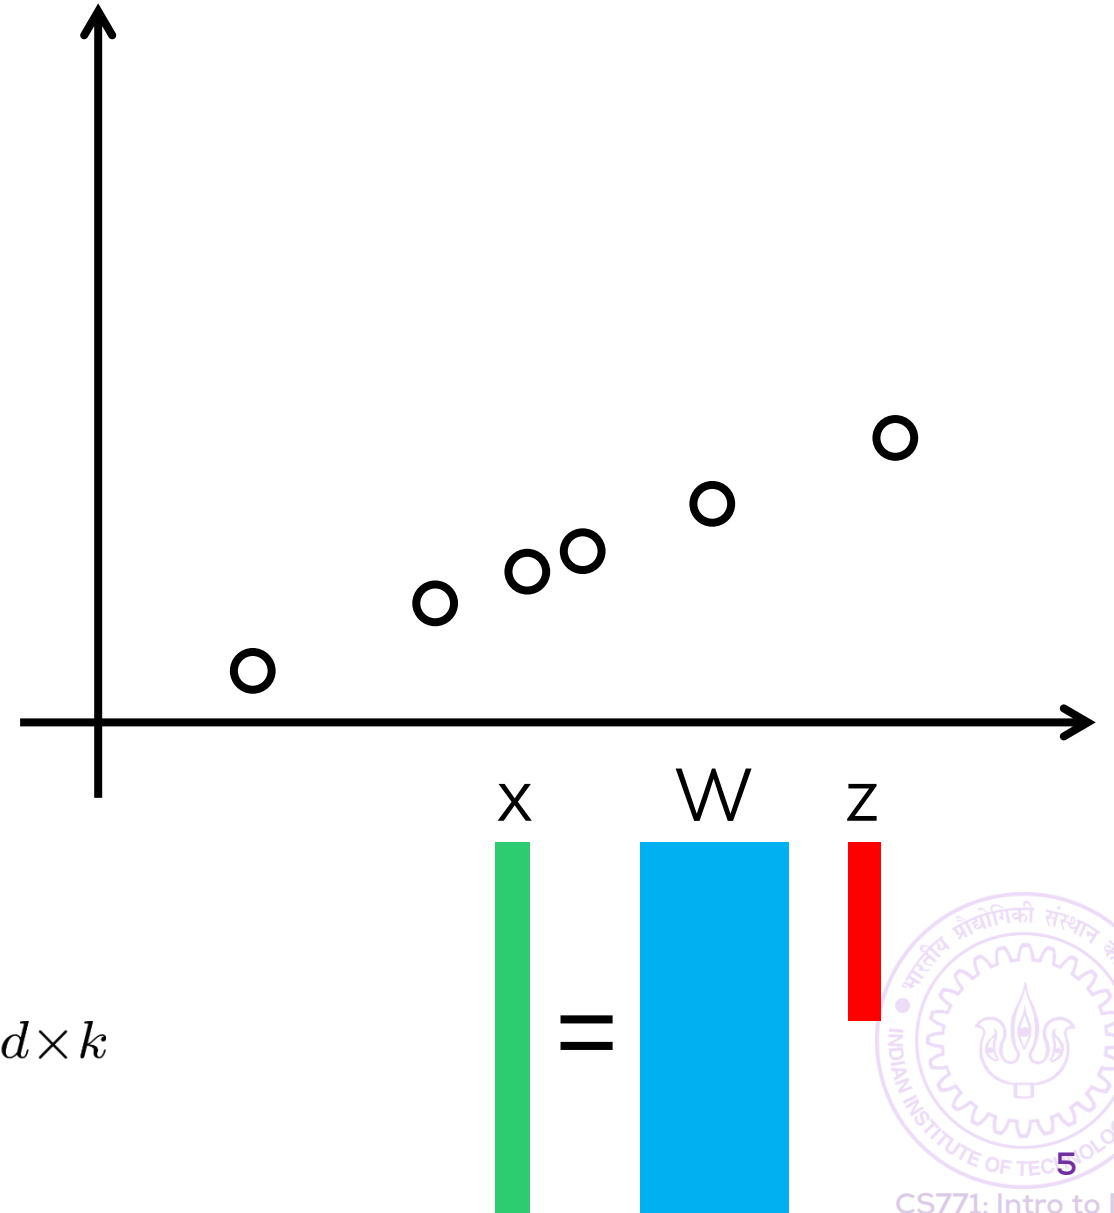
# Low-dimensional Structure in Data



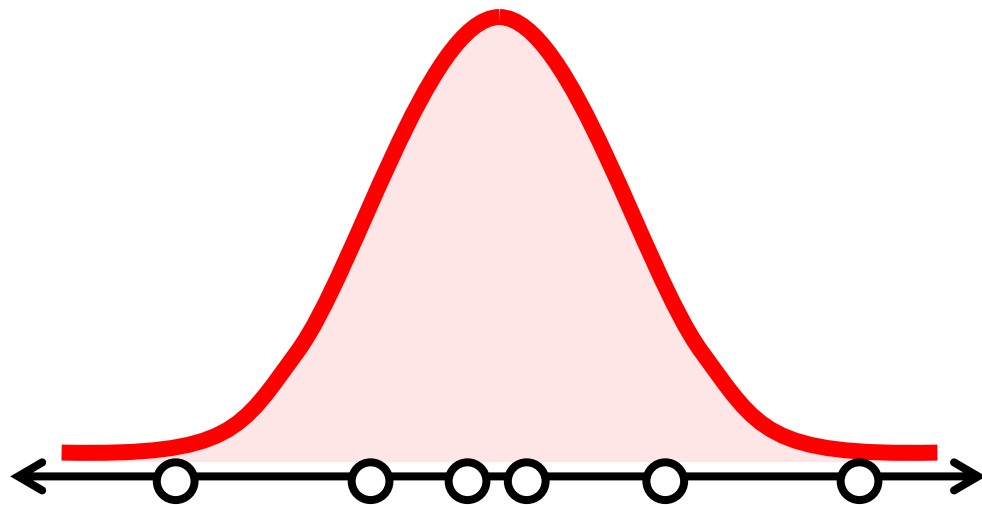
$$\mathbf{z}^i \sim \mathcal{N}(\mathbf{0}, I_k)$$

$$\mathbf{x}^i = W \mathbf{z}^i$$

$$W \in \mathbb{R}^{d \times k}$$



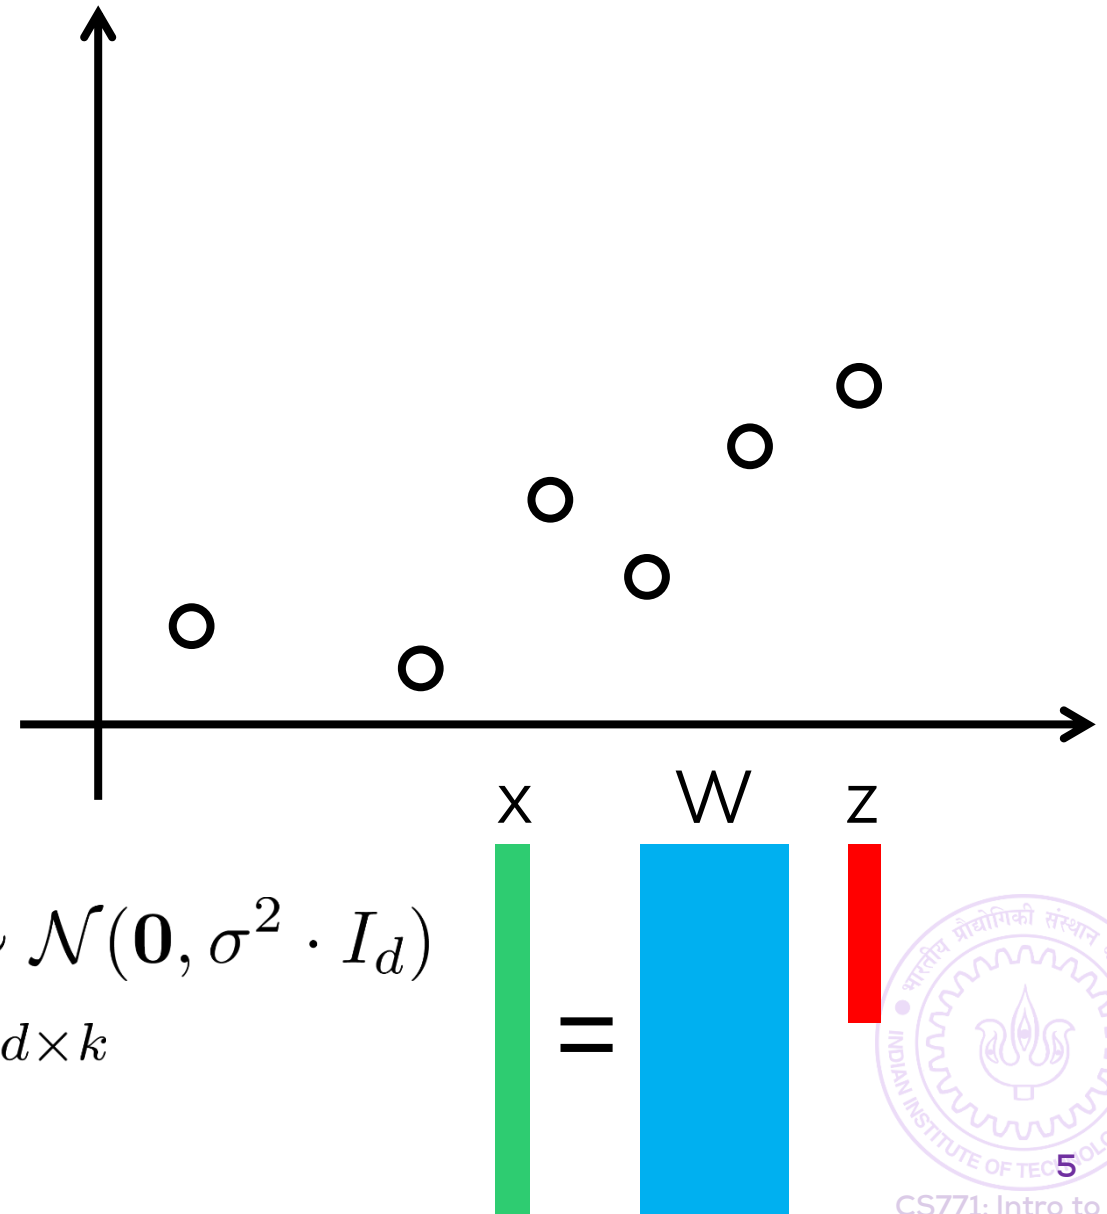
# Low-dimensional Structure in Data



$$\mathbf{z}^i \sim \mathcal{N}(\mathbf{0}, I_k)$$

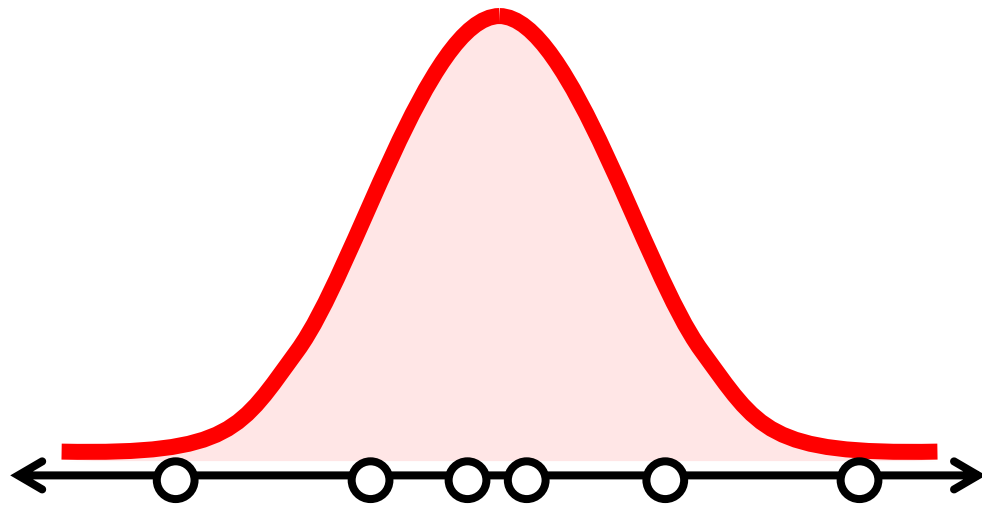
$$\mathbf{x}^i = W\mathbf{z}^i + \boldsymbol{\epsilon}^i, \quad \boldsymbol{\epsilon}^i \sim \mathcal{N}(\mathbf{0}, \sigma^2 \cdot I_d)$$

$$W \in \mathbb{R}^{d \times k}$$





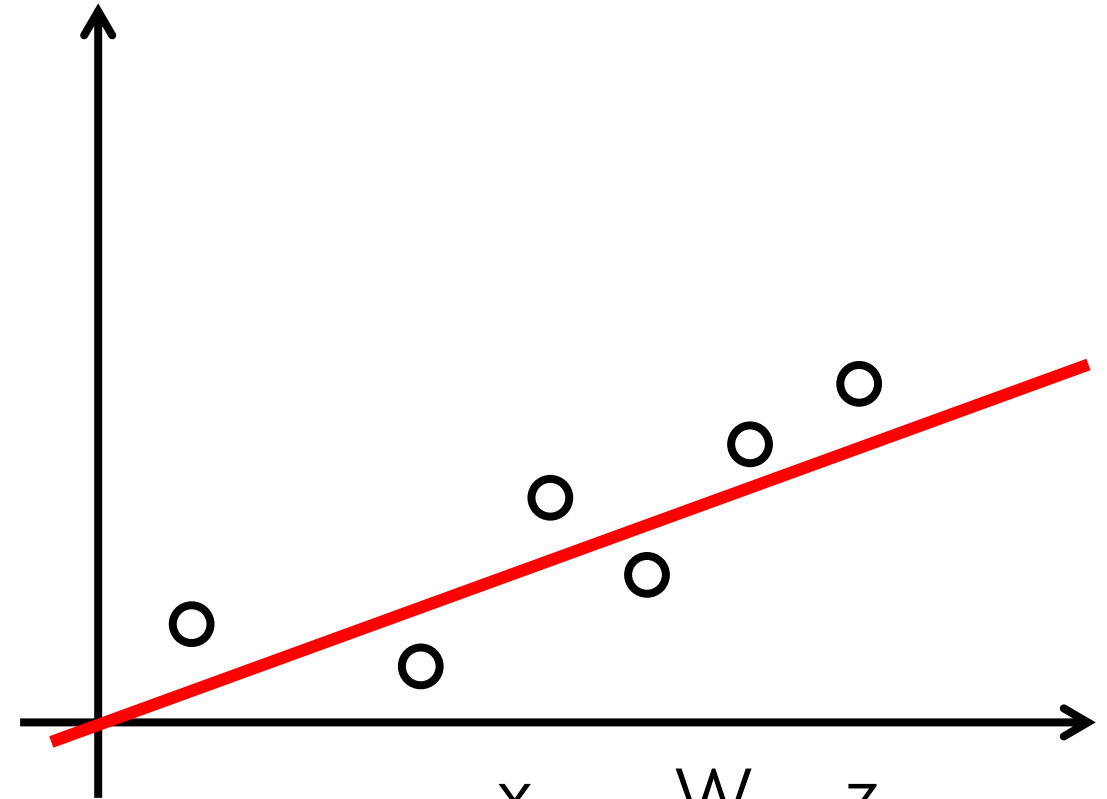
# Low-dimensional Structure in Data



$$\mathbf{z}^i \sim \mathcal{N}(\mathbf{0}, I_k)$$

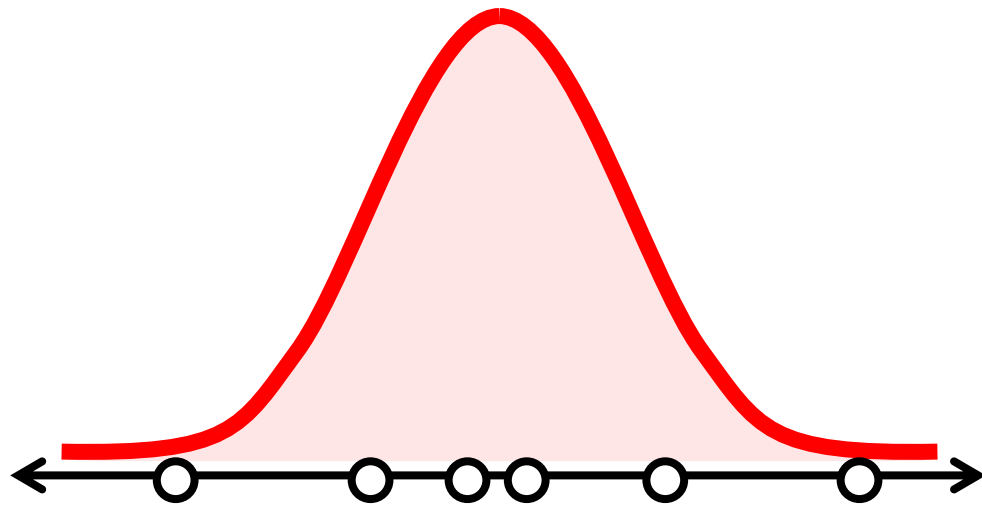
$$\mathbf{x}^i = W\mathbf{z}^i + \epsilon^i, \quad \epsilon^i \sim \mathcal{N}(\mathbf{0}, \sigma^2 \cdot I_d)$$

$$W \in \mathbb{R}^{d \times k}$$



$$\begin{array}{c} \mathbf{x} \\ \text{green bar} \end{array} = \begin{array}{c} W \\ \text{blue bar} \end{array} \begin{array}{c} \mathbf{z} \\ \text{red bar} \end{array}$$

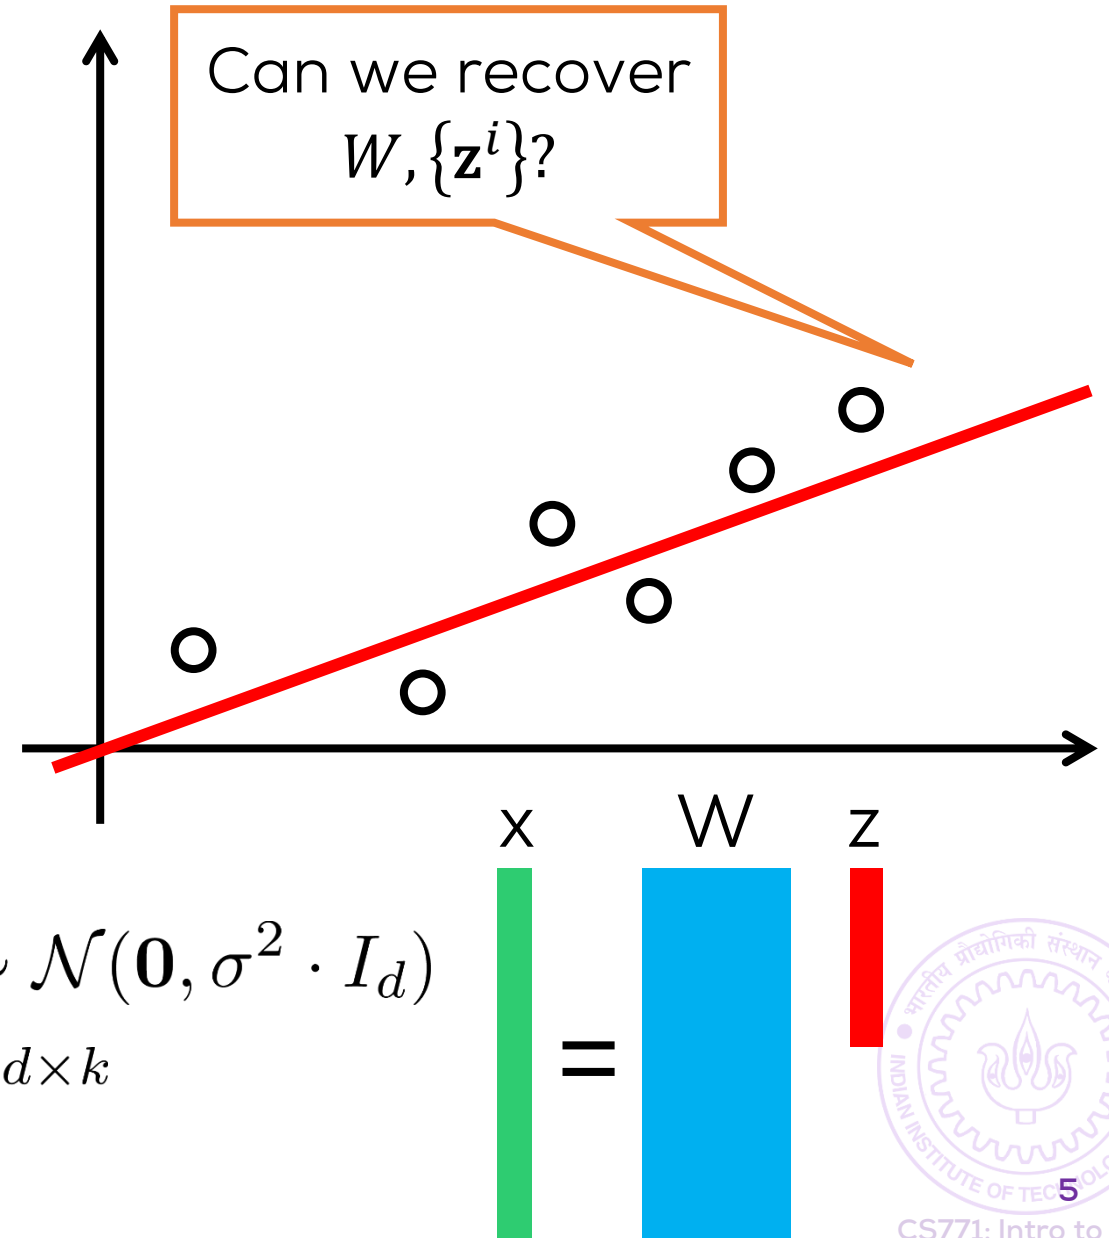
# Low-dimensional Structure in Data



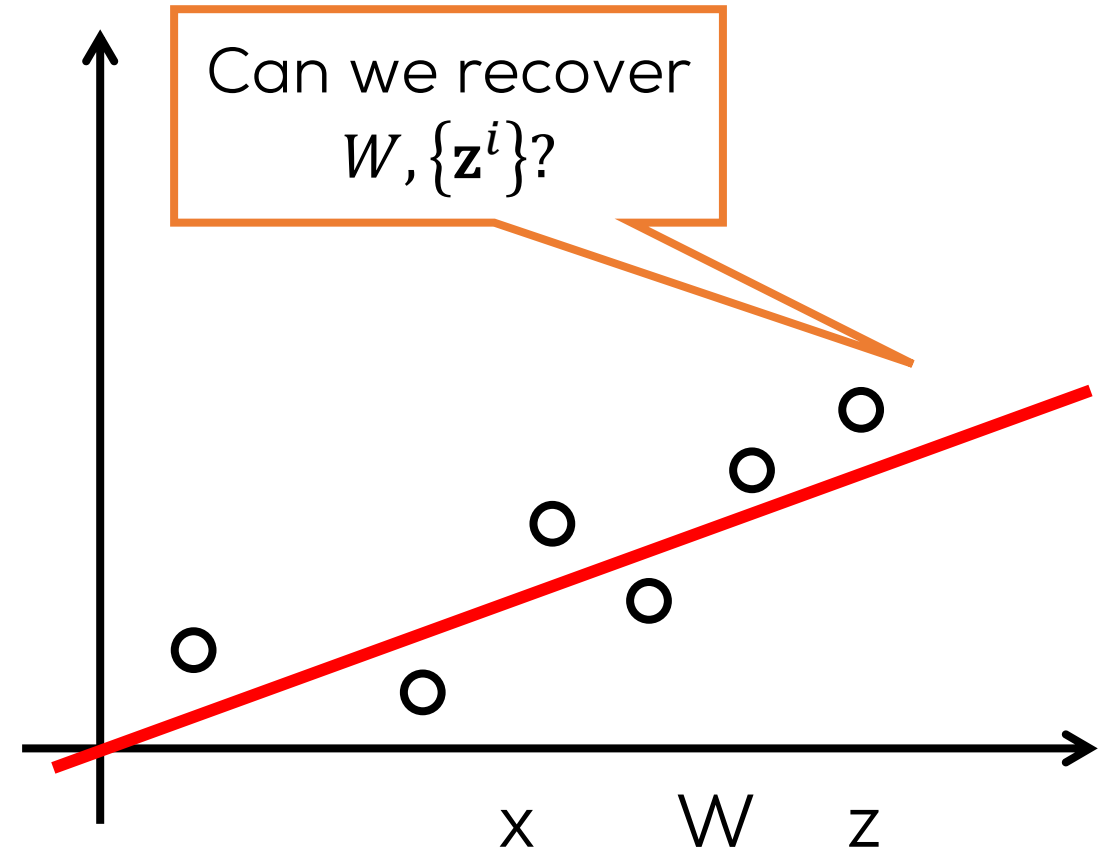
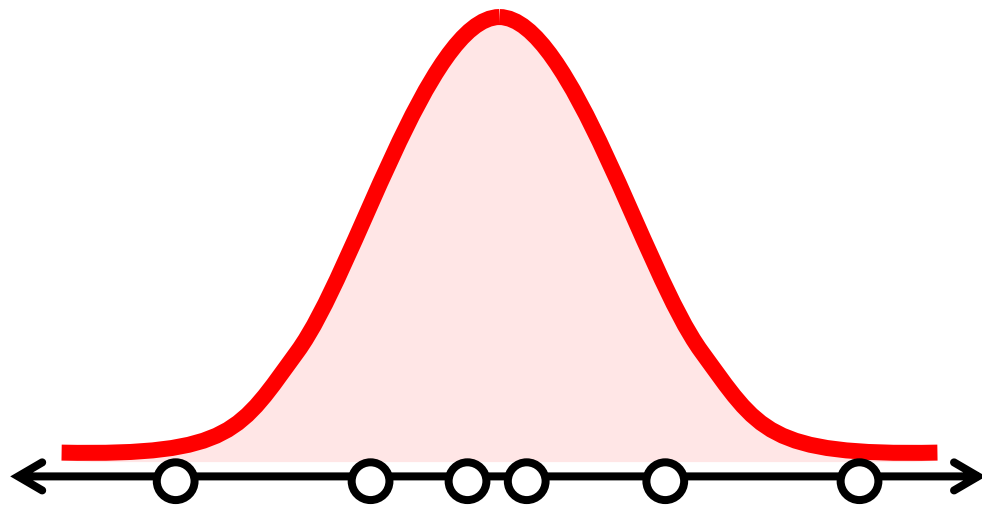
$$\mathbf{z}^i \sim \mathcal{N}(\mathbf{0}, I_k)$$

$$\mathbf{x}^i = W\mathbf{z}^i + \boldsymbol{\epsilon}^i, \quad \boldsymbol{\epsilon}^i \sim \mathcal{N}(\mathbf{0}, \sigma^2 \cdot I_d)$$

$$W \in \mathbb{R}^{d \times k}$$



# Low-dimensional Structure in Data



$$\mathbf{z}^i \sim \mathcal{N}(\mathbf{0}, I_k) \quad \mathbf{x}^i = W \mathbf{z}^i + \epsilon^i, \quad \epsilon^i \sim \mathcal{N}(\mathbf{0}, \sigma^2 \cdot I_d)$$

Dictionary/Factor  
Loading matrix

$$W \in \mathbb{R}^{d \times k}$$

# Isn't this exactly Linear Regression?

- No, subtle differences exist
- If we write things in the same notation, then

## Linear Regression

- $\mathbf{z}^i \in \mathbb{R}^k$ ,
- $y^i = \langle \mathbf{w}^*, \mathbf{z}^i \rangle + \epsilon^i$
- $\mathbf{w}^* \in \mathbb{R}^k$
- $\epsilon^i \sim \mathcal{N}(0, \sigma^2) \in \mathbb{R}$
- Observed data  $\mathbf{x}^i = (\mathbf{z}^i, y^i) \in \mathbb{R}^{k+1}$
- In linear regression,  $\mathbf{z}^i$  is visible, in low-rank data it is latent!

## Low-rank Modelling

- $\mathbf{z}^i \in \mathbb{R}^k$ ,
- $\mathbf{y}^i = W\mathbf{z}^i + \epsilon^i$
- $W \in \mathbb{R}^{d \times k}$
- $\epsilon^i \sim \mathcal{N}(\mathbf{0}, \sigma^2 \cdot I_d) \in \mathbb{R}^d$
- Observed data  $\mathbf{x}^i = \mathbf{y}^i \in \mathbb{R}^d$

# Applications

- Space savings: store  $k$ -dim  $\mathbf{z}^i$  instead of  $d$ -dim  $\mathbf{x}^i$ ,  $k \ll d$
- Discover hidden structure in data:  $W$  captures structure in data

# Applications

- Space savings: store  $k$ -dim  $\mathbf{z}^i$  instead of  $d$ -dim  $\mathbf{x}^i$ ,  $k \ll d$
- Discover hidden structure in data:  $W$  captures structure in data

Original Collection of Images



Credits: Piyush Rai, CS771, 2016-17-I





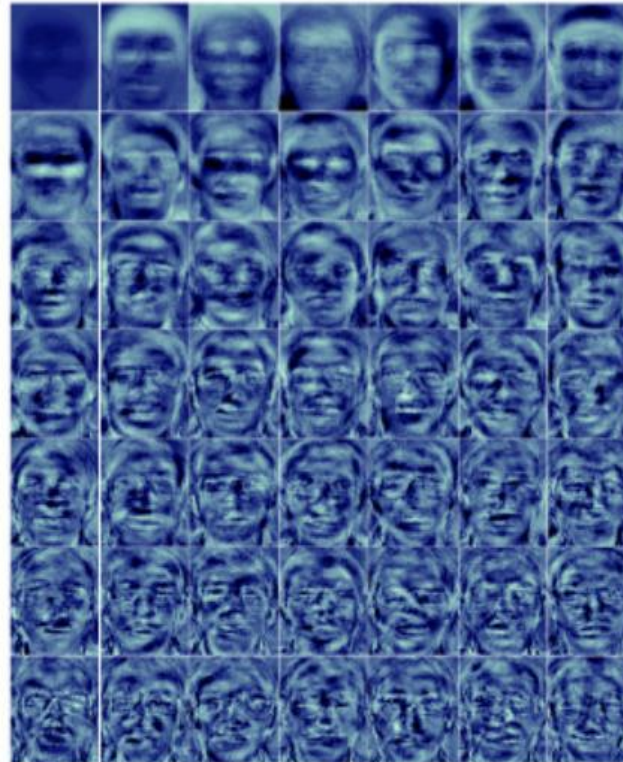
# Applications

- Space savings: store  $k$ -dim  $\mathbf{z}^i$  instead of  $d$ -dim  $\mathbf{x}^i$ ,  $k \ll d$
- Discover hidden structure in data:  $W$  captures structure in data

Original Collection of Images



K=49 Eigenvectors  
("eigenfaces") learned  
by PCA on this data



[7-1]





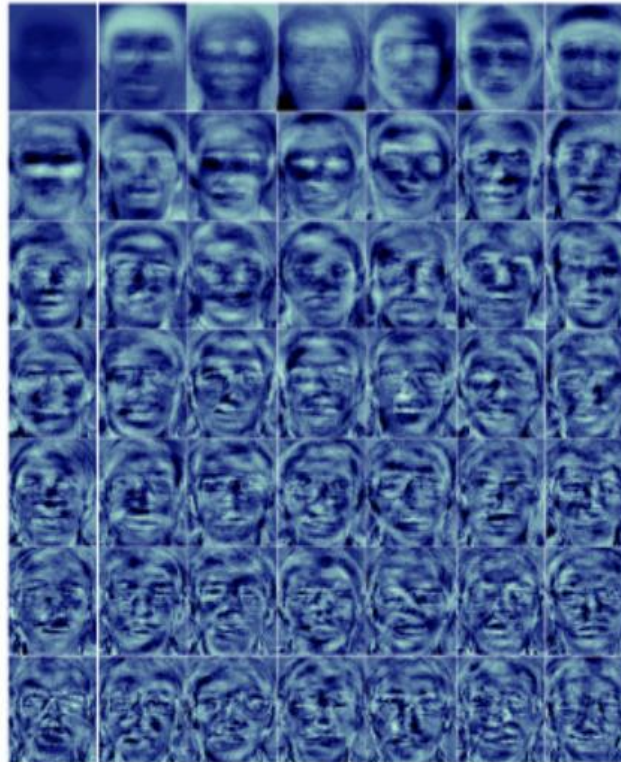
# Applications

- Space savings: store  $k$ -dim  $\mathbf{z}^i$  instead of  $d$ -dim  $\mathbf{x}^i$ ,  $k \ll d$
- Discover hidden structure in data:  $W$  captures structure in data

Original Collection of Images



K=49 Eigenvectors  
("eigenfaces") learned  
by PCA on this data



Each image's reconstructed version



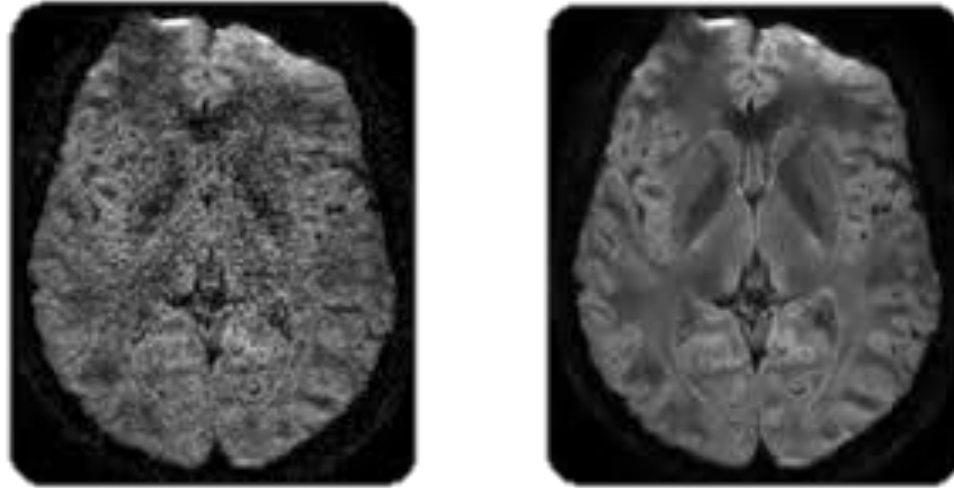


# Applications

- Noise removal:  $\mathbf{z}^i$  contains all the useful info, rest is noise

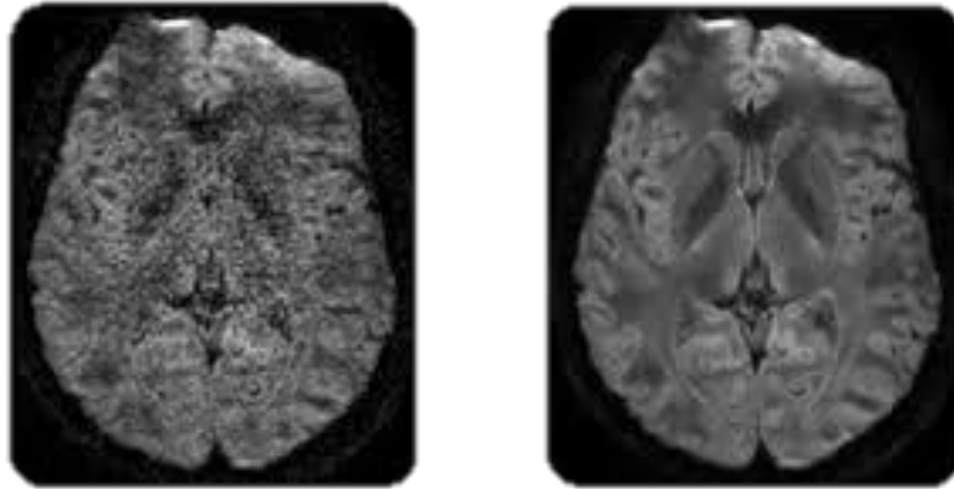
# Applications

- Noise removal:  $\mathbf{z}^i$  contains all the useful info, rest is noise



# Applications

- Noise removal:  $\mathbf{z}^i$  contains all the useful info, rest is noise



=



+



# Modelling Low-rank Data

- As discussed,  $\mathbf{z}^i \sim \mathcal{N}(\mathbf{0}, I_k)$
- As discussed  $\mathbf{x}^i | \mathbf{z}^i \sim \mathcal{N}(W\mathbf{z}^i, \sigma^2 \cdot I_d)$
- Not the only possible choice - others possible - Factor Analysis
- Things are not that bad here

$$\mathbb{P}[\mathbf{x}^i | \sigma, W] = \int_{\mathbf{z}} \mathbb{P}[\mathbf{x}^i | \mathbf{z}, \sigma, W] \cdot \mathbb{P}[\mathbf{z}] d\mathbf{z} = \mathcal{N}(0, \sigma^2 \cdot I_d + WW^\top)$$

- Note:  $\mathbb{P}[\mathbf{z} | \sigma, W] = \mathbb{P}[\mathbf{z}]$  by our definition
- Hmm ... so  $\mathbb{P}[\mathbf{x}^i | \sigma, W] = \mathcal{N}(0, \Sigma)$  where  $\Sigma = \sigma^2 \cdot I_d + WW^\top$
- But I know how to estimate  $\Sigma$  give many samples of  $\mathbf{x}$

# Approach 1

## Direct Estimation

Sept 15, 2017



# Direct Estimation

- If we have  $\mathbf{x} \sim \mathcal{N}(0, \Sigma)$ , then given many (many) samples  $\mathbf{x}^i$

$$\hat{\Sigma} = \frac{1}{n} \sum_{i=1}^n \mathbf{x}^i (\mathbf{x}^i)^\top$$

- So done ???
- Yeah ... No...
- How do we extract  $\sigma, W$  from  $\hat{\Sigma}$ ? (Remember  $\Sigma = \sigma^2 \cdot I_d + WW^\top$ )
- More importantly,  $\Sigma$  has  $d^2$  parameters in it ( $\Sigma \in \mathbb{R}^{d \times d}$ )
- To estimate it reliably, will need  $n \approx d^2$  samples ... too much
- Moreover, there are actually only  $\approx dk + 1$  parameters ( $W \in \mathbb{R}^{d \times k}$  and  $\sigma \in \mathbb{R}$ ). Should need only  $n \approx dk$  samples

# Approach 2

MLE Estimation for  $\sigma$  and  $W$

# Elementary Matrix Algebra

## The Singular Value Decomposition Theorem

- Every real matrix  $M \in \mathbb{R}^{m \times n}$  can be decomposed as
$$M = U\Lambda V^T$$
- $U = [\mathbf{u}^1, \dots, \mathbf{u}^m] \in \mathbb{R}^{m \times m}$  is an orthonormal matrix  $UU^T = U^T U = I$
- $\langle \mathbf{u}^i, \mathbf{u}^j \rangle = 1$  if  $i = j$ , 0 otherwise
- Columns of  $U$  are the *left singular vectors* of  $M$
- $V = [\mathbf{v}^1, \dots, \mathbf{v}^m] \in \mathbb{R}^{n \times n}$  is an orthonormal matrix  $VV^T = V^T V = I$
- Columns of  $V$  are the *right singular vectors* of  $M$
- $\Lambda = \text{diag}(\lambda_1, \dots, \lambda_{\min(m,n)}) \in \mathbb{R}_+^{m \times n}$  is a diagonal matrix
- We order  $\lambda_1 \geq \lambda_2 \geq \dots$
- Diagonal entries of  $\Lambda$  are the *singular values* of  $M$

$\lambda_1$	0	0	0
0	$\lambda_2$	0	0
0	0	$\lambda_3$	0



# Elementary Matrix Algebra

## The Singular Value Decomposition Theorem

- Every real matrix  $M \in \mathbb{R}^{m \times n}$  can be decomposed as

$$M = U\Lambda V^T = \sum_{i=1}^{\min(m,n)} \lambda_i \cdot \mathbf{u}^i (\mathbf{v}^i)^T$$

- For all  $i = 1, \dots, \min(m, n)$ ,  $M\mathbf{v}^i = \lambda_i \mathbf{u}^i$
- Why? Because  $\langle \mathbf{v}^i, \mathbf{v}^j \rangle = 1$  if  $i = j$ , 0 otherwise, will be very useful
- $U$  forms a basis for  $\mathbb{R}^m$
- Every vector  $\mathbf{x} \in \mathbb{R}^m$  can be written as  $\mathbf{x} = \sum_{i=1}^n \alpha_i \mathbf{u}^i$
- $\alpha_i$  can be (uniquely) found as  $\alpha_i = \langle \mathbf{x}, \mathbf{u}^i \rangle$
- $V$  similarly forms a basis for  $\mathbb{R}^n$



# Elementary Matrix Algebra

- If the matrix  $M \in \mathbb{R}^{m \times m}$  is symmetric (and hence square) then we can instead write the matrix as

$$M = U\Lambda U^\top = \sum_{i=1}^m \lambda_i \cdot \mathbf{u}^i (\mathbf{u}^i)^\top$$

- Columns of  $U$  are the *eigenvectors* of  $M$
- Diagonal entries of  $\Lambda$  are the *eigenvalues* of  $M$  (they are real)
- If all eigenvalues are  $\geq 0$ , then the matrix is called positive semi-definite (PSD)
- $\sigma^2 \cdot I$  is PSD, matrices of the form  $M = XX^\top$  or  $X^\top X$  are PSD
- If  $A, B$  are PSD then  $A + B$  is PSD too!

# MLE Estimation

- Given samples  $\mathbf{x}^1, \mathbf{x}^2, \dots, \mathbf{x}^n$  generated from  $\mathcal{N}(0, \sigma^2 \cdot I_d + WW^\top)$   
$$\log \mathbb{P}[X | W, \sigma] = \frac{n}{2} (d \log 2\pi + \log |C| + \text{tr}(C^{-1}S))$$
where  $C = WW^\top + \sigma^2 \cdot I_d$ , and  $S = \frac{1}{n} \sum_{i=1}^n \mathbf{x}^i (\mathbf{x}^i)^\top$
- $\text{tr}(M) = \sum_i M_{i,i}$
- For any  $A, B \in \mathbb{R}^{m \times n}$ ,  $\text{tr}(A^\top B) = \sum_{i,j} A_{i,j} B_{i,j} = \text{tr}(B^\top A)$
- Let  $S = U\Lambda U^\top$  be the eigen-decomposition of  $S$
- Alternately, let  $X = U\sqrt{\Lambda}V^\top$  be the singular decomposition of  $X$
- $\Lambda = \text{diag}(\lambda_1, \dots, \lambda_d)$ ,  $\lambda_i \geq 0$  (Why?),  $\sqrt{\Lambda} = \text{diag}(\sqrt{\lambda_1}, \dots, \sqrt{\lambda_d})$  (notation)
- In general if  $A = \text{diag}(a_1, \dots, a_d)$ , then  $\sqrt{A} = \text{diag}(\sqrt{a_1}, \dots, \sqrt{a_d})$

# MLE Estimation

- Given samples  $\mathbf{x}^1, \mathbf{x}^2, \dots, \mathbf{x}^n$  generated from  $\mathcal{N}(0, \sigma^2 \cdot I_d + WW^\top)$

$$\log \mathbb{P}[X | W, \sigma] = \frac{n}{2} (d \log 2\pi + \log |C| + \text{tr}(C^{-1}S))$$

where  $C = WW^\top + \sigma^2 \cdot I_d$ , and  $S = \frac{1}{n} \sum_{i=1}^n \mathbf{x}^i (\mathbf{x}^i)^\top$

- Let  $S = U\Lambda U^\top$  be the eigen-decomposition of  $S$

- $\hat{\sigma}_{\text{MLE}} = \frac{1}{d-k} \sum_{j=k+1}^d \lambda_j$

- Remember, we order  $\lambda_1 \geq \lambda_2 \geq \dots$

- $\hat{W}_{\text{MLE}} = U_k \sqrt{\Lambda_k - \hat{\sigma}_{\text{MLE}}^2 \cdot I}$

- where  $U_k = [u^1, \dots, u^k]$  and  $\Lambda_k = [\lambda_1, \dots, \lambda_k]$

- Top  $k$  eigenvalues and eigenvectors

# Principal Component Analysis

Noiseless dimensionality reduction

# The PCA estimate

- Let  $\sigma = 0$ , then the MLE looks like (no need to estimate  $\sigma$ )

$$\hat{W}_{\text{MLE}} = U_k \sqrt{\Lambda_k}$$

- So we need to find the  $k$  leading eigenvalues/vectors of  $S$
- Recall  $S = \frac{1}{n} \sum_{i=1}^n \mathbf{x}^i (\mathbf{x}^i)^\top$
- In general it takes  $O(d^3)$  time to find all  $d$  eigenvectors/values
- Much faster method to find top  $k$  in  $O(d^2 k)$  time

# The PCA estimate

Beautiful FA interpretation  
– next time!!

- Let  $\sigma = 0$ , then the MLE looks like (no need to estimate  $\sigma$ )

$$\hat{W}_{\text{MLE}} = U_k \sqrt{\Lambda_k}$$

- So we need to find the  $k$  leading eigenvalues/vectors of  $S$
- Recall  $S = \frac{1}{n} \sum_{i=1}^n \mathbf{x}^i (\mathbf{x}^i)^\top$
- In general it takes  $O(d^3)$  time to find all  $d$  eigenvectors/values
- Much faster method to find top  $k$  in  $O(d^2 k)$  time

# The Power Method

- Let  $S = U\Lambda U^\top = \sum_{i=1}^d \lambda_i \mathbf{u}^i (\mathbf{u}^i)^\top$  and  $\lambda_1 > \lambda_2 \geq \dots$  (strict separation)
- The above condition can be relaxed to handle cases  $\lambda_1 = \lambda_2$
- But makes life more complicated
- Key idea:  $U$  forms a basis for  $\mathbb{R}^d$ , every  $\mathbf{x} \in \mathbb{R}^d$  is  $\mathbf{x} = \sum_{i=1}^d \alpha_i \mathbf{u}^i$
- Assume that we have a vector  $\mathbf{x}$  so that  $\alpha_i = \frac{1}{\sqrt{d}}$  for all  $i \in [d]$
- Then what is the vector  $S\mathbf{x}$ ?
- Since  $\langle \mathbf{u}^i, \mathbf{u}^j \rangle = 1$  if  $i = j$ , 0 otherwise

$$S\mathbf{x} = \sum_{i=1}^d \alpha_i \lambda_i \cdot \mathbf{u}^i$$

Notice  $\alpha_1 \lambda_1 > \alpha_i \lambda_i$   
for all  $i \neq 1$

Amplifies  
component along  
leading eigenvector



# The Power Method

- Continuing this way, we can show that

$$S^t \mathbf{x} = \underbrace{SSSSSS}_t \mathbf{x} = \sum_{i=1}^d \alpha_i \lambda_i^t \cdot \mathbf{u}^i$$

- Even if  $\lambda_1 = 1.01$  and  $\lambda_2 = 1.005$ , after  $t=1000$  iterations,  $\lambda_1^t > 20000$  whereas  $\lambda_2 < 150$ . Tiny differences get amplified greatly!!
- Even if  $\lambda_1 = 0.995$  and  $\lambda_2 = 0.99$ , after  $t=1000$  iterations,  $\lambda_1^t > 0.005$  whereas  $\lambda_2 < 0.00005$ . The difference is still amplified!!
- No need to have  $\alpha_i$  equal. Even if  $\alpha_1$  is smaller than other  $\alpha_i$ , soon we will have  $\alpha_1 \lambda_1^t$  much much larger than  $\alpha_i \lambda_i^t$  for  $i \neq 1$ .
- The only thing we need to be careful about is to not have  $\alpha_1 = 0$ .  
The above procedure fails if  $\alpha_1 = 0$

# The Power Method

## THE POWER METHOD

1. Matrix  $S$
2. Initialize  $\mathbf{x}^0$  randomly  $\sim \mathcal{N}(\mathbf{0}, I)$
3. For  $t = 1, 2, \dots, T$

$$\mathbf{y}^t = S\mathbf{x}^{t-1}$$

$$\mathbf{x}^t = \frac{\mathbf{y}^t}{\|\mathbf{y}^t\|_2}$$

4. Repeat until convergence
5. Return eigenvector estimate as  $\mathbf{x}^T$
6. Return eigenvalue estimate as  $\|S\mathbf{x}^T\|_2$

# The Power Method

## THE POWER METHOD

Ensures  $\alpha_1 \neq 0$   
with high  
probability

1. Matrix  $S$
2. Initialize  $\mathbf{x}^0$  randomly  $\sim \mathcal{N}(\mathbf{0}, I)$
3. For  $t = 1, 2, \dots, T$

$$\mathbf{y}^t = S\mathbf{x}^{t-1}$$

$$\mathbf{x}^t = \frac{\mathbf{y}^t}{\|\mathbf{y}^t\|_2}$$

4. Repeat until convergence
5. Return eigenvector estimate as  $\mathbf{x}^T$
6. Return eigenvalue estimate as  $\|S\mathbf{x}^T\|_2$

# The Power Method

## THE POWER METHOD

1. Matrix  $S$
2. Initialize  $\mathbf{x}^0$  randomly  $\sim \mathcal{N}(\mathbf{0}, I)$
3. For  $t = 1, 2, \dots, T$

$$\mathbf{y}^t = S\mathbf{x}^{t-1}$$
$$\mathbf{x}^t = \frac{\mathbf{y}^t}{\|\mathbf{y}^t\|_2}$$

4. Repeat until convergence
5. Return eigenvector estimate as  $\mathbf{x}^T$
6. Return eigenvalue estimate as  $\|S\mathbf{x}^T\|_2$

Ensures  $\alpha_1 \neq 0$   
with high  
probability

Takes only  
 $O(d^2)$  time

# The Power Method

## THE POWER METHOD

1. Matrix  $S$
2. Initialize  $\mathbf{x}^0$  randomly  $\sim \mathcal{N}(\mathbf{0}, I)$
3. For  $t = 1, 2, \dots, T$

$$\mathbf{y}^t = S\mathbf{x}^{t-1}$$

$$\mathbf{x}^t = \frac{\mathbf{y}^t}{\|\mathbf{y}^t\|_2}$$

4. Repeat until convergence
5. Return eigenvector estimate as  $\mathbf{x}^T$
6. Return eigenvalue estimate as  $\|S\mathbf{x}^T\|_2$

Ensures  $\alpha_1 \neq 0$   
with high  
probability

Takes only  
 $O(d^2)$  time

Eigenvectors are  
all unit norm

# The Power Method

## THE POWER METHOD

1. Matrix  $S$
2. Initialize  $\mathbf{x}^0$  randomly  $\sim \mathcal{N}(\mathbf{0}, I)$
3. For  $t = 1, 2, \dots, T$

$$\mathbf{y}^t = S\mathbf{x}^{t-1}$$

$$\mathbf{x}^t = \frac{\mathbf{y}^t}{\|\mathbf{y}^t\|_2}$$

4. Repeat until convergence
5. Return eigenvector estimate as  $\mathbf{x}^T$
6. Return eigenvalue estimate as  $\|S\mathbf{x}^T\|_2$

Ensures  $\alpha_1 \neq 0$   
with high  
probability

Takes only  
 $O(d^2)$  time

Eigenvectors are  
all unit norm

Why is this  
appropriate  
??

# The Power Method

## THE POWER METHOD

1. Matrix  $S$
2. Initialize  $\mathbf{x}^0$  randomly  $\sim \mathcal{N}(\mathbf{0}, I)$
3. For  $t = 1, 2, \dots, T$

$$\mathbf{y}^t = S\mathbf{x}^{t-1}$$

$$\mathbf{x}^t = \frac{\mathbf{y}^t}{\|\mathbf{y}^t\|_2}$$

4. Repeat until convergence
5. Return eigenvector estimate as  $\mathbf{x}^T$
6. Return eigenvalue estimate as  $\|S\mathbf{x}^T\|_2$

Ensures  $\alpha_1 \neq 0$   
with high  
probability

Takes only  
 $O(d^2)$  time

Eigenvectors are  
all unit norm

Why is this  
appropriate  
??

Overall  
 $O(d^2)$  time

# Principal Component Analysis

## THE PCA METHOD

1. Matrix  $S$
2. Initialize  $S^0 \leftarrow S$
3. For  $j = 1, \dots, k$ 
  1. Let  $(\hat{\lambda}_j, \hat{\mathbf{u}}_j) \leftarrow \text{POWER-METHOD}(S^{j-1})$
  2. Let  $S^j \leftarrow S^{j-1} - \hat{\lambda}_j \cdot \hat{\mathbf{u}}_j (\hat{\mathbf{u}}_j)^\top$
4. Return  $\hat{W}_{\text{MLE}} = [\sqrt{\hat{\lambda}_j} \cdot \hat{\mathbf{u}}_j]_{j=1, \dots, k}$



# Principal Component Analysis

## THE PCA METHOD

1. Matrix  $S$
2. Initialize  $S^0 \leftarrow S$
3. For  $j = 1, \dots, k$ 
  1. Let  $(\hat{\lambda}_j, \hat{\mathbf{u}}_j) \leftarrow \text{POWER-METHOD}(S^{j-1})$
  2. Let  $S^j \leftarrow S^{j-1} - \hat{\lambda}_j \cdot \hat{\mathbf{u}}_j (\hat{\mathbf{u}}_j)^\top$
4. Return  $\hat{W}_{\text{MLE}} = [\sqrt{\hat{\lambda}_j} \cdot \hat{\mathbf{u}}_j]_{j=1, \dots, k}$

The peeling method

# Principal Component Analysis

## THE PCA METHOD

1. Matrix  $S$
2. Initialize  $S^0 \leftarrow S$
3. For  $j = 1, \dots, k$ 
  1. Let  $(\hat{\lambda}_j, \hat{\mathbf{u}}_j) \leftarrow \text{POWER-METHOD}(S^{j-1})$
  2. Let  $S^j \leftarrow S^{j-1} - \hat{\lambda}_j \cdot \hat{\mathbf{u}}_j(\hat{\mathbf{u}}_j)^\top$
4. Return  $\hat{W}_{\text{MLE}} = [\sqrt{\hat{\lambda}_j} \cdot \hat{\mathbf{u}}_j]_{j=1, \dots, k}$

The peeling method

$$S = \lambda_1 \mathbf{u}^1 (\mathbf{u}^1)^\top + \lambda_2 \mathbf{u}^2 (\mathbf{u}^2)^\top + \lambda_3 \mathbf{u}^3 (\mathbf{u}^3)^\top + \lambda_4 \mathbf{u}^4 (\mathbf{u}^4)^\top$$

# Principal Component Analysis

## THE PCA METHOD

1. Matrix  $S$
2. Initialize  $S^0 \leftarrow S$
3. For  $j = 1, \dots, k$ 
  1. Let  $(\hat{\lambda}_j, \hat{\mathbf{u}}_j) \leftarrow \text{POWER-METHOD}(S^{j-1})$
  2. Let  $S^j \leftarrow S^{j-1} - \hat{\lambda}_j \cdot \hat{\mathbf{u}}_j (\hat{\mathbf{u}}_j)^\top$
4. Return  $\hat{W}_{\text{MLE}} = [\sqrt{\hat{\lambda}_j} \cdot \hat{\mathbf{u}}_j]_{j=1, \dots, k}$

The peeling method

$$S = \lambda_1 \mathbf{u}^1 (\mathbf{u}^1)^\top + \lambda_2 \mathbf{u}^2 (\mathbf{u}^2)^\top + \lambda_3 \mathbf{u}^3 (\mathbf{u}^3)^\top + \lambda_4 \mathbf{u}^4 (\mathbf{u}^4)^\top$$

# Principal Component Analysis

## THE PCA METHOD

1. Matrix  $S$
2. Initialize  $S^0 \leftarrow S$
3. For  $j = 1, \dots, k$ 
  1. Let  $(\hat{\lambda}_j, \hat{\mathbf{u}}_j) \leftarrow \text{POWER-METHOD}(S^{j-1})$
  2. Let  $S^j \leftarrow S^{j-1} - \hat{\lambda}_j \cdot \hat{\mathbf{u}}_j (\hat{\mathbf{u}}_j)^\top$
4. Return  $\hat{W}_{\text{MLE}} = [\sqrt{\hat{\lambda}_j} \cdot \hat{\mathbf{u}}_j]_{j=1, \dots, k}$

The peeling method

$$S = \lambda_1 \mathbf{u}^1 (\mathbf{u}^1)^\top + \lambda_2 \mathbf{u}^2 (\mathbf{u}^2)^\top + \lambda_3 \mathbf{u}^3 (\mathbf{u}^3)^\top + \lambda_4 \mathbf{u}^4 (\mathbf{u}^4)^\top$$

# Principal Component Analysis

## THE PCA METHOD

1. Matrix  $S$
2. Initialize  $S^0 \leftarrow S$
3. For  $j = 1, \dots, k$ 
  1. Let  $(\hat{\lambda}_j, \hat{\mathbf{u}}_j) \leftarrow \text{POWER-METHOD}(S^{j-1})$
  2. Let  $S^j \leftarrow S^{j-1} - \hat{\lambda}_j \cdot \hat{\mathbf{u}}_j (\hat{\mathbf{u}}_j)^\top$
4. Return  $\hat{W}_{\text{MLE}} = [\sqrt{\hat{\lambda}_j} \cdot \hat{\mathbf{u}}_j]_{j=1, \dots, k}$

The peeling method

$$S = \lambda_1 \mathbf{u}^1 (\mathbf{u}^1)^\top + \lambda_2 \mathbf{u}^2 (\mathbf{u}^2)^\top + \lambda_3 \mathbf{u}^3 (\mathbf{u}^3)^\top + \lambda_4 \mathbf{u}^4 (\mathbf{u}^4)^\top$$

# Principal Component Analysis

## THE PCA METHOD

1. Matrix  $S$
2. Initialize  $S^0 \leftarrow S$
3. For  $j = 1, \dots, k$ 
  1. Let  $(\hat{\lambda}_j, \hat{\mathbf{u}}_j) \leftarrow \text{POWER-METHOD}(S^{j-1})$
  2. Let  $S^j \leftarrow S^{j-1} - \hat{\lambda}_j \cdot \hat{\mathbf{u}}_j (\hat{\mathbf{u}}_j)^\top$
4. Return  $\hat{W}_{\text{MLE}} = [\sqrt{\hat{\lambda}_j} \cdot \hat{\mathbf{u}}_j]_{j=1}^k$

The peeling method

Some residue might still be left due to inaccurate estimation of  $\lambda_i, \mathbf{u}^i$  but usually small

$$S = \lambda_1 \mathbf{u}^1 (\mathbf{u}^1)^\top + \lambda_2 \mathbf{u}^2 (\mathbf{u}^2)^\top + \lambda_3 \mathbf{u}^3 (\mathbf{u}^3)^\top + \lambda_4 \mathbf{u}^4 (\mathbf{u}^4)^\top$$

# Principal Component Analysis

## THE PCA METHOD

1. Matrix  $S$
2. Initialize  $S^0 \leftarrow S$
3. For  $j = 1, \dots, k$ 
  1. Let  $(\hat{\lambda}_j, \hat{\mathbf{u}}_j) \leftarrow \text{POWER-METHOD}(S^{j-1})$
  2. Let  $S^j \leftarrow S^{j-1} - \hat{\lambda}_j \cdot \hat{\mathbf{u}}_j (\hat{\mathbf{u}}_j)^\top$
4. Return  $\hat{W}_{\text{MLE}} = [\sqrt{\hat{\lambda}_j} \cdot \hat{\mathbf{u}}_j]_{j=1}^k$

Overall  
 $O(d^2k)$  time

The peeling  
method

Some residue might still  
be left due to  
inaccurate estimation of  
 $\lambda_i, \mathbf{u}^i$  but usually small

$$S = \lambda_1 \mathbf{u}^1 (\mathbf{u}^1)^\top + \lambda_2 \mathbf{u}^2 (\mathbf{u}^2)^\top + \lambda_3 \mathbf{u}^3 (\mathbf{u}^3)^\top + \lambda_4 \mathbf{u}^4 (\mathbf{u}^4)^\top$$

# Probabilistic Principal Component Analysis

## THE PPCA METHOD

1. Matrix  $S$
2. Initialize  $S^0 \leftarrow S$
3. For  $j = 1, \dots, d$ 
  1. Let  $(\hat{\lambda}_j, \hat{\mathbf{u}}_j) \leftarrow \text{POWER-METHOD}(S^{j-1})$
  2. Let  $S^j \leftarrow S^{j-1} - \hat{\lambda}_j \cdot \hat{\mathbf{u}}_j (\hat{\mathbf{u}}_j)^\top$
4. Let  $\hat{\sigma}_{\text{MLE}} = \frac{1}{d-k} \sum_{j=k+1}^d \hat{\lambda}_j$
5. Return  $\hat{W}_{\text{MLE}} = \left[ \sqrt{\hat{\lambda}_j - \hat{\sigma}_{\text{MLE}}^2} \cdot \hat{\mathbf{u}}_j \right]_{j=1, \dots, k}$



# Probabilistic Principal Component Analysis

## THE PPCA METHOD

1. Matrix  $S$
2. Initialize  $S^0 \leftarrow S$
3. For  $j = 1, \dots, d$ 
  1. Let  $(\hat{\lambda}_j, \hat{\mathbf{u}}_j) \leftarrow \text{POWER-METHOD}(S^{j-1})$
  2. Let  $S^j \leftarrow S^{j-1} - \hat{\lambda}_j \cdot \hat{\mathbf{u}}_j (\hat{\mathbf{u}}_j)^\top$
4. Let  $\hat{\sigma}_{\text{MLE}}^2 = \frac{1}{d-k} \sum_{j=k+1}^d \hat{\lambda}_j$
5. Return  $\hat{W}_{\text{MLE}} = \left[ \sqrt{\hat{\lambda}_j - \hat{\sigma}_{\text{MLE}}^2} \cdot \hat{\mathbf{u}}_j \right]_{j=1, \dots, k}$

Recall that

$$\hat{W}_{\text{MLE}} = U_k \sqrt{\Lambda_k - \hat{\sigma}_{\text{MLE}}^2 \cdot I}$$

# Probabilistic Principal Component Analysis

## THE PPCA METHOD

1. Matrix  $S$
2. Initialize  $S^0 \leftarrow S$
3. For  $j = 1, \dots, d$

Recall that

$$\hat{W}_{\text{MLE}} = U_k \sqrt{\Lambda_k - \hat{\sigma}_{\text{MLE}}^2 \cdot I}$$

1. Let  $(\hat{\lambda}_j, \hat{\mathbf{u}}_j) \leftarrow \text{POWER-METHOD}(S^{j-1})$

2. Let  $S^j \leftarrow S^{j-1} - \hat{\lambda}_j \cdot \hat{\mathbf{u}}_j (\hat{\mathbf{u}}_j)^\top$

Takes  $O(d^3)$   
time ☹

4. Let  $\hat{\sigma}_{\text{MLE}} = \frac{1}{d-k} \sum_{j=k+1}^d \hat{\lambda}_j$

5. Return  $\hat{W}_{\text{MLE}} = \left[ \sqrt{\hat{\lambda}_j - \hat{\sigma}_{\text{MLE}}^2} \cdot \hat{\mathbf{u}}_j \right]_{j=1, \dots, k}$

# Probabilistic Principal Component Analysis

## THE PPCA METHOD

1. Matrix  $S$
2. Initialize  $S^0 \leftarrow S$
3. For  $j = 1, \dots, d$ 
  1. Let  $(\hat{\lambda}_j, \hat{\mathbf{u}}_j) \leftarrow \text{POWER-METHOD}(S^{j-1})$
  2. Let  $S^j \leftarrow S^{j-1} - \hat{\lambda}_j \cdot \hat{\mathbf{u}}_j (\hat{\mathbf{u}}_j)^\top$
4. Let  $\hat{\sigma}_{\text{MLE}} = \frac{1}{d-k} \sum_{j=k+1}^d \hat{\lambda}_j$
5. Return  $\hat{W}_{\text{MLE}} = \left[ \sqrt{\hat{\lambda}_j - \hat{\sigma}_{\text{MLE}}^2} \cdot \hat{\mathbf{u}}_j \right]_{j=1, \dots, k}$

Recall that

$$\hat{W}_{\text{MLE}} = U_k \sqrt{\Lambda_k - \hat{\sigma}_{\text{MLE}}^2 \cdot I}$$

Takes  $O(d^3)$   
time ☹

Can we do  
better?

# Probabilistic Principal Component Analysis

## THE PPCA METHOD

1. Matrix  $S$
2. Initialize  $S^0 \leftarrow S$
3. For  $j = 1, \dots, d$ 
  1. Let  $(\hat{\lambda}_j, \hat{\mathbf{u}}_j) \leftarrow \text{POWER-METHOD}(S^{j-1})$

Recall that

$$\hat{W}_{\text{MLE}} = U_k \sqrt{\Lambda_k - \hat{\sigma}_{\text{MLE}}^2 \cdot I}$$

2. Let  $S^j \leftarrow S^{j-1} - \hat{\lambda}_j \cdot \hat{\mathbf{u}}_j (\hat{\mathbf{u}}_j)^\top$

4. Let  $\hat{\sigma}_{\text{MLE}} = \frac{1}{d-k} \sum_{j=k+1}^d \hat{\lambda}_j$

Takes  $O(d^3)$   
time ☹

After we  
reconvene

5. Return  $\hat{W}_{\text{MLE}} = \left[ \sqrt{\hat{\lambda}_j - \hat{\sigma}_{\text{MLE}}^2} \cdot \hat{\mathbf{u}}_j \right]_{j=1, \dots, k}$

Can we do  
better?

# A Few Thoughts

- Many extensions possible
  - Factor analysis  $\boldsymbol{\epsilon}^i \sim \mathcal{N}(0, \Sigma_x)$
  - Non-centered data  $\mathbf{z}^i \sim \mathcal{N}(\boldsymbol{\mu}_z, \Sigma_z)$
  - Non-centered noise  $\mathbf{x}^i = W\mathbf{z}^i + \boldsymbol{\epsilon}^i, \boldsymbol{\epsilon}^i \sim \mathcal{N}(\boldsymbol{\mu}_\epsilon, \Sigma_x)$
- Handle missing data, as can most generative models (GMM etc)
  - $\mathbf{x}^i = [\mathbf{x}_{\text{obs}}^i, \mathbf{x}_{\text{miss}}^i], \mathbb{P}[\mathbf{x}^i] = \mathbb{P}[\mathbf{x}_{\text{miss}}^i | \mathbf{x}_{\text{obs}}^i] \cdot \mathbb{P}[\mathbf{x}_{\text{obs}}^i]$
- Mixture of PPCA? Mixture of GMMs?
- Sequential models: Kalman filters, Hidden Markov models
- Hierarchical models

# A Few Thoughts

- PPCA, PCA do not do well on data with non-linear structure

# A Few Thoughts

- PPCA, PCA do not do well on data with non-linear structure

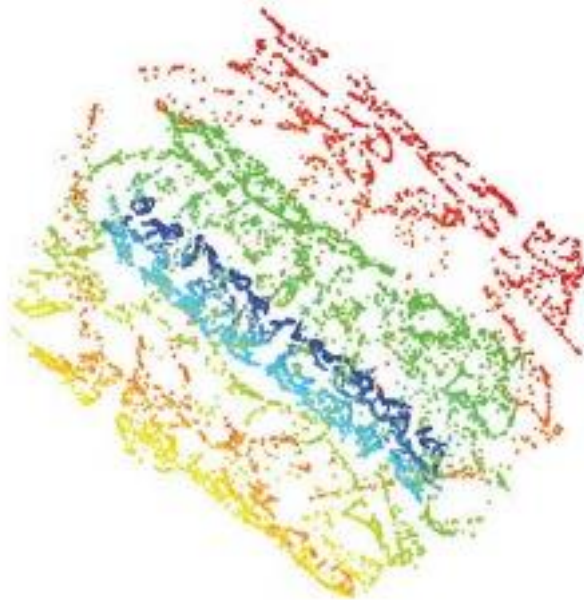
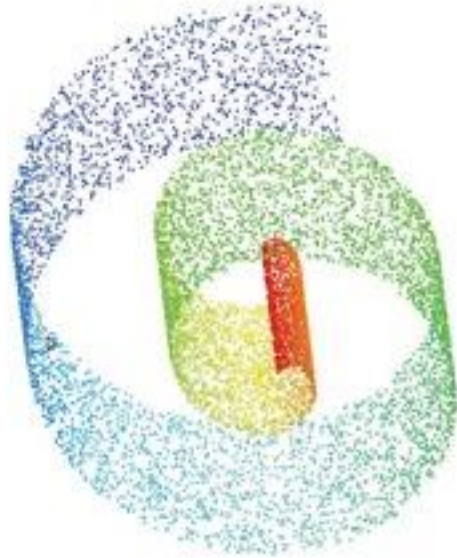
"Swiss Roll"  
data



# A Few Thoughts

- PPCA, PCA do not do well on data with non-linear structure

"Swiss Roll"  
data



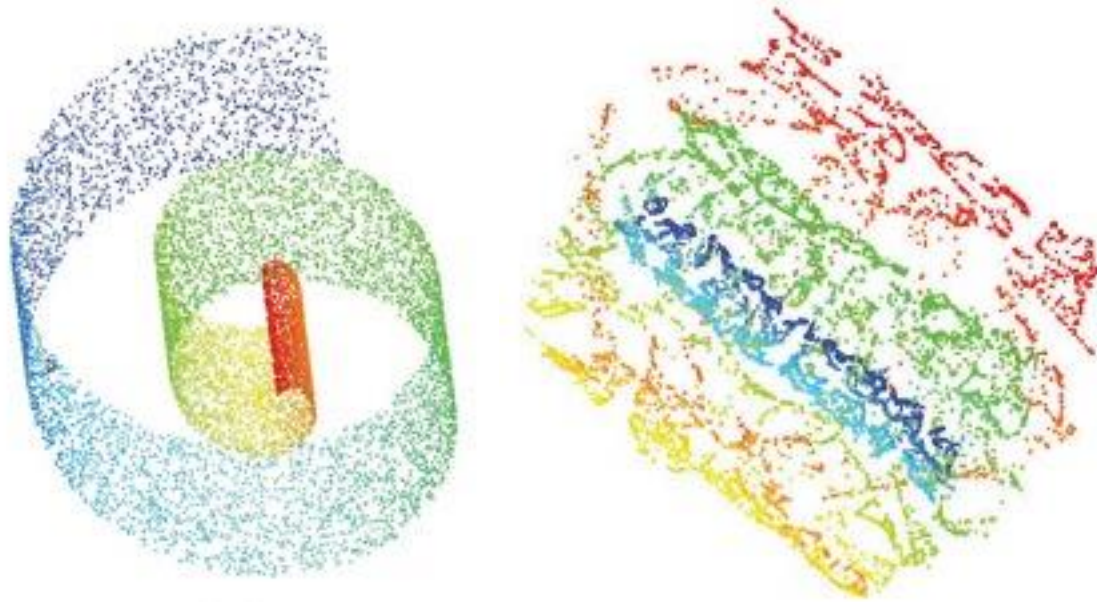
What PCA/PPCA  
will give us



# A Few Thoughts

- PPCA, PCA do not do well on data with non-linear structure

"Swiss Roll" data



What we really want

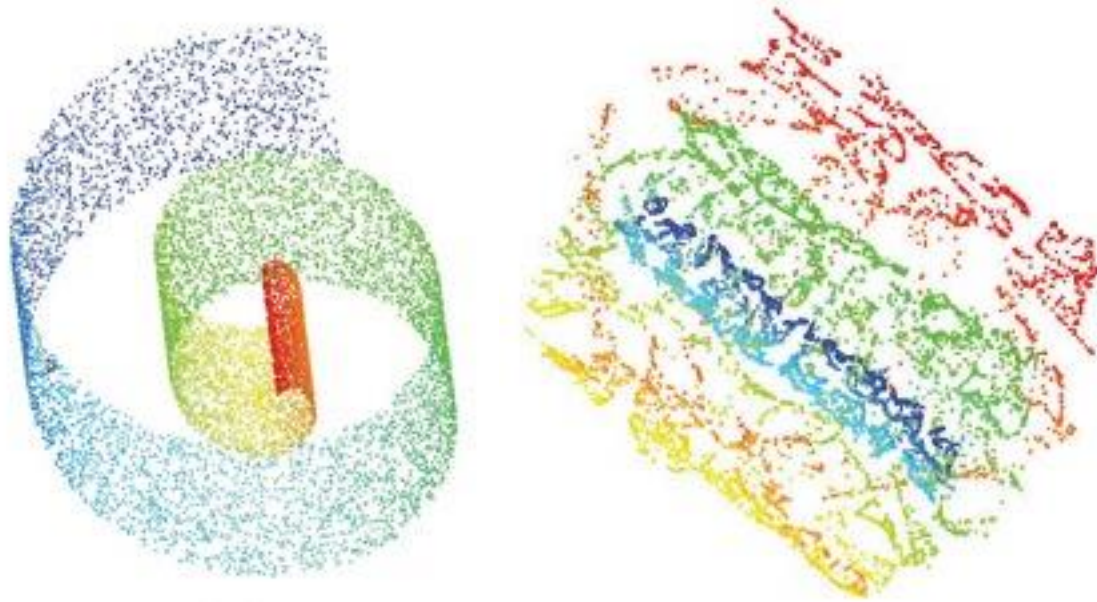


What PCA/PPCA will give us

# A Few Thoughts

- PPCA, PCA do not do well on data with non-linear structure

"Swiss Roll"  
data



What we  
really want



Non-linear  
dimensionality reduction  
Kernel PCA,  
Autoencoders

What PCA/PPCA  
will give us

# Please give your Feedback

<http://tinyurl.com/ml17-18afb>