

# Non-linear Models-V

CS771: Introduction to Machine Learning

Purushottam Kar



# Outline of discussion

- More on architecture of neural networks
- Training methods for feedforward networks

# Answering the Fan Mail

- How does the notion of similarity (kernels) over notion of distance?
  - Two sides of the same coin
  - Kernels have nice math properties that algorithms exploit
- Please relate toy examples in Lec 17 to accelerated kernel methods
  - Better idea: download data from <https://goo.gl/JXEQjr>
  - Create feature maps for Gaussian kernel <https://goo.gl/hBsX1E>
  - See the magic happen!

# Neural Networks

Oct 25, 2017




# The “neuron” in Neural Networks


Oct 25, 2017



# The “neuron” in Neural Networks



# The “neuron” in Neural Networks




$\mathbf{x}_1$

$\mathbf{x}_2$

...


$\mathbf{x}_d$

# The “neuron” in Neural Networks






# The “neuron” in Neural Networks



# The "neuron" in Neural Networks

## Input

Input layer: no i/p  
Hidden/output layer: i/p  
from other neurons




# The "neuron" in Neural Networks

## Input

Input layer: no i/p  
Hidden/output layer: i/p  
from other neurons

Some input items can  
be a constant e.g. 1



# The "neuron" in Neural Networks


## Activation

Input/output layer: id  
Hidden layer: non-linear

## Input

Input layer: no i/p  
Hidden/output layer: i/p  
from other neurons

Some input items can  
be a constant e.g. 1



# in Neural Networks

## Output

Output layer: final o/p  
Input/hidden layer: o/p  
to other neurons


## Activation

Input/output layer: id  
Hidden layer: non-linear

## Input

Input layer: no i/p  
Hidden/output layer: i/p  
from other neurons

Some input items can  
be a constant e.g. 1



# Neural Networks

## Output

Output layer: final o/p  
Input/hidden layer: o/p  
to other neurons

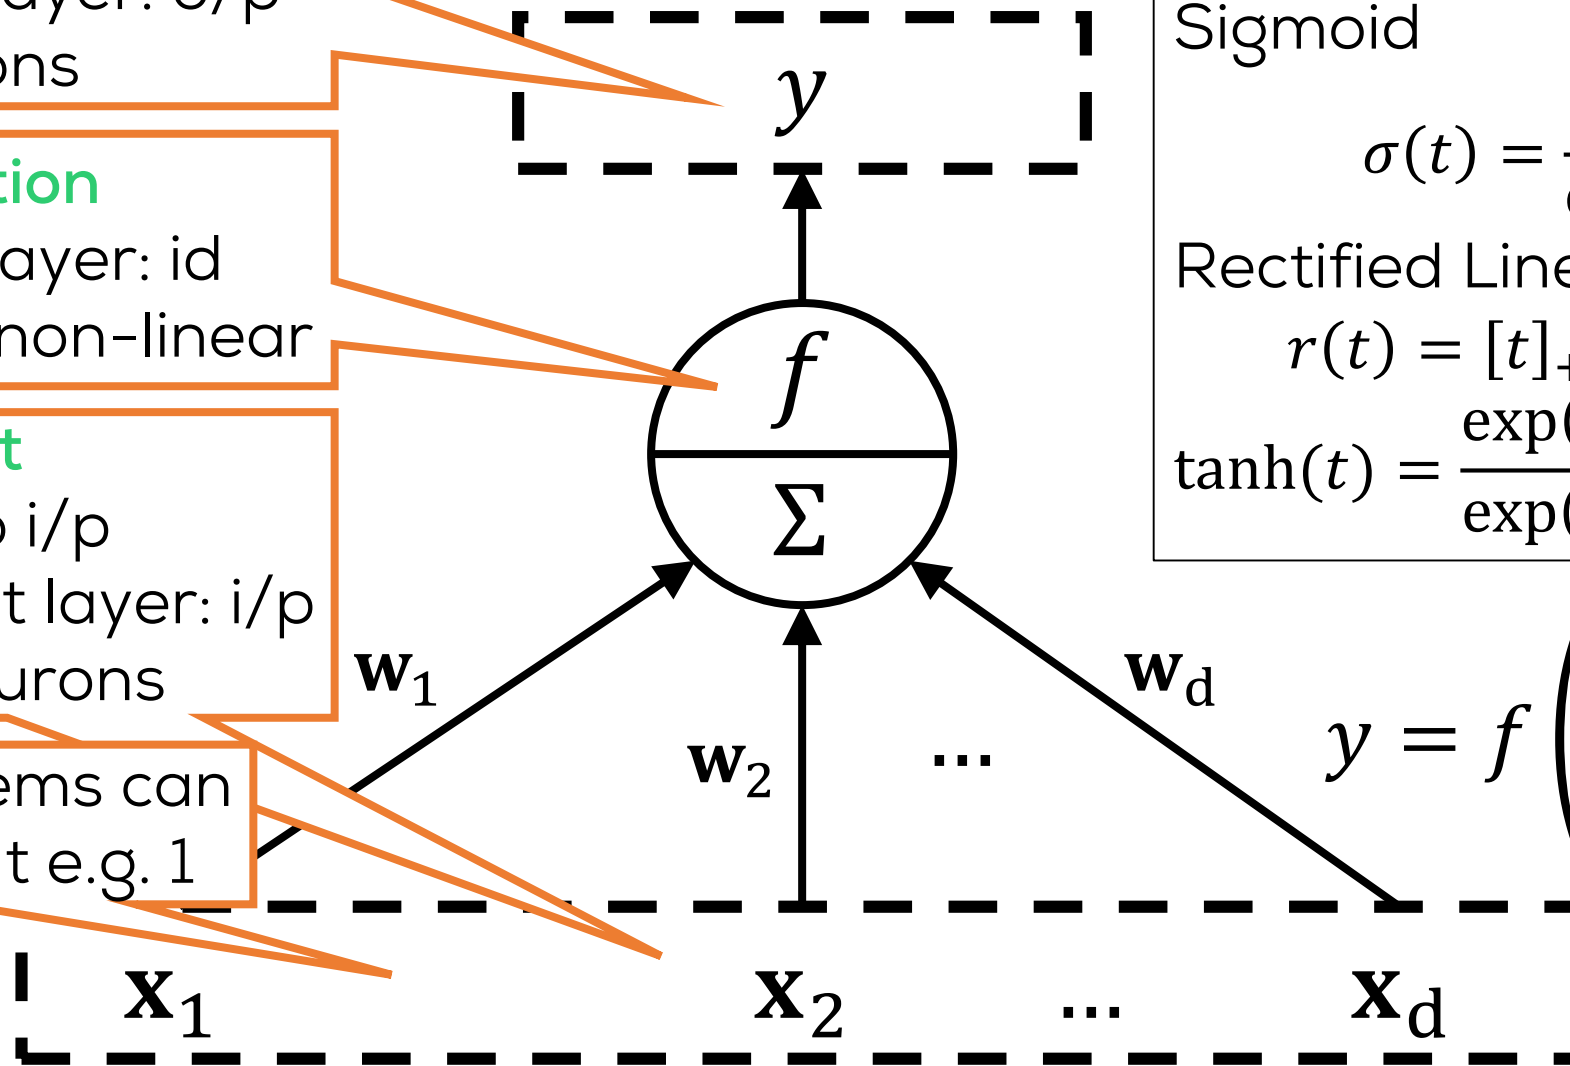
## Activation

Input/output layer: id  
Hidden layer: non-linear

## Input

Input layer: no i/p  
Hidden/output layer: i/p  
from other neurons

Some input items can  
be a constant e.g. 1



## Common "activation" fns $f$

Sigmoid

$$\sigma(t) = \frac{\exp(t)}{\exp(t) + 1}$$

Rectified Linear Unit (ReLU)

$$r(t) = [t]_+ = \max(t, 0)$$

$$\tanh(t) = \frac{\exp(2t) - 1}{\exp(2t) + 1}$$

# Neural Networks

## Output

Output layer: final o/p  
Input/hidden layer: o/p to other neurons

## Activation

Input/output layer: id  
Hidden layer: non-linear

## Input

Input layer: no i/p  
Hidden/output layer: i/p from other neurons

Some input items can be a constant e.g. 1


## Common "activation" fns $f$

Sigmoid

A single neuron also called a **perceptron**. Using non-linear  $f$  gives generalized linear models

$$f(t) = [t]_+ = \max(t, 0)$$

$$\tanh(t) = \frac{\exp(2t) - 1}{\exp(2t) + 1}$$



## Output

Output layer: final o/p  
Input/hidden layer: o/p to other neurons

## Activation

Input/output layer: id  
Hidden layer: non-linear

## Input

Input layer: no i/p  
Hidden/output layer: i/p from other neurons

Some input items can be a constant e.g. 1

Sometimes output layer is given a non-id activation. Matter of convention


## Common "activation" fns $f$

Sigmoid

A single neuron also called a **perceptron**. Using non-linear  $f$  gives generalized linear models


$$\text{ReLU}(t) = [t]_+ = \max(t, 0)$$

$$\tanh(t) = \frac{\exp(2t) - 1}{\exp(2t) + 1}$$





# Feedforward Neural Network



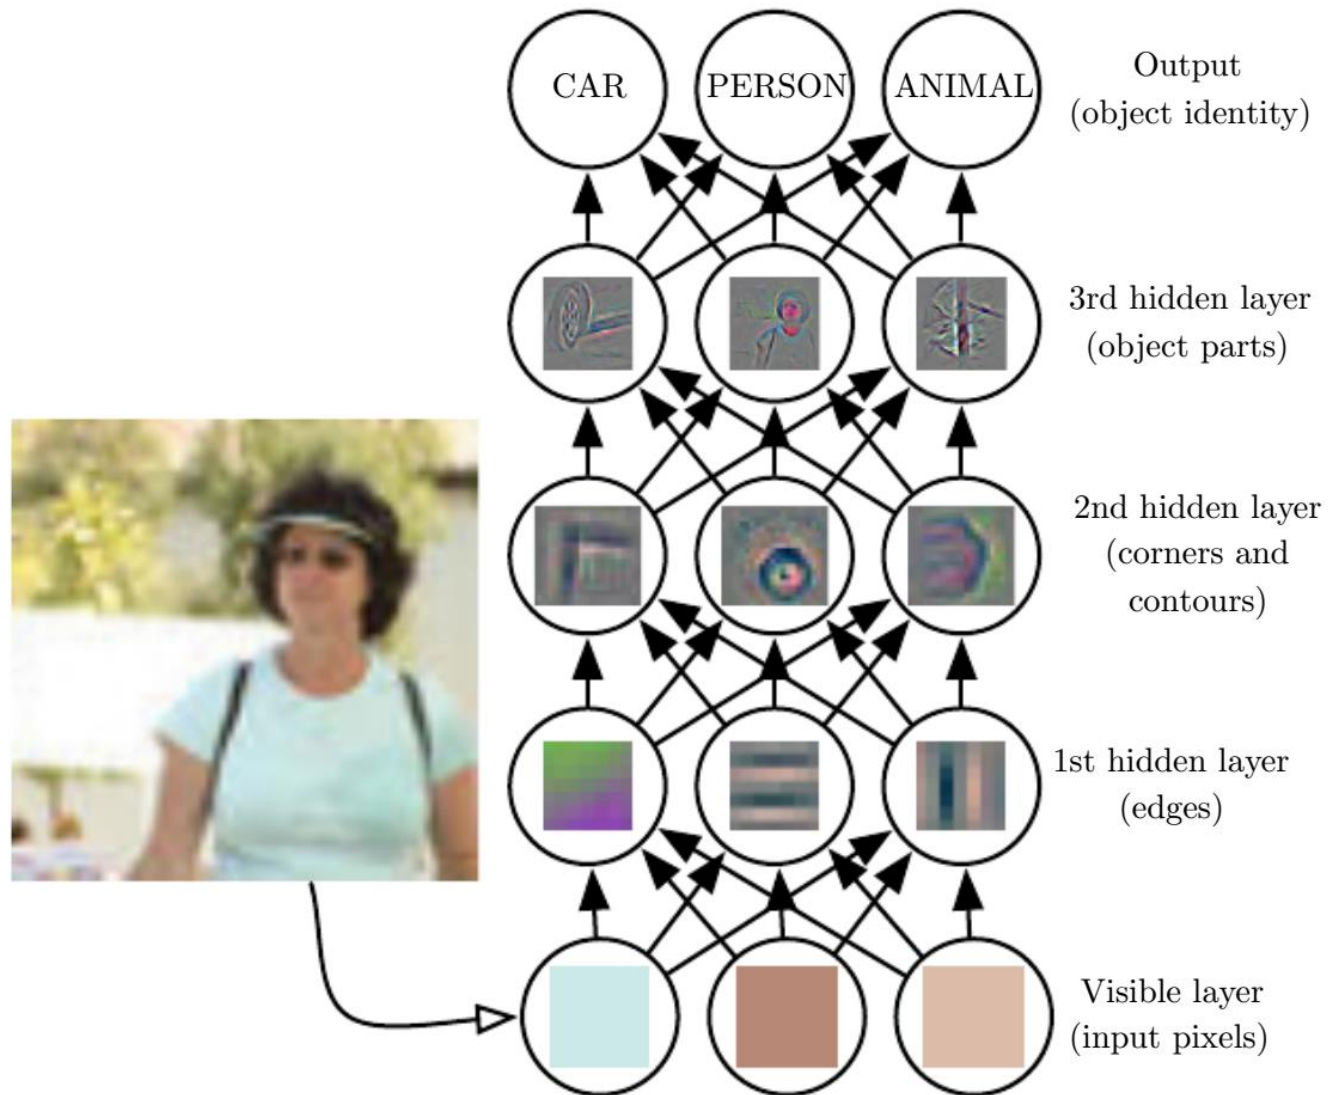
# Neural Networks as Feature Learners

Oct 25, 2017

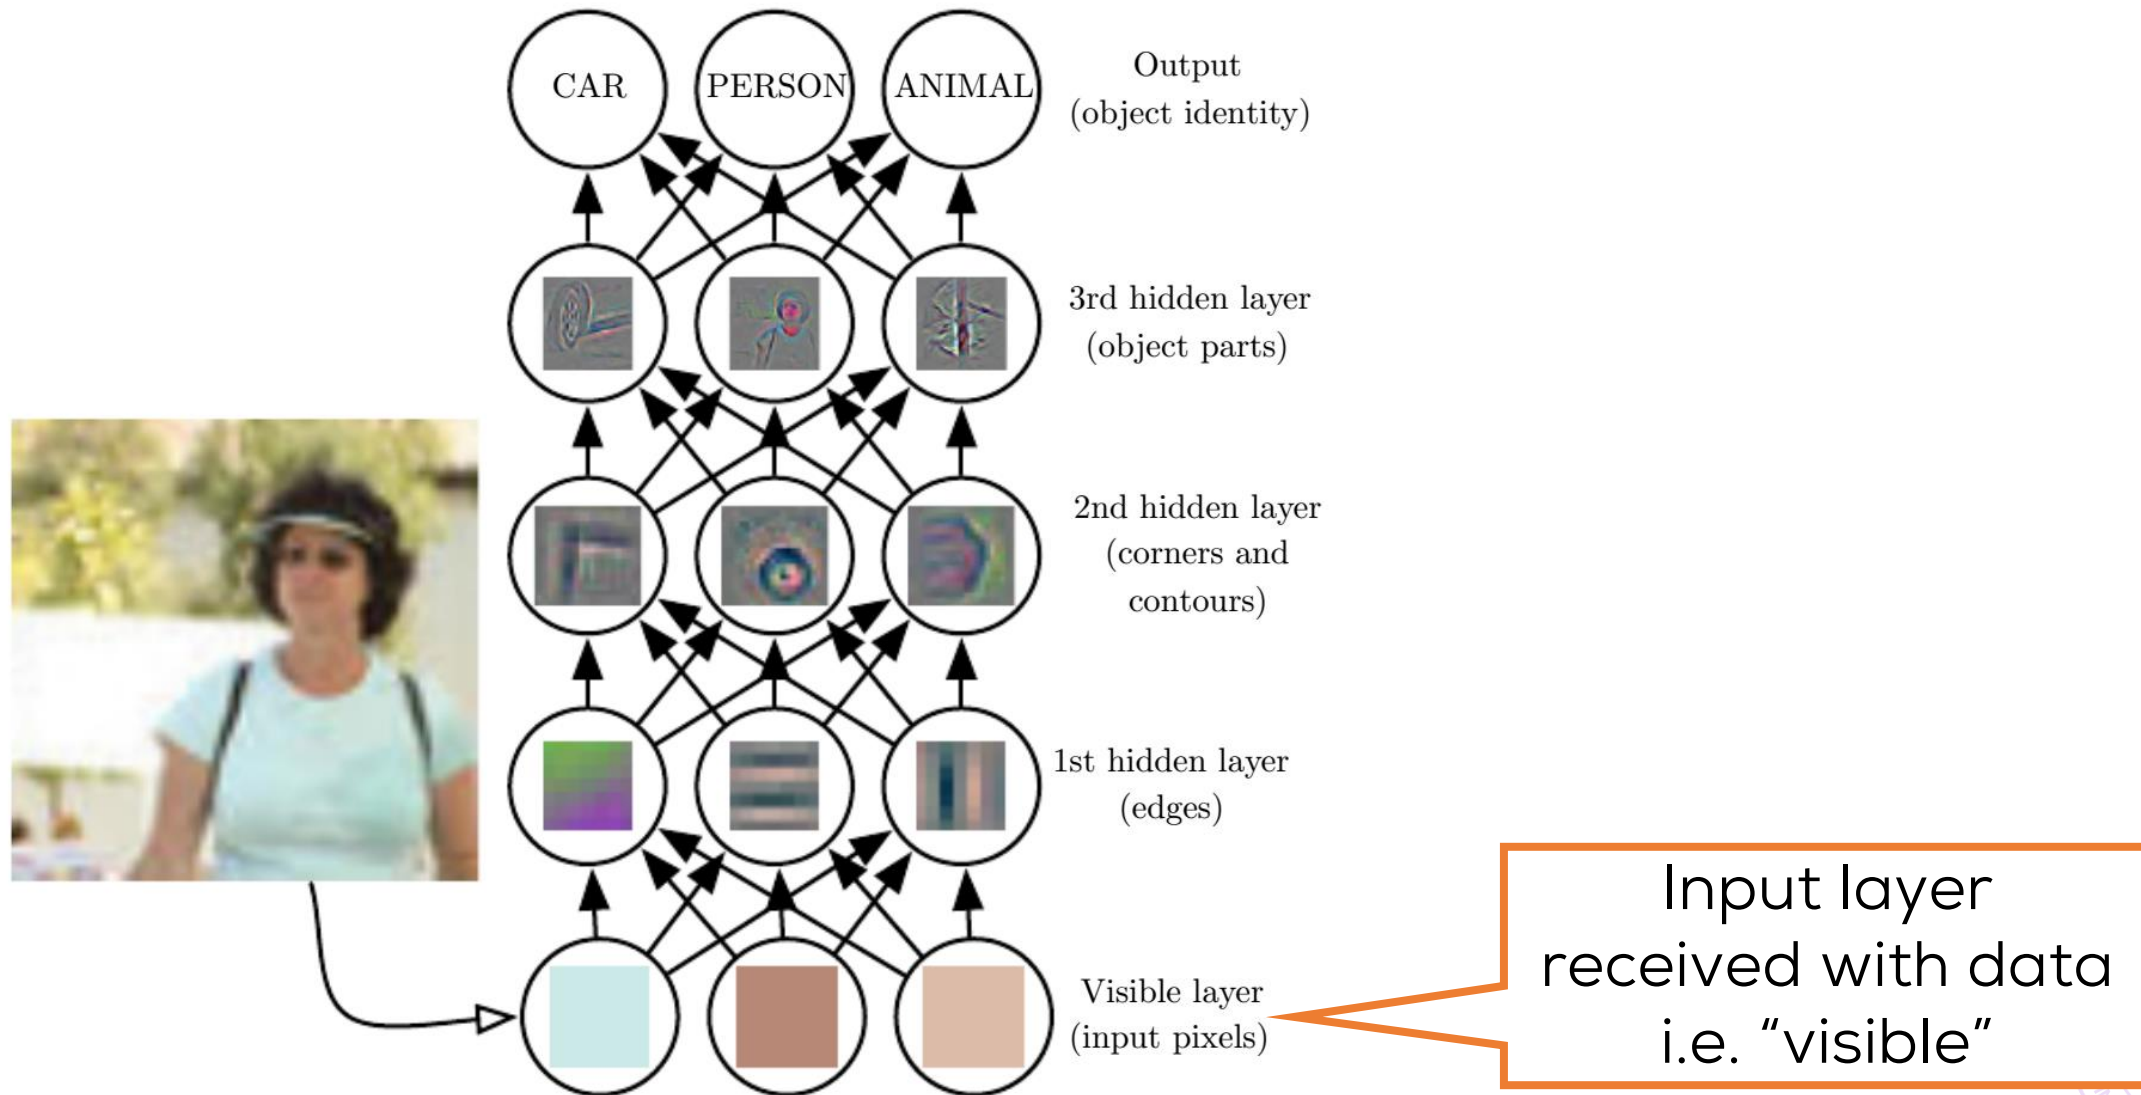
Zeiler, M. D. and Fergus, R. ECCV 2014



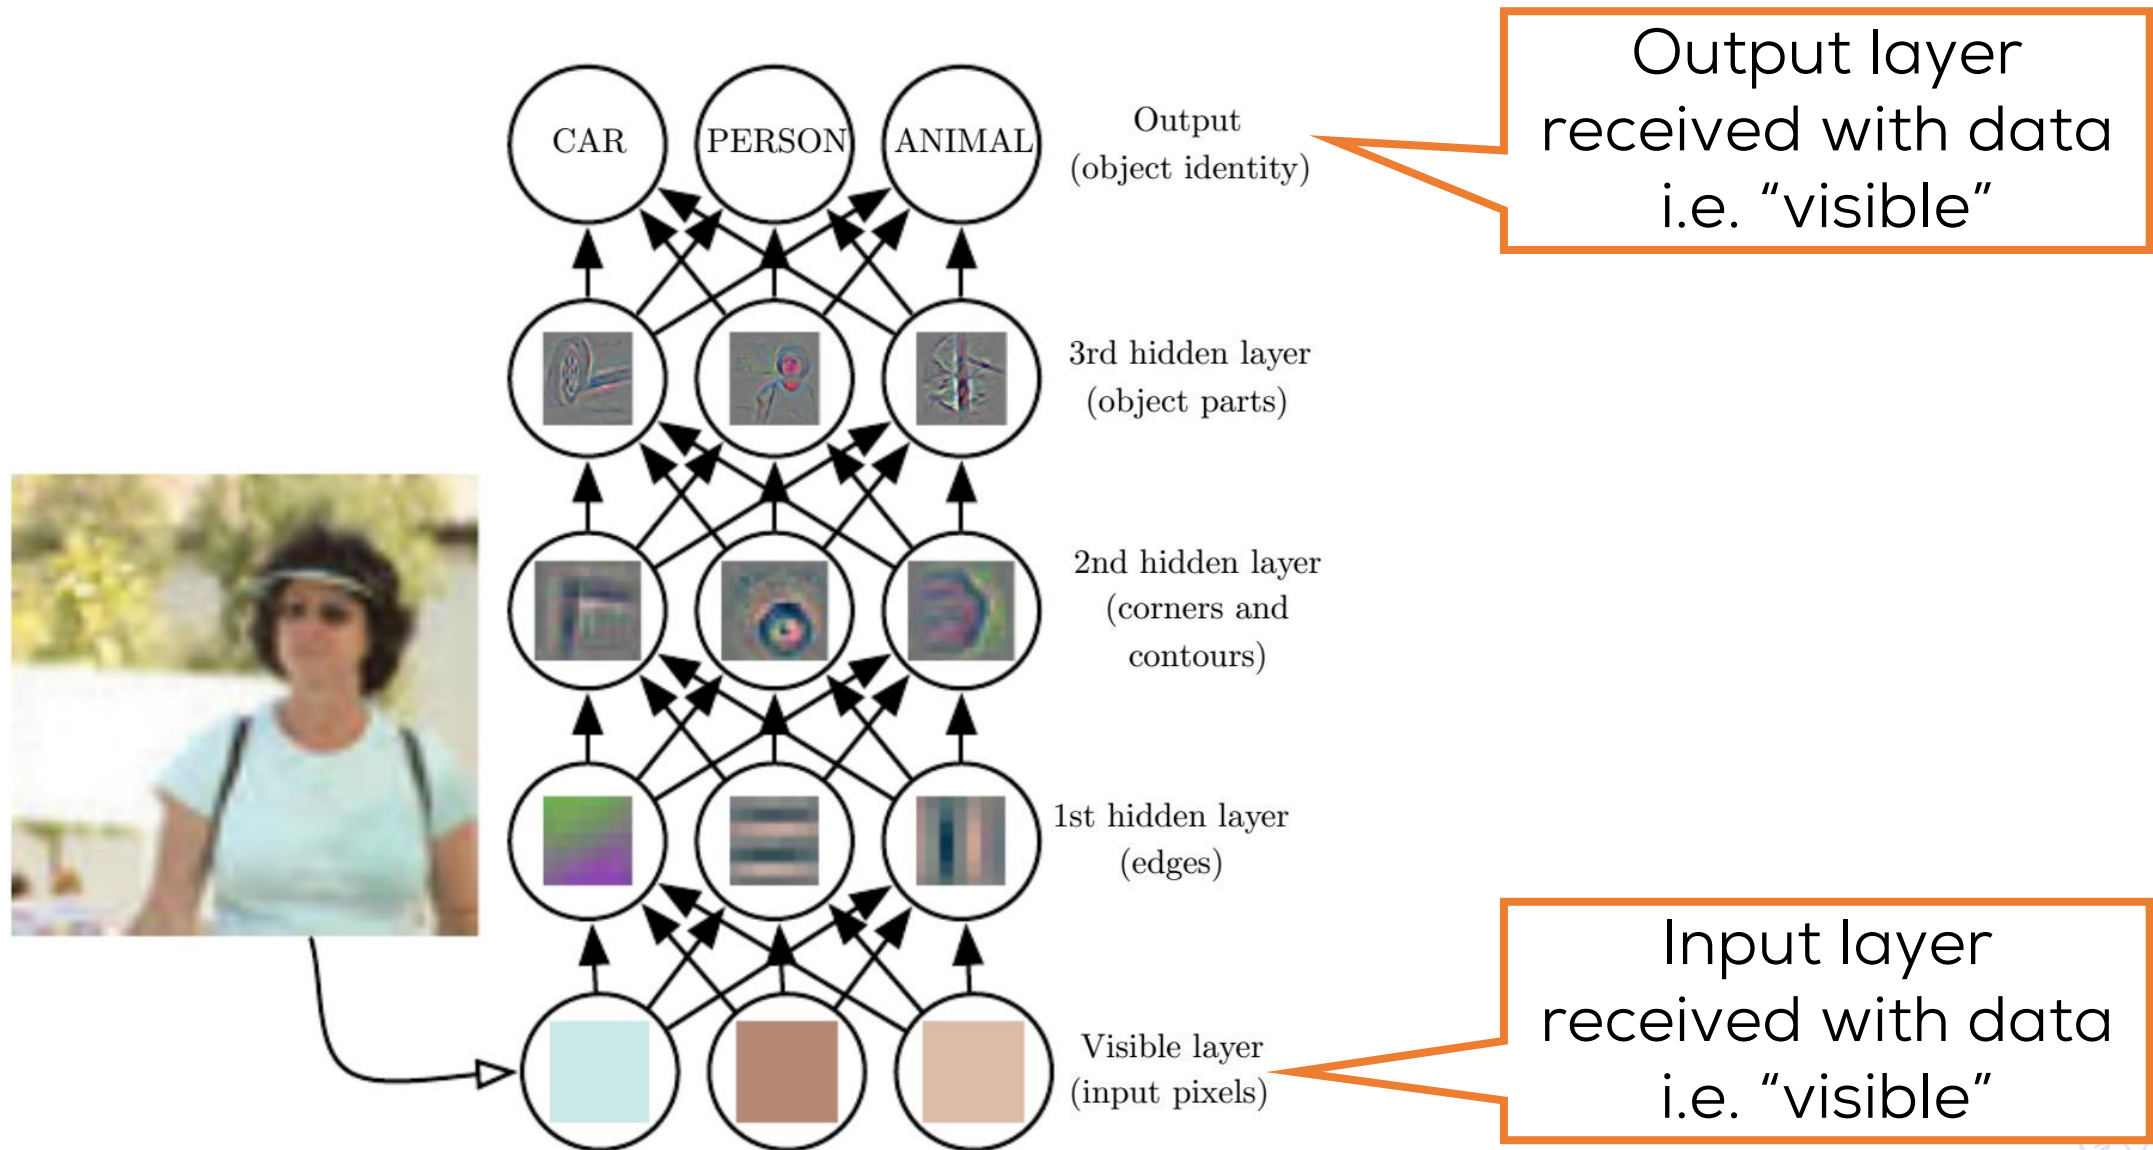
# Neural Networks as Feature Learners



# Neural Networks as Feature Learners

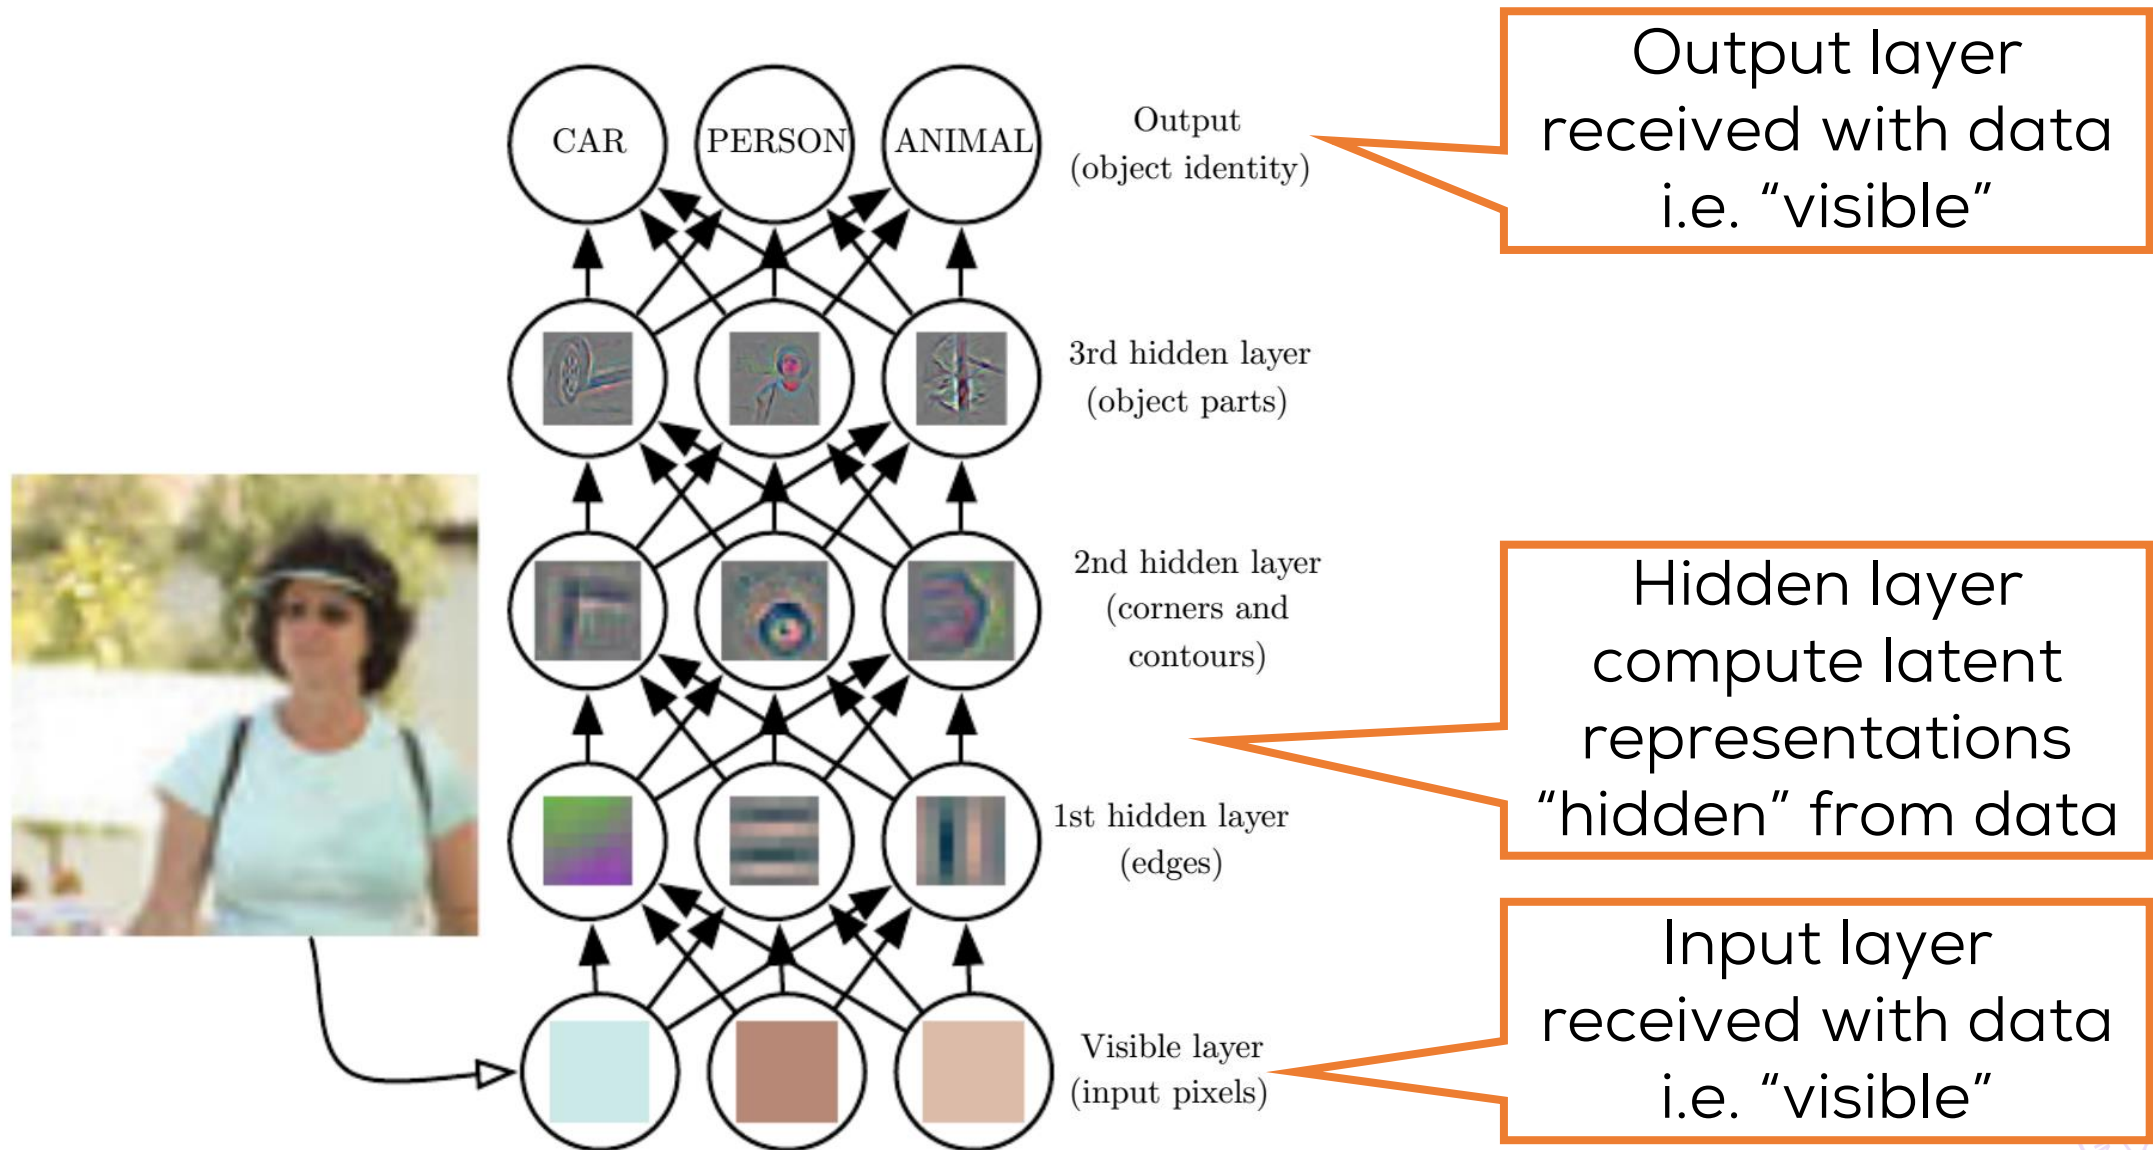


# Neural Networks as Feature Learners

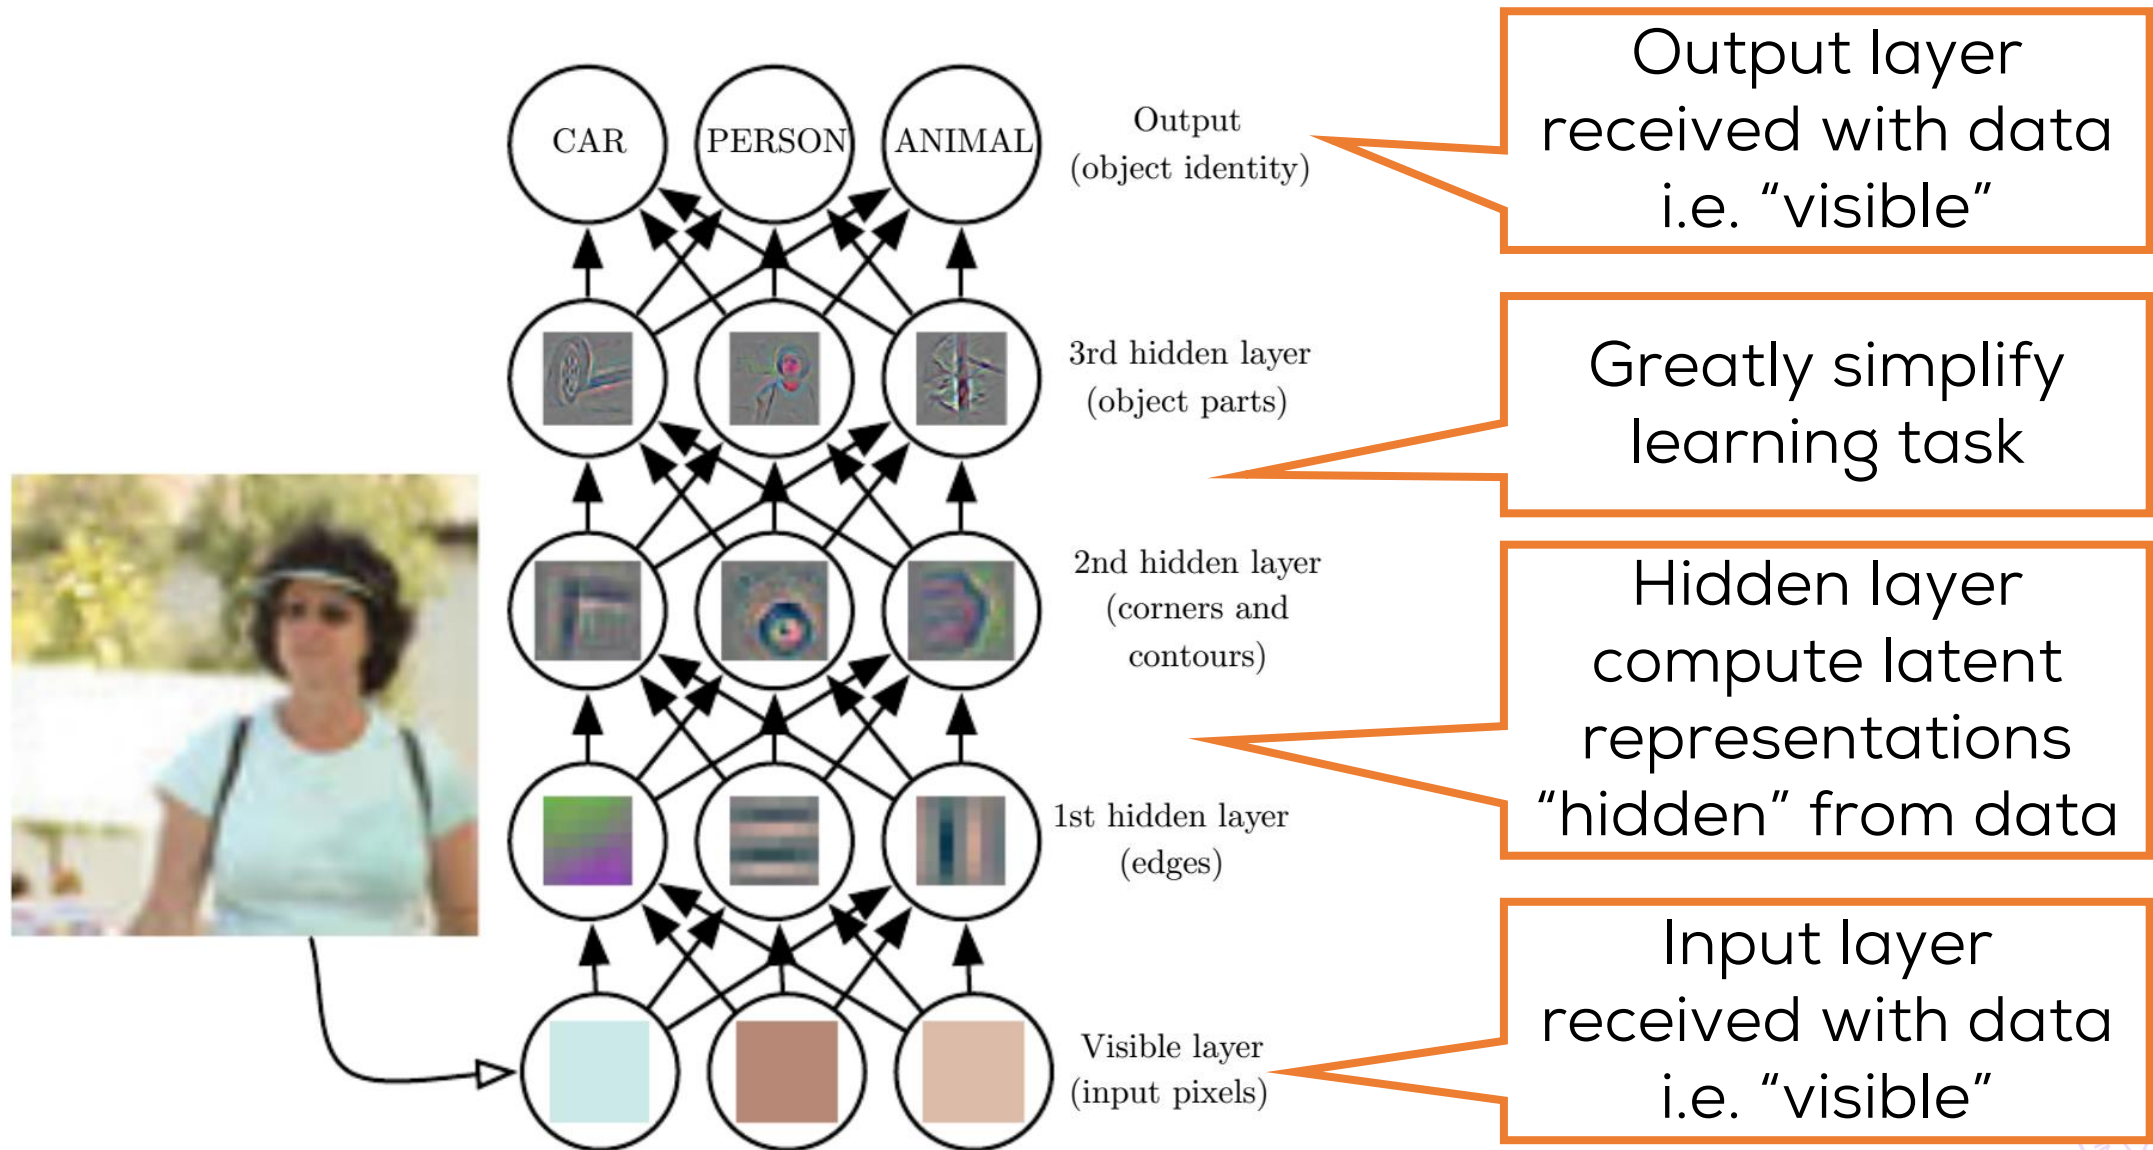




# Neural Networks as Feature Learners




# Neural Networks as Feature Learners



# Activation/link functions




# Activation/link functions



$$\sigma(t) = \frac{\exp(t)}{\exp(t) + 1}$$


Sigmoid

# Activation/link functions



$$\sigma(t) = \frac{\exp(t)}{\exp(t) + 1}$$


Sigmoid



$$\tanh(t) = \frac{\exp(2t) - 1}{\exp(2t) + 1} = 2\sigma(2t) - 1$$


Hyperbolic Tangent

# Activation/link functions




$$\sigma(t) = \frac{\exp(t)}{\exp(t) + 1}$$

Sigmoid



$$f(t) = \max(t, 0)$$


Rectified Linear Unit  
(ReLU)



$$\tanh(t) = \frac{\exp(2t) - 1}{\exp(2t) + 1} = 2\sigma(2t) - 1$$


Hyperbolic Tangent

# Activation/link functions




$$\sigma(t) = \frac{\exp(t)}{\exp(t) + 1}$$

Sigmoid




$$f(t) = \max(t, 0)$$

Rectified Linear Unit  
(ReLU)



$$\tanh(t) = \frac{\exp(2t) - 1}{\exp(2t) + 1} = 2\sigma(2t) - 1$$

Hyperbolic Tangent




$$f(t) = \log(1 + \exp(t))$$

Softplus

# Multi-output Networks




# Multi-output Networks




# Multi-output Networks

Multi-class/label  
classification,  
Vector regression



# Multi-output Networks


Multi-class/label  
classification,  
Vector regression





# Multi-output Networks

Multi-class/label  
classification,  
Vector regression




$$\mathbf{z}_i = \frac{\exp(\mathbf{t}_i)}{\sum_{j=1}^K \exp(\mathbf{t}_j)}$$

Softmax



# Multi-output Networks



Multi-class/label  
classification,  
Vector regression


$$\mathbf{z}_i = \frac{\exp(\mathbf{t}_i)}{\sum_{j=1}^K \exp(\mathbf{t}_j)}$$

Softmax

Like sigmoid, converts real  
values to probability values



# Multi-output Networks



Multi-class/label  
classification,  
Vector regression

$$\mathbf{z}_i = \frac{\exp(\mathbf{t}_i)}{\sum_{j=1}^K \exp(\mathbf{t}_j)}$$


Softmax

Like sigmoid, converts real  
values to probability values

Useful in modelling  
likelihood maximization  
problems using NN



# Multi-output Networks



Multi-class/label  
classification,  
Vector regression

$$z_i = \frac{\exp(\mathbf{t}_i)}{\sum_{j=1}^K \exp(\mathbf{t}_j)}$$


Softmax

Like sigmoid, converts real  
values to probability values

Useful in modelling  
likelihood maximization  
problems using NN

Normalize before use  
 $\tilde{\mathbf{t}}_i = \mathbf{t}_i - \max_j \mathbf{t}_j$

# Multi-output Networks



Multi-class/label  
classification,  
Vector regression

$$\mathbf{z}_i = \frac{\exp(\mathbf{t}_i)}{\sum_{j=1}^K \exp(\mathbf{t}_j)}$$

Softmax

Like sigmoid, converts real  
values to probability values

Useful in modelling  
likelihood maximization  
problems using NN

Normalize before use  
 $\tilde{\mathbf{t}}_i = \mathbf{t}_i - \max_j \mathbf{t}_j$

Note that  $f_{\text{SM}}(\mathbf{t}) = f_{\text{SM}}(\tilde{\mathbf{t}})$

# Loss/Cost Functions

- Squared loss

$$\ell_{\text{LS}}(\hat{y}, y) = (\hat{y} - y)^2$$

- Absolute difference

$$\ell_{\text{ABS}}(\hat{y}, y) = |\hat{y} - y|$$

- Negative log-likelihood loss

$$y \in [K], \hat{\mathbf{y}} = f_{\text{SM}}(\mathbf{t}), \mathbf{t} \in \mathbb{R}^K$$

$$\ell_{\text{NLL}}(\hat{\mathbf{y}}, y) = -\log(\hat{\mathbf{y}}_y)$$

- Cross-entropy

$$\ell_{\text{CE}}(\hat{y}, y) = y \cdot \log \hat{y} + (1 - y) \cdot \log(1 - \hat{y})$$

- Hinge loss

$$\ell_{\text{Hinge}}(\hat{y}, y) = [1 - y\hat{y}]_+$$

# Loss/Cost Functions

- Squared loss

$$\ell_{\text{LS}}(\hat{y}, y) = (\hat{y} - y)^2$$

LS used with  
identity/ReLU  
activation

- Absolute difference

$$\ell_{\text{ABS}}(\hat{y}, y) = |\hat{y} - y|$$

Sigmoid/softmax  
flatten out quickly  
so LS doesn't work

- Negative log-likelihood loss

$$y \in [K], \hat{\mathbf{y}} = f_{\text{SM}}(\mathbf{t}), \mathbf{t} \in \mathbb{R}^K$$

$$\ell_{\text{NLL}}(\hat{\mathbf{y}}, y) = -\log(\hat{\mathbf{y}}_y)$$

NLL, CE used with  
sigmoid/softmax  
activations

- Cross-entropy

$$\ell_{\text{CE}}(\hat{y}, y) = y \cdot \log \hat{y} + (1 - y) \cdot \log(1 - \hat{y})$$

- Hinge loss

$$\ell_{\text{Hinge}}(\hat{y}, y) = [1 - y\hat{y}]_+$$

CE/Hinge usually  
used when  $y \in \{0,1\}$   
is binary

# Training a Perceptron

Oct 25, 2017





# The Generalized Perceptron

- Simply a linear model with a wrapper thrown around it
- Makes predictions as

$$\hat{y} = f(\langle \mathbf{w}, \mathbf{x} \rangle)$$

- Given lots of data points


$$(\mathbf{x}^1, y^1), \dots, (\mathbf{x}^n, y^n), \mathbf{x}^i \in \mathbb{R}^d$$

- ... and a loss function

$$\ell: \mathbb{R} \times \mathbb{R} \rightarrow \mathbb{R}_+$$

- ... training a perceptron involves finding

$$\arg \min_{\mathbf{w} \in \mathbb{R}^d} \frac{1}{n} \sum_i^n \ell(f(\langle \mathbf{w}, \mathbf{x}^i \rangle), y^i) =: \arg \min_{\mathbf{w} \in \mathbb{R}^d} F(\mathbf{w})$$



# Gradient Descent Revisited

Oct 25, 2017



# Gradient Descent Revisited

## GRADIENT DESCENT

1. Initialize  $\mathbf{w}^0$
2. For  $t = 1, 2, \dots$ 
  1. Obtain a descent direction  $\mathbf{g}^t$
  2. Update  $\mathbf{w}^{t+1} \leftarrow \mathbf{w}^t - \eta_t \cdot \mathbf{g}^t$
3. Repeat until convergence

# Gradient Descent Revisited

## GRADIENT DESCENT

1. Initialize  $\mathbf{w}^0$
2. For  $t = 1, 2, \dots$ 
  1. Obtain a descent direction  $\mathbf{g}^t$
  2. Update  $\mathbf{w}^{t+1} \leftarrow \mathbf{w}^t - \eta_t \cdot \mathbf{g}^t$
3. Repeat until convergence

- How to find a descent direction?

# Gradient Descent Revisited

## GRADIENT DESCENT

1. Initialize  $\mathbf{w}^0$
2. For  $t = 1, 2, \dots$ 
  1. Obtain a descent direction  $\mathbf{g}^t$
  2. Update  $\mathbf{w}^{t+1} \leftarrow \mathbf{w}^t - \eta_t \cdot \mathbf{g}^t$
3. Repeat until convergence

- How to find a descent direction?
- How to choose a step length?

# Gradient Descent Revisited

## GRADIENT DESCENT

1. Initialize  $\mathbf{w}^0$
2. For  $t = 1, 2, \dots$ 
  1. Obtain a descent direction  $\mathbf{g}^t$
  2. Update  $\mathbf{w}^{t+1} \leftarrow \mathbf{w}^t - \eta_t \cdot \mathbf{g}^t$
3. Repeat until convergence

- How to find a descent direction?
- How to choose a step length?
- How to detect convergence?

# Gradient Descent Revisited

## GRADIENT DESCENT

1. Initialize  $\mathbf{w}^0$
2. For  $t = 1, 2, \dots$ 
  1. Obtain a descent direction  $\mathbf{g}^t$
  2. Update  $\mathbf{w}^{t+1} \leftarrow \mathbf{w}^t - \eta_t \cdot \mathbf{g}^t$
3. Repeat until convergence

- How to find a descent direction?
- How to choose a step length?
- How to detect convergence?
- How to avoid overfitting?

# Gradient Descent Revisited

## GRADIENT DESCENT

1. Initialize  $\mathbf{w}^0$
2. For  $t = 1, 2, \dots$ 
  1. Obtain a descent direction  $\mathbf{g}^t$
  2. Update  $\mathbf{w}^{t+1} \leftarrow \mathbf{w}^t - \eta_t \cdot \mathbf{g}^t$
3. Repeat until convergence

- How to find a descent direction?
- How to choose a step length?
- How to detect convergence?
- How to avoid overfitting?

Have to be more careful than earlier since now, problems are not nicely behaved



# Choosing a descent direction

- Batch gradient

$$\mathbf{g}^t = \nabla F(\mathbf{w}^t) = \frac{1}{n} \sum_{i=1}^n \ell'(f(\langle \mathbf{w}^t, \mathbf{x}^i \rangle), y^i) \cdot f'(\langle \mathbf{w}^t, \mathbf{x}^i \rangle) \cdot \mathbf{x}^i$$

- Mini-batch gradient: choose a mini-batch  $I_1^t, I_2^t, \dots, I_B^t \sim [n]$

$$\mathbf{g}^t = \frac{1}{B} \sum_{j=1}^B \ell' \left( f \left( \langle \mathbf{w}^t, \mathbf{x}^{I_j^t} \rangle \right), y^i \right) \cdot f' \left( \langle \mathbf{w}^t, \mathbf{x}^{I_j^t} \rangle \right) \cdot \mathbf{x}^{I_j^t}$$

- Newton's method

$$\mathbf{g}^t = \left( \nabla^2 F(\mathbf{w}^t) \right)^{-1} \nabla F(\mathbf{w}^t)$$

# Choosing a descent direction

- Batch gradient

$$\mathbf{g}^t = \nabla F(\mathbf{w}^t) = \frac{1}{n} \sum_{i=1}^n \ell'(f(\langle \mathbf{w}^t, \mathbf{x}^i \rangle), y^i) \cdot f'(\langle \mathbf{w}^t, \mathbf{x}^i \rangle) \cdot \mathbf{x}^i$$

Chain rule!

Very small batch sizes usually not used for deep networks

- Mini-batch gradient: choose a mini-batch  $I_1^t, I_2^t, \dots, I_B^t \sim [n]$

$$\mathbf{g}^t = \frac{1}{B} \sum_{j=1}^B \ell' \left( f \left( \langle \mathbf{w}^t, \mathbf{x}^{I_j^t} \rangle \right), y^i \right) \cdot f' \left( \langle \mathbf{w}^t, \mathbf{x}^{I_j^t} \rangle \right) \cdot \mathbf{x}^{I_j^t}$$

- Newton's method

$$\mathbf{g}^t = \left( \nabla^2 F(\mathbf{w}^t) \right)^{-1} \nabla F(\mathbf{w}^t)$$

For a NN with  $E$  edges,  $\mathcal{O}(E^3)$  time per iteration!

Expensive!  $\mathcal{O}(d^3)$  time per iteration

# How to detect convergence

- Tolerance technique
  - For a predecided tolerance value  $\epsilon$ , if  $F(\mathbf{w}^t) < \epsilon$ , stop
- Zero-th order technique
  - If function value has not changed too much between iterations, stop!
$$|F(\mathbf{w}^{t+1}) - F(\mathbf{w}^t)| < \tau$$
- First order technique
  - If gradient is too “small”  $\|\nabla F(\mathbf{w}^t)\|_2 < \delta$ , stop!
- Primal dual
  - If primal and dual objective values are close, stop
  - Does not work every where – reliable for convex problems

# How to decide step length?

- Simply rule of thumb – naïve but fast  
Choose  $\eta_t \rightarrow 0$  (diminishing) and  $\sum \eta_t \rightarrow \infty$  (infinite travel)
- Example  $\eta_t = C/\sqrt{t}$  or  $\eta_t = C/t$  for some  $C > 0$
- Line search – super careful but expensive
$$\eta_t = \arg \min_{\eta \geq 0} F(\mathbf{w}^t - \eta \cdot \mathbf{g}^t)$$
- Can we do something more adaptive?
- Can we let each coordinate of  $\mathbf{w}$  get its own step length?
- $\mathbf{w}^{t+1} \leftarrow \mathbf{w}^t - \eta \cdot H^{-1} \mathbf{g}^t$  where  $H = \text{diag}(h_1, h_2, \dots, h_d)$
- Wait ... this looks like the Newton method!
- Indeed this is an approximate Newton method – wait a bit!

# Momentum Methods

- Introduce a velocity term to push GD along, avoid oscillations


$$\begin{aligned}\mathbf{v}^t &= \gamma \cdot \mathbf{v}^{t-1} + \eta \cdot \nabla F(\mathbf{w}^t) \\ \mathbf{w}^{t+1} &\leftarrow \mathbf{w}^t - \mathbf{v}^t\end{aligned}$$

- **Nesterov's accelerated gradient (NAG)**

- Does a “look-ahead”

$$\begin{aligned}\mathbf{v}^t &= \gamma \cdot \mathbf{v}^{t-1} + \eta \cdot \nabla F(\mathbf{w}^t - \eta \cdot \mathbf{v}^{t-1}) \\ \mathbf{w}^{t+1} &\leftarrow \mathbf{w}^t - \mathbf{v}^t\end{aligned}$$

- For “smooth” convex problems, NAG ensures  $\epsilon$ -convergence in just  $\mathcal{O}(1/\epsilon)$  steps hence the name “accelerated” gradient
- Don't have very deep insights why it works ☺



# Adaptive Learning Rates

$$\mathbf{w}^{t+1} \leftarrow \mathbf{w}^t - \eta \cdot (H^t)^{-1} \mathbf{g}^t$$
$$H^t = \text{diag}(h_1^t, h_2^t, \dots, h_d^t)$$

- Adagrad (Duchi et al. 2011)

$$h_i^t = \sqrt{\epsilon + \sum_{\tau=1}^t (\mathbf{g}_i^\tau)^2}$$

- Note that if all coordinates get roughly similar gradients then we have  $h_i^t \approx t$  and Adagrad behaves as if we had set  $\eta_t \approx \eta/t$
- However, if some coordinate getting updated very vigorously,  $|\mathbf{g}_i^\tau| \gg 0$  for all  $\tau$ , Adagrad slows it down a bit
- If some coordinate is static  $\mathbf{g}_i^\tau \equiv 0$  for all  $\tau$ , then  $h_i^t = \epsilon$ , no effect

# Adaptive Learning Rates

$$\mathbf{w}^{t+1} \leftarrow \mathbf{w}^t - \eta \cdot (H^t)^{-1} \mathbf{g}^t$$
$$H^t = \text{diag}(h_1^t, h_2^t, \dots, h_d^t)$$

- RMSProp (Hinton 2012)

$$h_i^t = \sqrt{\epsilon + v_i^t}$$

$$v_i^t = \gamma \cdot v_i^{t-1} + (1 - \gamma) \cdot (\mathbf{g}_i^t)^2$$

- Adagrad can be too aggressive in forcing step sizes down
- RMSProp has better performance in non-convex settings
- Hinton suggests  $\gamma \approx 0.9$
- May be combined with NAG as well!

# Adaptive Learning Rates

$$\mathbf{w}^{t+1} \leftarrow \mathbf{w}^t - \eta \cdot (H^t)^{-1} \mathbf{u}^t$$
$$H^t = \text{diag}(h_1^t, h_2^t, \dots, h_d^t)$$

- Adam (Kingma and Ba 2014)

$$h_i^t = \sqrt{\epsilon + v_i^t}$$
$$\mathbf{u}^t = \gamma_1 \cdot \mathbf{u}^{t-1} + (1 - \gamma_1) \cdot \mathbf{g}^t$$
$$v_i^t = \gamma_2 \cdot v_i^{t-1} + (1 - \gamma_2) \cdot (\mathbf{g}_i^t)^2$$

- Keeps track of past gradients as well as squared gradients
- Actually does a bias correction step before using  $\mathbf{u}^t$  and  $H^t$
- Details can be found in Deep Learning textbook



# How to prevent overfitting?

- Add a regularization term  $L_2/L_1$  to the objective

$$\arg \min_{\mathbf{w} \in \mathbb{R}^d} F(\mathbf{w}) + \lambda \cdot \|\mathbf{w}\|_2^2$$

- Gradients calculations change very slightly
- Constraint the weights of the network to satisfy  $|\mathbf{w}_i| < r$

$$\arg \min_{\|\mathbf{w}\|_\infty < r} F(\mathbf{w})$$

- Sometimes gradient coordinates are also *clipped* this way
- Noise injection in output
  - For binary classification  $y^i = 0 \rightarrow y^i = \epsilon, y^i = 1 \rightarrow y^i = 1 - \epsilon$
  - For regression problems,  $y^i \rightarrow y^i + \epsilon^i$ , where  $\epsilon^i \sim \mathcal{N}(0, \sigma^2)$
  - Can be shown to be equivalent to regularization in nice cases

# How to prevent overfitting?

- Early stopping – return model with best validation set performance rather than best training set performance
- Can use many of these strategies in combination
- Parameter sharing – add constraints of the form

$$\mathbf{w}_i = \mathbf{w}_j$$

- Sparse recovery – constrain, say at least 10% weights to be zero

$$\|\mathbf{w}\|_0 \leq k \ll d$$

- Dropout
  - Effectively trains on multiple sparse networks in parallel
  - While executing a GD update, randomly remove edges or entire nodes from network so they do not participate
  - Can be shown to be equivalent to  $L_2$  reg. in nice settings

# Other techniques used to train NNs

- Pre-training
- Batch normalization
- Curriculum learning
- Pre-training
- Conjugate gradient descent
- Normalized gradient descent
- Approximate Newton method L-BFGS
- Some of these developed earlier for convex opt. problems
- Some we have discussed earlier in context of linear models  
Coordinate descent, Model averaging (Ruppert-Polyak method)

Nice discussion in the  
Deep Learning book

# Up Next

- Training multi-layered perceptrons – backpropagation
- Autoencoders
- RNNs
- CNNs
- GANs
- Cannot go into too many details but will cover basics 😊

# Please give your Feedback

<http://tinyurl.com/ml17-18afb>