

Latent Variable Models for Dimensionality Reduction

Piyush Rai

Probabilistic Machine Learning (CS772A)

September 7, 2017

A Simple Additive Model for Data Compression

- Consider a set of observations $\mathbf{x}_1, \dots, \mathbf{x}_N$, with $\mathbf{x}_n \in \mathbb{R}^D$

A Simple Additive Model for Data Compression

- Consider a set of observations $\mathbf{x}_1, \dots, \mathbf{x}_N$, with $\mathbf{x}_n \in \mathbb{R}^D$
- Let's approximate each \mathbf{x}_n by a **linear combination** of K vectors $\mathbf{w}_1, \mathbf{w}_2, \dots, \mathbf{w}_K$ ($K \ll D$)

$$\mathbf{x}_n \approx \sum_{k=1}^K z_{nk} \mathbf{w}_k$$

A Simple Additive Model for Data Compression

- Consider a set of observations $\mathbf{x}_1, \dots, \mathbf{x}_N$, with $\mathbf{x}_n \in \mathbb{R}^D$
- Let's approximate each \mathbf{x}_n by a **linear combination** of K vectors $\mathbf{w}_1, \mathbf{w}_2, \dots, \mathbf{w}_K$ ($K \ll D$)

$$\mathbf{x}_n \approx \sum_{k=1}^K z_{nk} \mathbf{w}_k \quad \text{or} \quad \mathbf{x}_n \approx \mathbf{W} \mathbf{z}_n$$

A Simple Additive Model for Data Compression

- Consider a set of observations $\mathbf{x}_1, \dots, \mathbf{x}_N$, with $\mathbf{x}_n \in \mathbb{R}^D$
- Let's approximate each \mathbf{x}_n by a **linear combination** of K vectors $\mathbf{w}_1, \mathbf{w}_2, \dots, \mathbf{w}_K$ ($K \ll D$)

$$\mathbf{x}_n \approx \sum_{k=1}^K z_{nk} \mathbf{w}_k \quad \text{or} \quad \mathbf{x}_n \approx \mathbf{W} \mathbf{z}_n$$

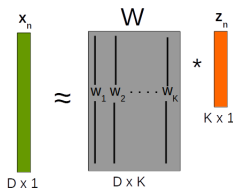
where $\mathbf{W} = [\mathbf{w}_1 \dots \mathbf{w}_K]$ is $D \times K$, each $\mathbf{w}_k \in \mathbb{R}^D$, and $\mathbf{z}_n = [z_{n1} \dots z_{nK}] \in \mathbb{R}^K$

A Simple Additive Model for Data Compression

- Consider a set of observations $\mathbf{x}_1, \dots, \mathbf{x}_N$, with $\mathbf{x}_n \in \mathbb{R}^D$
- Let's approximate each \mathbf{x}_n by a **linear combination** of K vectors $\mathbf{w}_1, \mathbf{w}_2, \dots, \mathbf{w}_K$ ($K \ll D$)

$$\mathbf{x}_n \approx \sum_{k=1}^K z_{nk} \mathbf{w}_k \quad \text{or} \quad \mathbf{x}_n \approx \mathbf{W} \mathbf{z}_n$$

where $\mathbf{W} = [\mathbf{w}_1 \dots \mathbf{w}_K]$ is $D \times K$, each $\mathbf{w}_k \in \mathbb{R}^D$, and $\mathbf{z}_n = [z_{n1} \dots z_{nK}] \in \mathbb{R}^K$

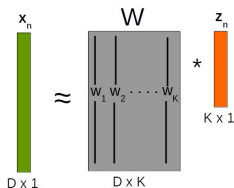


A Simple Additive Model for Data Compression

- Consider a set of observations $\mathbf{x}_1, \dots, \mathbf{x}_N$, with $\mathbf{x}_n \in \mathbb{R}^D$
- Let's approximate each \mathbf{x}_n by a **linear combination** of K vectors $\mathbf{w}_1, \mathbf{w}_2, \dots, \mathbf{w}_K$ ($K \ll D$)

$$\mathbf{x}_n \approx \sum_{k=1}^K z_{nk} \mathbf{w}_k \quad \text{or} \quad \mathbf{x}_n \approx \mathbf{W} \mathbf{z}_n$$

where $\mathbf{W} = [\mathbf{w}_1 \dots \mathbf{w}_K]$ is $D \times K$, each $\mathbf{w}_k \in \mathbb{R}^D$, and $\mathbf{z}_n = [z_{n1} \dots z_{nK}] \in \mathbb{R}^K$



A Simple Additive Model for Data Compression

- Consider a set of observations $\mathbf{x}_1, \dots, \mathbf{x}_N$, with $\mathbf{x}_n \in \mathbb{R}^D$
- Let's approximate each \mathbf{x}_n by a **linear combination** of K vectors $\mathbf{w}_1, \mathbf{w}_2, \dots, \mathbf{w}_K$ ($K \ll D$)

$$\mathbf{x}_n \approx \sum_{k=1}^K z_{nk} \mathbf{w}_k \quad \text{or} \quad \mathbf{x}_n \approx \mathbf{W} \mathbf{z}_n$$

where $\mathbf{W} = [\mathbf{w}_1 \dots \mathbf{w}_K]$ is $D \times K$, each $\mathbf{w}_k \in \mathbb{R}^D$, and $\mathbf{z}_n = [z_{n1} \dots z_{nK}] \in \mathbb{R}^K$



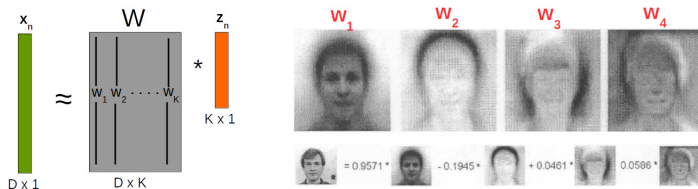
- z_{nk} tell us much of “component” \mathbf{w}_k is present in the observation \mathbf{x}_n

A Simple Additive Model for Data Compression

- Consider a set of observations $\mathbf{x}_1, \dots, \mathbf{x}_N$, with $\mathbf{x}_n \in \mathbb{R}^D$
- Let's approximate each \mathbf{x}_n by a **linear combination** of K vectors $\mathbf{w}_1, \mathbf{w}_2, \dots, \mathbf{w}_K$ ($K \ll D$)

$$\mathbf{x}_n \approx \sum_{k=1}^K z_{nk} \mathbf{w}_k \quad \text{or} \quad \mathbf{x}_n \approx \mathbf{W} \mathbf{z}_n$$

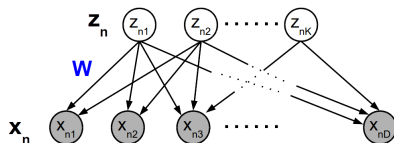
where $\mathbf{W} = [\mathbf{w}_1 \dots \mathbf{w}_K]$ is $D \times K$, each $\mathbf{w}_k \in \mathbb{R}^D$, and $\mathbf{z}_n = [z_{n1} \dots z_{nK}] \in \mathbb{R}^K$



- z_{nk} tell us much of “component” \mathbf{w}_k is present in the observation \mathbf{x}_n
- Can think of $\mathbf{z}_n \in \mathbb{R}^K$ as a “**compressed**” latent representation of $\mathbf{x}_n \in \mathbb{R}^D$ (would like to learn it)

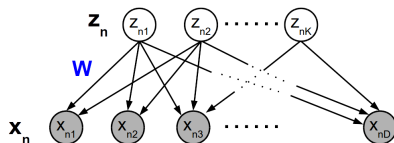
Zooming-in into the Model..

- Can think of latent z_n generating x_n via a **linear mapping** defined by $D \times K$ matrix **W**



Zooming-in into the Model..

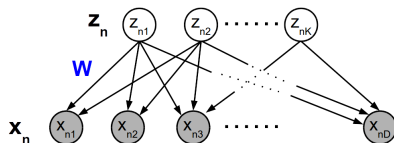
- Can think of latent z_n generating x_n via a **linear mapping** defined by $D \times K$ matrix **W**



- Note: Also possible to define a **nonlinear mapping** from z_n to x_n , e.g. **Variational Auto-encoders (VAE)**, **Gaussian Process Latent Variable Models (GPLVM)**

Zooming-in into the Model..

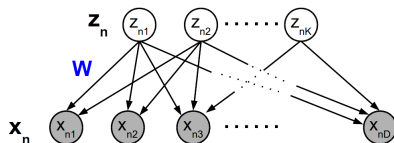
- Can think of latent z_n generating x_n via a **linear mapping** defined by $D \times K$ matrix **W**



- Note: Also possible to define a **nonlinear mapping** from z_n to x_n , e.g. **Variational Auto-encoders (VAE)**, **Gaussian Process Latent Variable Models (GPLVM)**
 - Such models are much more flexible but inference becomes harder though

Zooming-in into the Model..

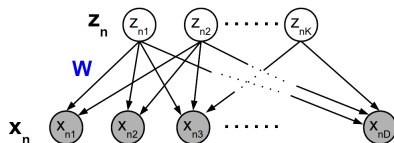
- Can think of latent z_n generating x_n via a **linear mapping** defined by $D \times K$ matrix **W**



- Note: Also possible to define a **nonlinear mapping** from z_n to x_n , e.g. **Variational Auto-encoders (VAE)**, **Gaussian Process Latent Variable Models (GPLVM)**
 - Such models are much more flexible but inference becomes harder though
 - Will look at some of these when discussing **deep generative models**

Zooming-in into the Model..

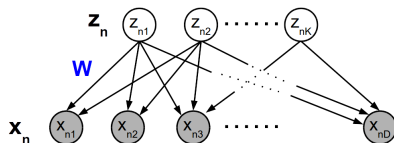
- Can think of latent z_n generating x_n via a **linear mapping** defined by $D \times K$ matrix **W**



- Note: Also possible to define a **nonlinear mapping** from z_n to x_n , e.g. **Variational Auto-encoders (VAE)**, **Gaussian Process Latent Variable Models (GPLVM)**
 - Such models are much more flexible but inference becomes harder though
 - Will look at some of these when discussing **deep generative models**
- **Today's focus** will be on the **linear case** (will also look at a **nonlinear extension via mixtures**)

Zooming-in into the Model..

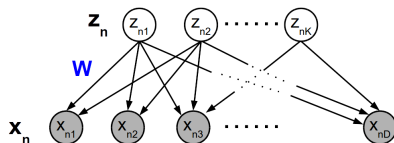
- Can think of latent \mathbf{z}_n generating \mathbf{x}_n via a **linear mapping** defined by $D \times K$ matrix \mathbf{W}



- Note: Also possible to define a **nonlinear mapping** from \mathbf{z}_n to \mathbf{x}_n , e.g. **Variational Auto-encoders (VAE)**, **Gaussian Process Latent Variable Models (GPLVM)**
 - Such models are much more flexible but inference becomes harder though
 - Will look at some of these when discussing **deep generative models**
- **Today's focus** will be on the **linear case** (will also look at a **nonlinear extension via mixtures**)
 - Linear models are simple but very powerful. Very easy to do inference in such models (e.g., using EM)

Zooming-in into the Model..

- Can think of latent \mathbf{z}_n generating \mathbf{x}_n via a **linear mapping** defined by $D \times K$ matrix \mathbf{W}



- Note: Also possible to define a **nonlinear mapping** from \mathbf{z}_n to \mathbf{x}_n , e.g. **Variational Auto-encoders (VAE)**, **Gaussian Process Latent Variable Models (GPLVM)**
 - Such models are much more flexible but inference becomes harder though
 - Will look at some of these when discussing **deep generative models**
- **Today's focus** will be on the **linear case** (will also look at a **nonlinear extension via mixtures**)
 - Linear models are simple but very powerful. Very easy to do inference in such models (e.g., using EM)
 - **Nice interpretability**. E.g., columns of $[\mathbf{w}_1 \dots \mathbf{w}_K]$ are like K “latent parts” that compose \mathbf{x}_n

Latent Factor Models for Dimensionality Reduction

- In the linear model, we represented \mathbf{x}_n approximately as $\mathbf{x}_n \approx \mathbf{W}\mathbf{z}_n$

Latent Factor Models for Dimensionality Reduction

- In the linear model, we represented \mathbf{x}_n approximately as $\mathbf{x}_n \approx \mathbf{W}\mathbf{z}_n$
- Let's assume the **data-point specific Gaussian noise** or “residual” $\epsilon_n \sim \mathcal{N}(0, \sigma^2 \mathbf{I}_D)$ and write

$$\mathbf{x}_n = \mathbf{W}\mathbf{z}_n + \epsilon_n$$

Latent Factor Models for Dimensionality Reduction

- In the linear model, we represented \mathbf{x}_n approximately as $\mathbf{x}_n \approx \mathbf{W}\mathbf{z}_n$
- Let's assume the **data-point specific Gaussian noise** or “residual” $\epsilon_n \sim \mathcal{N}(0, \sigma^2 \mathbf{I}_D)$ and write

$$\mathbf{x}_n = \mathbf{W}\mathbf{z}_n + \epsilon_n$$

- The distribution of \mathbf{x}_n , conditioned on \mathbf{z}_n will be

Latent Factor Models for Dimensionality Reduction

- In the linear model, we represented \mathbf{x}_n approximately as $\mathbf{x}_n \approx \mathbf{W}\mathbf{z}_n$
- Let's assume the **data-point specific Gaussian noise** or “residual” $\epsilon_n \sim \mathcal{N}(0, \sigma^2 \mathbf{I}_D)$ and write

$$\mathbf{x}_n = \mathbf{W}\mathbf{z}_n + \epsilon_n$$

- The distribution of \mathbf{x}_n , conditioned on \mathbf{z}_n will be

$$p(\mathbf{x}_n | \mathbf{z}_n, \mathbf{W}, \sigma^2) = \mathcal{N}(\mathbf{W}\mathbf{z}_n, \sigma^2 \mathbf{I}_D)$$

Latent Factor Models for Dimensionality Reduction

- In the linear model, we represented \mathbf{x}_n approximately as $\mathbf{x}_n \approx \mathbf{W}\mathbf{z}_n$
- Let's assume the **data-point specific Gaussian noise** or “residual” $\epsilon_n \sim \mathcal{N}(0, \sigma^2 \mathbf{I}_D)$ and write

$$\mathbf{x}_n = \mathbf{W}\mathbf{z}_n + \epsilon_n$$

- The distribution of \mathbf{x}_n , conditioned on \mathbf{z}_n will be

$$p(\mathbf{x}_n | \mathbf{z}_n, \mathbf{W}, \sigma^2) = \mathcal{N}(\mathbf{W}\mathbf{z}_n, \sigma^2 \mathbf{I}_D)$$

- If \mathbf{z}_n is a *priori* Gaussian distributed, e.g., $p(\mathbf{z}_n) = \mathcal{N}(\mathbf{0}, \mathbf{I}_K)$

Latent Factor Models for Dimensionality Reduction

- In the linear model, we represented \mathbf{x}_n approximately as $\mathbf{x}_n \approx \mathbf{W}\mathbf{z}_n$
- Let's assume the **data-point specific Gaussian noise** or “residual” $\epsilon_n \sim \mathcal{N}(0, \sigma^2 \mathbf{I}_D)$ and write

$$\mathbf{x}_n = \mathbf{W}\mathbf{z}_n + \epsilon_n$$

- The distribution of \mathbf{x}_n , conditioned on \mathbf{z}_n will be

$$p(\mathbf{x}_n | \mathbf{z}_n, \mathbf{W}, \sigma^2) = \mathcal{N}(\mathbf{W}\mathbf{z}_n, \sigma^2 \mathbf{I}_D)$$

- If \mathbf{z}_n is *a priori* Gaussian distributed, e.g., $p(\mathbf{z}_n) = \mathcal{N}(\mathbf{0}, \mathbf{I}_K) \Rightarrow$ **Linear Gaussian Model (LGM)**

Latent Factor Models for Dimensionality Reduction

- In the linear model, we represented \mathbf{x}_n approximately as $\mathbf{x}_n \approx \mathbf{W}\mathbf{z}_n$
- Let's assume the **data-point specific Gaussian noise** or “residual” $\epsilon_n \sim \mathcal{N}(0, \sigma^2 \mathbf{I}_D)$ and write

$$\mathbf{x}_n = \mathbf{W}\mathbf{z}_n + \epsilon_n$$

- The distribution of \mathbf{x}_n , conditioned on \mathbf{z}_n will be

$$p(\mathbf{x}_n | \mathbf{z}_n, \mathbf{W}, \sigma^2) = \mathcal{N}(\mathbf{W}\mathbf{z}_n, \sigma^2 \mathbf{I}_D)$$

- If \mathbf{z}_n is *a priori* Gaussian distributed, e.g., $p(\mathbf{z}_n) = \mathcal{N}(\mathbf{0}, \mathbf{I}_K) \Rightarrow$ **Linear Gaussian Model** (LGM)
- Note: Recall LGMs have nice properties, e.g., easy to compute $p(\mathbf{z}|\mathbf{x})$ or $p(\mathbf{x})$ (both Gaussians!)

Latent Factor Models for Dimensionality Reduction

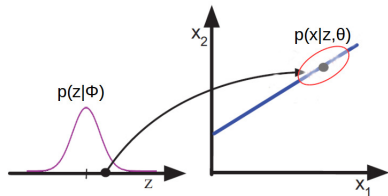
- In the linear model, we represented \mathbf{x}_n approximately as $\mathbf{x}_n \approx \mathbf{W}\mathbf{z}_n$
- Let's assume the **data-point specific Gaussian noise** or “residual” $\epsilon_n \sim \mathcal{N}(0, \sigma^2 \mathbf{I}_D)$ and write

$$\mathbf{x}_n = \mathbf{W}\mathbf{z}_n + \epsilon_n$$

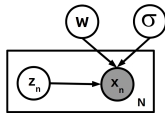
- The distribution of \mathbf{x}_n , conditioned on \mathbf{z}_n will be

$$p(\mathbf{x}_n | \mathbf{z}_n, \mathbf{W}, \sigma^2) = \mathcal{N}(\mathbf{W}\mathbf{z}_n, \sigma^2 \mathbf{I}_D)$$

- If \mathbf{z}_n is a *priori* Gaussian distributed, e.g., $p(\mathbf{z}_n) = \mathcal{N}(\mathbf{0}, \mathbf{I}_K) \Rightarrow$ **Linear Gaussian Model (LGM)**
- Note: Recall LGMs have nice properties, e.g., easy to compute $p(\mathbf{z}|\mathbf{x})$ or $p(\mathbf{x})$ (both Gaussians!)

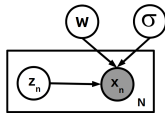


Latent Factor Models for Dimensionality Reduction



- Can think of the data generation process via the following [generative story](#) for each $\mathbf{x}_n \in \mathbb{R}^D$

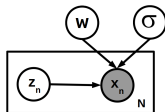
Latent Factor Models for Dimensionality Reduction



- Can think of the data generation process via the following **generative story** for each $\mathbf{x}_n \in \mathbb{R}^D$
 - Generate latent variables $\mathbf{z}_n \in \mathbb{R}^K$ as

$$\mathbf{z}_n \sim \mathcal{N}(0, \mathbf{I}_K)$$

Latent Factor Models for Dimensionality Reduction

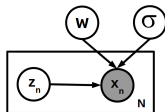


- Can think of the data generation process via the following **generative story** for each $\mathbf{x}_n \in \mathbb{R}^D$
 - Generate latent variables $\mathbf{z}_n \in \mathbb{R}^K$ as
 - Generate data \mathbf{x}_n conditioned on \mathbf{z}_n as

$$\mathbf{z}_n \sim \mathcal{N}(\mathbf{0}, \mathbf{I}_K)$$

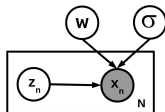
$$\mathbf{x}_n \sim \mathcal{N}(\mathbf{W}\mathbf{z}_n, \sigma^2 \mathbf{I}_D)$$

Latent Factor Models for Dimensionality Reduction



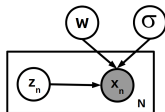
- Can think of the data generation process via the following **generative story** for each $\mathbf{x}_n \in \mathbb{R}^D$
 - Generate latent variables $\mathbf{z}_n \in \mathbb{R}^K$ as
$$\mathbf{z}_n \sim \mathcal{N}(\mathbf{0}, \mathbf{I}_K)$$
 - Generate data \mathbf{x}_n conditioned on \mathbf{z}_n as
$$\mathbf{x}_n \sim \mathcal{N}(\mathbf{W}\mathbf{z}_n, \sigma^2 \mathbf{I}_D)$$
- $\mathbf{W} = [\mathbf{w}_1, \dots, \mathbf{w}_K]$ is the $D \times K$ “**factor loading matrix**” (also known as “**dictionary**”)

Latent Factor Models for Dimensionality Reduction



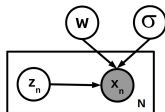
- Can think of the data generation process via the following **generative story** for each $\mathbf{x}_n \in \mathbb{R}^D$
 - Generate latent variables $\mathbf{z}_n \in \mathbb{R}^K$ as
$$\mathbf{z}_n \sim \mathcal{N}(\mathbf{0}, \mathbf{I}_K)$$
 - Generate data \mathbf{x}_n conditioned on \mathbf{z}_n as
$$\mathbf{x}_n \sim \mathcal{N}(\mathbf{W}\mathbf{z}_n, \sigma^2 \mathbf{I}_D)$$
- $\mathbf{W} = [\mathbf{w}_1, \dots, \mathbf{w}_K]$ is the $D \times K$ “**factor loading matrix**” (also known as “**dictionary**”)
- \mathbf{z}_n called **latent features** or **latent factors** or **factor scores** of \mathbf{x}_n w.r.t. dictionary \mathbf{W}

Latent Factor Models for Dimensionality Reduction



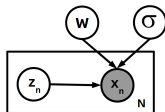
- Can think of the data generation process via the following **generative story** for each $\mathbf{x}_n \in \mathbb{R}^D$
 - Generate latent variables $\mathbf{z}_n \in \mathbb{R}^K$ as
$$\mathbf{z}_n \sim \mathcal{N}(\mathbf{0}, \mathbf{I}_K)$$
 - Generate data \mathbf{x}_n conditioned on \mathbf{z}_n as
$$\mathbf{x}_n \sim \mathcal{N}(\mathbf{W}\mathbf{z}_n, \sigma^2 \mathbf{I}_D)$$
- $\mathbf{W} = [\mathbf{w}_1, \dots, \mathbf{w}_K]$ is the $D \times K$ “**factor loading matrix**” (also known as “**dictionary**”)
- \mathbf{z}_n called **latent features** or **latent factors** or **factor scores** of \mathbf{x}_n w.r.t. dictionary \mathbf{W}
- This model is popularly known as **Probabilistic PCA (PPCA)**

Latent Factor Models for Dimensionality Reduction



- Can think of the data generation process via the following **generative story** for each $\mathbf{x}_n \in \mathbb{R}^D$
 - Generate latent variables $\mathbf{z}_n \in \mathbb{R}^K$ as
$$\mathbf{z}_n \sim \mathcal{N}(\mathbf{0}, \mathbf{I}_K)$$
 - Generate data \mathbf{x}_n conditioned on \mathbf{z}_n as
$$\mathbf{x}_n \sim \mathcal{N}(\mathbf{W}\mathbf{z}_n, \sigma^2 \mathbf{I}_D)$$
- $\mathbf{W} = [\mathbf{w}_1, \dots, \mathbf{w}_K]$ is the $D \times K$ “**factor loading matrix**” (also known as “**dictionary**”)
- \mathbf{z}_n called **latent features** or **latent factors** or **factor scores** of \mathbf{x}_n w.r.t. dictionary \mathbf{W}
- This model is popularly known as **Probabilistic PCA (PPCA)**
- **Note:** The Gaussians on \mathbf{z} and \mathbf{x} can be replaced by other distributions (e.g., Exp. Family)

Latent Factor Models for Dimensionality Reduction



- Can think of the data generation process via the following **generative story** for each $\mathbf{x}_n \in \mathbb{R}^D$
 - Generate latent variables $\mathbf{z}_n \in \mathbb{R}^K$ as

$$\mathbf{z}_n \sim \mathcal{N}(\mathbf{0}, \mathbf{I}_K)$$

- Generate data \mathbf{x}_n conditioned on \mathbf{z}_n as

$$\mathbf{x}_n \sim \mathcal{N}(\mathbf{W}\mathbf{z}_n, \sigma^2 \mathbf{I}_D)$$

- $\mathbf{W} = [\mathbf{w}_1, \dots, \mathbf{w}_K]$ is the $D \times K$ “**factor loading matrix**” (also known as “**dictionary**”)
- \mathbf{z}_n called **latent features** or **latent factors** or **factor scores** of \mathbf{x}_n w.r.t. dictionary \mathbf{W}
- This model is popularly known as **Probabilistic PCA (PPCA)**
- **Note:** The Gaussians on \mathbf{z} and \mathbf{x} can be replaced by other distributions (e.g., Exp. Family)
- In the Gaussian case, if $p(\epsilon) = \mathcal{N}(\mathbf{0}, \Psi)$ where Ψ is diagonal, it's called **Factor Analysis (FA)**

PPCA: Marginal and Posterior Distributions

- Suppose we're modeling D -dim data using a (say zero mean) Gaussian

$$p(\mathbf{x}) = \mathcal{N}(\mathbf{0}, \mathbf{\Sigma})$$

where $\mathbf{\Sigma}$ is a $D \times D$ p.s.d. cov. matrix, $\mathcal{O}(D^2)$ parameters needed

PPCA: Marginal and Posterior Distributions

- Suppose we're modeling D -dim data using a (say zero mean) Gaussian

$$p(\mathbf{x}) = \mathcal{N}(\mathbf{0}, \mathbf{\Sigma})$$

where $\mathbf{\Sigma}$ is a $D \times D$ p.s.d. cov. matrix, $\mathcal{O}(D^2)$ parameters needed

- Consider modeling the same data using PPCA: $p(\mathbf{x}|\mathbf{z}) = \mathcal{N}(\mathbf{W}\mathbf{z}, \sigma^2\mathbf{I}_D)$, $p(\mathbf{z}) = \mathcal{N}(\mathbf{0}, \mathbf{I}_K)$

PPCA: Marginal and Posterior Distributions

- Suppose we're modeling D -dim data using a (say zero mean) Gaussian

$$p(\mathbf{x}) = \mathcal{N}(0, \mathbf{\Sigma})$$

where $\mathbf{\Sigma}$ is a $D \times D$ p.s.d. cov. matrix, $\mathcal{O}(D^2)$ parameters needed

- Consider modeling the same data using PPCA: $p(\mathbf{x}|\mathbf{z}) = \mathcal{N}(\mathbf{W}\mathbf{z}, \sigma^2\mathbf{I}_D)$, $p(\mathbf{z}) = \mathcal{N}(0, \mathbf{I}_K)$
- For this Gaussian PPCA, the **marginal** distribution $p(\mathbf{x}) = \int p(\mathbf{x}, \mathbf{z})d\mathbf{z}$ is

$$p(\mathbf{x}) = \mathcal{N}(0, \mathbf{W}\mathbf{W}^\top + \sigma^2\mathbf{I}_D)$$

(recall the Gaussian marginal result; verify)

PPCA: Marginal and Posterior Distributions

- Suppose we're modeling D -dim data using a (say zero mean) Gaussian

$$p(\mathbf{x}) = \mathcal{N}(\mathbf{0}, \mathbf{\Sigma})$$

where $\mathbf{\Sigma}$ is a $D \times D$ p.s.d. cov. matrix, $\mathcal{O}(D^2)$ parameters needed

- Consider modeling the same data using PPCA: $p(\mathbf{x}|\mathbf{z}) = \mathcal{N}(\mathbf{W}\mathbf{z}, \sigma^2\mathbf{I}_D)$, $p(\mathbf{z}) = \mathcal{N}(\mathbf{0}, \mathbf{I}_K)$
- For this Gaussian PPCA, the **marginal** distribution $p(\mathbf{x}) = \int p(\mathbf{x}, \mathbf{z})d\mathbf{z}$ is

$$\boxed{p(\mathbf{x}) = \mathcal{N}(\mathbf{0}, \mathbf{W}\mathbf{W}^\top + \sigma^2\mathbf{I}_D)} \quad (\text{recall the Gaussian marginal result; verify})$$

- Cov. matrix is close to **low-rank** as $\sigma^2 \rightarrow 0$. Only $(DK + 1)$ parameters needed (**nice when $D \gg N$**)

PPCA: Marginal and Posterior Distributions

- Suppose we're modeling D -dim data using a (say zero mean) Gaussian

$$p(\mathbf{x}) = \mathcal{N}(0, \mathbf{\Sigma})$$

where $\mathbf{\Sigma}$ is a $D \times D$ p.s.d. cov. matrix, $\mathcal{O}(D^2)$ parameters needed

- Consider modeling the same data using PPCA: $p(\mathbf{x}|\mathbf{z}) = \mathcal{N}(\mathbf{W}\mathbf{z}, \sigma^2\mathbf{I}_D)$, $p(\mathbf{z}) = \mathcal{N}(0, \mathbf{I}_K)$
- For this Gaussian PPCA, the **marginal** distribution $p(\mathbf{x}) = \int p(\mathbf{x}, \mathbf{z})d\mathbf{z}$ is

$$\boxed{p(\mathbf{x}) = \mathcal{N}(0, \mathbf{W}\mathbf{W}^\top + \sigma^2\mathbf{I}_D)} \quad (\text{recall the Gaussian marginal result; verify})$$

- Cov. matrix is close to **low-rank** as $\sigma^2 \rightarrow 0$. Only $(DK + 1)$ parameters needed (**nice when $D \gg N$**)
 - **PPCA = Low-rank Gaussian**. Fewer parameters to learn; less chance of overfitting

PPCA: Marginal and Posterior Distributions

- Suppose we're modeling D -dim data using a (say zero mean) Gaussian

$$p(\mathbf{x}) = \mathcal{N}(\mathbf{0}, \mathbf{\Sigma})$$

where $\mathbf{\Sigma}$ is a $D \times D$ p.s.d. cov. matrix, $\mathcal{O}(D^2)$ parameters needed

- Consider modeling the same data using PPCA: $p(\mathbf{x}|\mathbf{z}) = \mathcal{N}(\mathbf{W}\mathbf{z}, \sigma^2\mathbf{I}_D)$, $p(\mathbf{z}) = \mathcal{N}(\mathbf{0}, \mathbf{I}_K)$
- For this Gaussian PPCA, the **marginal** distribution $p(\mathbf{x}) = \int p(\mathbf{x}, \mathbf{z})d\mathbf{z}$ is

$p(\mathbf{x}) = \mathcal{N}(\mathbf{0}, \mathbf{W}\mathbf{W}^\top + \sigma^2\mathbf{I}_D)$

 (recall the Gaussian marginal result; verify)

- Cov. matrix is close to **low-rank** as $\sigma^2 \rightarrow 0$. Only $(DK + 1)$ parameters needed (**nice when $D \gg N$**)
 - **PPCA = Low-rank Gaussian**. Fewer parameters to learn; less chance of overfitting
- **Posterior** of \mathbf{z} is also Gaussian (recall the Gaussian posterior result of linear Gaussian model)

$p(\mathbf{z}|\mathbf{x}, \mathbf{W}, \sigma^2) = \mathcal{N}(\mathbf{M}^{-1}\mathbf{W}^\top \mathbf{x}, \sigma^2\mathbf{M}^{-1})$

 (where $\mathbf{M} = \mathbf{W}^\top \mathbf{W} + \sigma^2\mathbf{I}_K$)

Identifiability

- Note that for PPCA, we have $p(\mathbf{x}_n) = \mathcal{N}(\mathbf{0}, \mathbf{W}\mathbf{W}^\top + \sigma^2\mathbf{I}_D)$

Identifiability

- Note that for PPCA, we have $p(\mathbf{x}_n) = \mathcal{N}(\mathbf{0}, \mathbf{W}\mathbf{W}^\top + \sigma^2\mathbf{I}_D)$
- If we replace \mathbf{W} by $\tilde{\mathbf{W}} = \mathbf{W}\mathbf{R}$ for some $K \times K$ orthogonal rotation matrix \mathbf{R} then

$$p(\mathbf{x}_n) = \mathcal{N}(\mathbf{0}, \tilde{\mathbf{W}}\tilde{\mathbf{W}}^\top + \sigma^2\mathbf{I}_D)$$

Identifiability

- Note that for PPCA, we have $p(\mathbf{x}_n) = \mathcal{N}(\mathbf{0}, \mathbf{W}\mathbf{W}^\top + \sigma^2\mathbf{I}_D)$
- If we replace \mathbf{W} by $\tilde{\mathbf{W}} = \mathbf{W}\mathbf{R}$ for some $K \times K$ orthogonal rotation matrix \mathbf{R} then

$$\begin{aligned} p(\mathbf{x}_n) &= \mathcal{N}(\mathbf{0}, \tilde{\mathbf{W}}\tilde{\mathbf{W}}^\top + \sigma^2\mathbf{I}_D) \\ &= \mathcal{N}(\mathbf{0}, \mathbf{W}\mathbf{R}\mathbf{R}^\top\mathbf{W}^\top + \sigma^2\mathbf{I}_D) \end{aligned}$$

Identifiability

- Note that for PPCA, we have $p(\mathbf{x}_n) = \mathcal{N}(\mathbf{0}, \mathbf{W}\mathbf{W}^\top + \sigma^2\mathbf{I}_D)$
- If we replace \mathbf{W} by $\tilde{\mathbf{W}} = \mathbf{W}\mathbf{R}$ for some $K \times K$ orthogonal rotation matrix \mathbf{R} then

$$\begin{aligned} p(\mathbf{x}_n) &= \mathcal{N}(\mathbf{0}, \tilde{\mathbf{W}}\tilde{\mathbf{W}}^\top + \sigma^2\mathbf{I}_D) \\ &= \mathcal{N}(\mathbf{0}, \mathbf{W}\mathbf{R}\mathbf{R}^\top\mathbf{W}^\top + \sigma^2\mathbf{I}_D) \\ &= \mathcal{N}(\mathbf{0}, \mathbf{W}\mathbf{W}^\top + \sigma^2\mathbf{I}_D) \end{aligned}$$

Identifiability

- Note that for PPCA, we have $p(\mathbf{x}_n) = \mathcal{N}(\mathbf{0}, \mathbf{W}\mathbf{W}^\top + \sigma^2\mathbf{I}_D)$
- If we replace \mathbf{W} by $\tilde{\mathbf{W}} = \mathbf{W}\mathbf{R}$ for some $K \times K$ orthogonal rotation matrix \mathbf{R} then

$$\begin{aligned} p(\mathbf{x}_n) &= \mathcal{N}(\mathbf{0}, \tilde{\mathbf{W}}\tilde{\mathbf{W}}^\top + \sigma^2\mathbf{I}_D) \\ &= \mathcal{N}(\mathbf{0}, \mathbf{W}\mathbf{R}\mathbf{R}^\top\mathbf{W}^\top + \sigma^2\mathbf{I}_D) \\ &= \mathcal{N}(\mathbf{0}, \mathbf{W}\mathbf{W}^\top + \sigma^2\mathbf{I}_D) \end{aligned}$$

- Thus PPCA solution isn't unique (for every \mathbf{W} , there is a $\tilde{\mathbf{W}} = \mathbf{W}\mathbf{R}$ that gives the same solution)

Identifiability

- Note that for PPCA, we have $p(\mathbf{x}_n) = \mathcal{N}(\mathbf{0}, \mathbf{W}\mathbf{W}^\top + \sigma^2\mathbf{I}_D)$
- If we replace \mathbf{W} by $\tilde{\mathbf{W}} = \mathbf{W}\mathbf{R}$ for some $K \times K$ orthogonal rotation matrix \mathbf{R} then

$$\begin{aligned} p(\mathbf{x}_n) &= \mathcal{N}(\mathbf{0}, \tilde{\mathbf{W}}\tilde{\mathbf{W}}^\top + \sigma^2\mathbf{I}_D) \\ &= \mathcal{N}(\mathbf{0}, \mathbf{W}\mathbf{R}\mathbf{R}^\top\mathbf{W}^\top + \sigma^2\mathbf{I}_D) \\ &= \mathcal{N}(\mathbf{0}, \mathbf{W}\mathbf{W}^\top + \sigma^2\mathbf{I}_D) \end{aligned}$$

- Thus PPCA solution isn't unique (for every \mathbf{W} , there is a $\tilde{\mathbf{W}} = \mathbf{W}\mathbf{R}$ that gives the same solution)
- Thus the PPCA/FA model is **not uniquely identifiable**

Identifiability

- Note that for PPCA, we have $p(\mathbf{x}_n) = \mathcal{N}(\mathbf{0}, \mathbf{W}\mathbf{W}^\top + \sigma^2\mathbf{I}_D)$
- If we replace \mathbf{W} by $\tilde{\mathbf{W}} = \mathbf{W}\mathbf{R}$ for some $K \times K$ orthogonal rotation matrix \mathbf{R} then

$$\begin{aligned} p(\mathbf{x}_n) &= \mathcal{N}(\mathbf{0}, \tilde{\mathbf{W}}\tilde{\mathbf{W}}^\top + \sigma^2\mathbf{I}_D) \\ &= \mathcal{N}(\mathbf{0}, \mathbf{W}\mathbf{R}\mathbf{R}^\top\mathbf{W}^\top + \sigma^2\mathbf{I}_D) \\ &= \mathcal{N}(\mathbf{0}, \mathbf{W}\mathbf{W}^\top + \sigma^2\mathbf{I}_D) \end{aligned}$$

- Thus PPCA solution isn't unique (for every \mathbf{W} , there is a $\tilde{\mathbf{W}} = \mathbf{W}\mathbf{R}$ that gives the same solution)
- Thus the PPCA/FA model is **not uniquely identifiable**
- Usually this is not a problem, unless we want to very strictly interpret \mathbf{W}

Identifiability

- Note that for PPCA, we have $p(\mathbf{x}_n) = \mathcal{N}(\mathbf{0}, \mathbf{W}\mathbf{W}^\top + \sigma^2\mathbf{I}_D)$
- If we replace \mathbf{W} by $\tilde{\mathbf{W}} = \mathbf{W}\mathbf{R}$ for some $K \times K$ orthogonal rotation matrix \mathbf{R} then

$$\begin{aligned} p(\mathbf{x}_n) &= \mathcal{N}(\mathbf{0}, \tilde{\mathbf{W}}\tilde{\mathbf{W}}^\top + \sigma^2\mathbf{I}_D) \\ &= \mathcal{N}(\mathbf{0}, \mathbf{W}\mathbf{R}\mathbf{R}^\top\mathbf{W}^\top + \sigma^2\mathbf{I}_D) \\ &= \mathcal{N}(\mathbf{0}, \mathbf{W}\mathbf{W}^\top + \sigma^2\mathbf{I}_D) \end{aligned}$$

- Thus PPCA solution isn't unique (for every \mathbf{W} , there is a $\tilde{\mathbf{W}} = \mathbf{W}\mathbf{R}$ that gives the same solution)
- Thus the PPCA/FA model is **not uniquely identifiable**
- Usually this is not a problem, unless we want to very strictly interpret \mathbf{W}
- To ensure identifiability, there are some alternatives

Identifiability

- Note that for PPCA, we have $p(\mathbf{x}_n) = \mathcal{N}(\mathbf{0}, \mathbf{W}\mathbf{W}^\top + \sigma^2\mathbf{I}_D)$
- If we replace \mathbf{W} by $\tilde{\mathbf{W}} = \mathbf{W}\mathbf{R}$ for some $K \times K$ orthogonal rotation matrix \mathbf{R} then

$$\begin{aligned} p(\mathbf{x}_n) &= \mathcal{N}(\mathbf{0}, \tilde{\mathbf{W}}\tilde{\mathbf{W}}^\top + \sigma^2\mathbf{I}_D) \\ &= \mathcal{N}(\mathbf{0}, \mathbf{W}\mathbf{R}\mathbf{R}^\top\mathbf{W}^\top + \sigma^2\mathbf{I}_D) \\ &= \mathcal{N}(\mathbf{0}, \mathbf{W}\mathbf{W}^\top + \sigma^2\mathbf{I}_D) \end{aligned}$$

- Thus PPCA solution isn't unique (for every \mathbf{W} , there is a $\tilde{\mathbf{W}} = \mathbf{W}\mathbf{R}$ that gives the same solution)
- Thus the PPCA/FA model is **not uniquely identifiable**
- Usually this is not a problem, unless we want to very strictly interpret \mathbf{W}
- To ensure identifiability, there are some alternatives
 - Constrain \mathbf{W} it to be a **lower-triangular** or **sparse** matrix

Identifiability

- Note that for PPCA, we have $p(\mathbf{x}_n) = \mathcal{N}(\mathbf{0}, \mathbf{W}\mathbf{W}^\top + \sigma^2\mathbf{I}_D)$
- If we replace \mathbf{W} by $\tilde{\mathbf{W}} = \mathbf{W}\mathbf{R}$ for some $K \times K$ orthogonal rotation matrix \mathbf{R} then

$$\begin{aligned}p(\mathbf{x}_n) &= \mathcal{N}(\mathbf{0}, \tilde{\mathbf{W}}\tilde{\mathbf{W}}^\top + \sigma^2\mathbf{I}_D) \\&= \mathcal{N}(\mathbf{0}, \mathbf{W}\mathbf{R}\mathbf{R}^\top\mathbf{W}^\top + \sigma^2\mathbf{I}_D) \\&= \mathcal{N}(\mathbf{0}, \mathbf{W}\mathbf{W}^\top + \sigma^2\mathbf{I}_D)\end{aligned}$$

- Thus PPCA solution isn't unique (for every \mathbf{W} , there is a $\tilde{\mathbf{W}} = \mathbf{W}\mathbf{R}$ that gives the same solution)
- Thus the PPCA/FA model is **not uniquely identifiable**
- Usually this is not a problem, unless we want to very strictly interpret \mathbf{W}
- To ensure identifiability, there are some alternatives
 - Constrain \mathbf{W} it to be a **lower-triangular** or **sparse** matrix
 - Use **Independent Component Analysis** (ICA): ICA uses a **non-Gaussian prior** on \mathbf{z} to get identifiability

$$p(\mathbf{z}) = \prod_{k=1}^K p_k(z_k) \quad (\text{each } p_k \text{ is a non-Gaussian distr. like Laplace})$$

Parameter Estimation for PPCA via MLE

- Given data $\mathbf{X} = \{\mathbf{x}_n\}_{n=1}^N$, estimate latent vars: $\mathbf{Z} = \{\mathbf{z}_n\}_{n=1}^N$ and parameters $\Theta = (\mathbf{W}, \sigma^2)$

[†] Probabilistic Principal Component Analysis (Tipping and Bishop, 1999)

Parameter Estimation for PPCA via MLE

- Given data $\mathbf{X} = \{\mathbf{x}_n\}_{n=1}^N$, estimate latent vars: $\mathbf{Z} = \{\mathbf{z}_n\}_{n=1}^N$ and parameters $\Theta = (\mathbf{W}, \sigma^2)$
- If we only want $\Theta = (\mathbf{W}$ and $\sigma^2)$, we can do MLE directly on $p(\mathbf{X}|\Theta) = \mathcal{N}(0, \mathbf{W}\mathbf{W}^\top + \sigma^2\mathbf{I}_D)$

[†] Probabilistic Principal Component Analysis (Tipping and Bishop, 1999)

Parameter Estimation for PPCA via MLE

- Given data $\mathbf{X} = \{\mathbf{x}_n\}_{n=1}^N$, estimate latent vars: $\mathbf{Z} = \{\mathbf{z}_n\}_{n=1}^N$ and parameters $\Theta = (\mathbf{W}, \sigma^2)$
- If we only want $\Theta = (\mathbf{W}$ and $\sigma^2)$, we can do MLE directly on $p(\mathbf{X}|\Theta) = \mathcal{N}(0, \mathbf{W}\mathbf{W}^\top + \sigma^2\mathbf{I}_D)$
- Closed-form MLE solution[†] can be obtained for \mathbf{W} and σ^2 by maximizing

$$\log p(\mathbf{X}|\Theta) = -\frac{N}{2}(D \log 2\pi + \log |\mathbf{C}| + \text{trace}(\mathbf{C}^{-1}\mathbf{S})) \quad (\text{Verify})$$

where \mathbf{S} is the data covariance matrix, $\mathbf{C}^{-1} = \sigma^{-1}\mathbf{I} - \sigma^{-1}\mathbf{W}\mathbf{M}^{-1}\mathbf{W}^\top$ and $\mathbf{M} = \mathbf{W}^\top\mathbf{W} + \sigma^2\mathbf{I}$

[†] Probabilistic Principal Component Analysis (Tipping and Bishop, 1999)

Parameter Estimation for PPCA via MLE

- Given data $\mathbf{X} = \{\mathbf{x}_n\}_{n=1}^N$, estimate latent vars: $\mathbf{Z} = \{\mathbf{z}_n\}_{n=1}^N$ and parameters $\Theta = (\mathbf{W}, \sigma^2)$
- If we only want $\Theta = (\mathbf{W}$ and $\sigma^2)$, we can do MLE directly on $p(\mathbf{X}|\Theta) = \mathcal{N}(0, \mathbf{W}\mathbf{W}^\top + \sigma^2\mathbf{I}_D)$
- Closed-form MLE solution[†] can be obtained for \mathbf{W} and σ^2 by maximizing

$$\log p(\mathbf{X}|\Theta) = -\frac{N}{2}(D \log 2\pi + \log |\mathbf{C}| + \text{trace}(\mathbf{C}^{-1}\mathbf{S})) \quad (\text{Verify})$$

where \mathbf{S} is the data covariance matrix, $\mathbf{C}^{-1} = \sigma^{-1}\mathbf{I} - \sigma^{-1}\mathbf{W}\mathbf{M}^{-1}\mathbf{W}^\top$ and $\mathbf{M} = \mathbf{W}^\top\mathbf{W} + \sigma^2\mathbf{I}$

- The MLE solution is given by

$$\begin{aligned}\mathbf{W}_{ML} &= \mathbf{U}_K(\mathbf{L}_K - \sigma_{ML}^2\mathbf{I})^{1/2}\mathbf{R} \\ \sigma_{ML}^2 &= \frac{1}{D-K} \sum_{k=K+1}^D \lambda_k\end{aligned}$$

[†] Probabilistic Principal Component Analysis (Tipping and Bishop, 1999)

Parameter Estimation for PPCA via MLE

- Given data $\mathbf{X} = \{\mathbf{x}_n\}_{n=1}^N$, estimate latent vars: $\mathbf{Z} = \{\mathbf{z}_n\}_{n=1}^N$ and parameters $\Theta = (\mathbf{W}, \sigma^2)$
- If we only want $\Theta = (\mathbf{W}$ and $\sigma^2)$, we can do MLE directly on $p(\mathbf{X}|\Theta) = \mathcal{N}(0, \mathbf{W}\mathbf{W}^\top + \sigma^2\mathbf{I}_D)$
- Closed-form MLE solution[†] can be obtained for \mathbf{W} and σ^2 by maximizing

$$\log p(\mathbf{X}|\Theta) = -\frac{N}{2}(D \log 2\pi + \log |\mathbf{C}| + \text{trace}(\mathbf{C}^{-1}\mathbf{S})) \quad (\text{Verify})$$

where \mathbf{S} is the data covariance matrix, $\mathbf{C}^{-1} = \sigma^{-1}\mathbf{I} - \sigma^{-1}\mathbf{W}\mathbf{M}^{-1}\mathbf{W}^\top$ and $\mathbf{M} = \mathbf{W}^\top\mathbf{W} + \sigma^2\mathbf{I}$

- The MLE solution is given by

$$\begin{aligned}\mathbf{W}_{ML} &= \mathbf{U}_K(\mathbf{L}_K - \sigma_{ML}^2\mathbf{I})^{1/2}\mathbf{R} \\ \sigma_{ML}^2 &= \frac{1}{D-K} \sum_{k=K+1}^D \lambda_k\end{aligned}$$

where \mathbf{U}_K is $D \times K$ matrix of **top K eigvecs** of \mathbf{S} ,

[†] Probabilistic Principal Component Analysis (Tipping and Bishop, 1999)

Parameter Estimation for PPCA via MLE

- Given data $\mathbf{X} = \{\mathbf{x}_n\}_{n=1}^N$, estimate latent vars: $\mathbf{Z} = \{\mathbf{z}_n\}_{n=1}^N$ and parameters $\Theta = (\mathbf{W}, \sigma^2)$
- If we only want $\Theta = (\mathbf{W}$ and $\sigma^2)$, we can do MLE directly on $p(\mathbf{X}|\Theta) = \mathcal{N}(0, \mathbf{W}\mathbf{W}^\top + \sigma^2\mathbf{I}_D)$
- Closed-form MLE solution[†] can be obtained for \mathbf{W} and σ^2 by maximizing

$$\log p(\mathbf{X}|\Theta) = -\frac{N}{2}(D \log 2\pi + \log |\mathbf{C}| + \text{trace}(\mathbf{C}^{-1}\mathbf{S})) \quad (\text{Verify})$$

where \mathbf{S} is the data covariance matrix, $\mathbf{C}^{-1} = \sigma^{-1}\mathbf{I} - \sigma^{-1}\mathbf{W}\mathbf{M}^{-1}\mathbf{W}^\top$ and $\mathbf{M} = \mathbf{W}^\top\mathbf{W} + \sigma^2\mathbf{I}$

- The MLE solution is given by

$$\begin{aligned}\mathbf{W}_{ML} &= \mathbf{U}_K(\mathbf{L}_K - \sigma_{ML}^2\mathbf{I})^{1/2}\mathbf{R} \\ \sigma_{ML}^2 &= \frac{1}{D-K} \sum_{k=K+1}^D \lambda_k\end{aligned}$$

where \mathbf{U}_K is $D \times K$ matrix of top K eigvecs of \mathbf{S} , \mathbf{L}_K : $K \times K$ diagonal matrix of top K eigvals $\lambda_1, \dots, \lambda_K$,

[†] Probabilistic Principal Component Analysis (Tipping and Bishop, 1999)

Parameter Estimation for PPCA via MLE

- Given data $\mathbf{X} = \{\mathbf{x}_n\}_{n=1}^N$, estimate latent vars: $\mathbf{Z} = \{\mathbf{z}_n\}_{n=1}^N$ and parameters $\Theta = (\mathbf{W}, \sigma^2)$
- If we only want $\Theta = (\mathbf{W}$ and $\sigma^2)$, we can do MLE directly on $p(\mathbf{X}|\Theta) = \mathcal{N}(0, \mathbf{W}\mathbf{W}^\top + \sigma^2\mathbf{I}_D)$
- Closed-form MLE solution[†] can be obtained for \mathbf{W} and σ^2 by maximizing

$$\log p(\mathbf{X}|\Theta) = -\frac{N}{2}(D \log 2\pi + \log |\mathbf{C}| + \text{trace}(\mathbf{C}^{-1}\mathbf{S})) \quad (\text{Verify})$$

where \mathbf{S} is the data covariance matrix, $\mathbf{C}^{-1} = \sigma^{-2}\mathbf{I} - \sigma^{-2}\mathbf{W}\mathbf{M}^{-1}\mathbf{W}^\top$ and $\mathbf{M} = \mathbf{W}^\top\mathbf{W} + \sigma^2\mathbf{I}$

- The MLE solution is given by

$$\begin{aligned}\mathbf{W}_{ML} &= \mathbf{U}_K(\mathbf{L}_K - \sigma_{ML}^2\mathbf{I})^{1/2}\mathbf{R} \\ \sigma_{ML}^2 &= \frac{1}{D-K} \sum_{k=K+1}^D \lambda_k\end{aligned}$$

where \mathbf{U}_K is $D \times K$ matrix of top K eigvecs of \mathbf{S} , \mathbf{L}_K : $K \times K$ diagonal matrix of top K eigvals $\lambda_1, \dots, \lambda_K$, \mathbf{R} is a $K \times K$ arbitrary rotation matrix

[†] Probabilistic Principal Component Analysis (Tipping and Bishop, 1999)

Parameter Estimation for PPCA via MLE

- Given data $\mathbf{X} = \{\mathbf{x}_n\}_{n=1}^N$, estimate latent vars: $\mathbf{Z} = \{\mathbf{z}_n\}_{n=1}^N$ and parameters $\Theta = (\mathbf{W}, \sigma^2)$
- If we only want $\Theta = (\mathbf{W}$ and $\sigma^2)$, we can do MLE directly on $p(\mathbf{X}|\Theta) = \mathcal{N}(0, \mathbf{W}\mathbf{W}^\top + \sigma^2\mathbf{I}_D)$
- Closed-form MLE solution[†] can be obtained for \mathbf{W} and σ^2 by maximizing

$$\log p(\mathbf{X}|\Theta) = -\frac{N}{2}(D \log 2\pi + \log |\mathbf{C}| + \text{trace}(\mathbf{C}^{-1}\mathbf{S})) \quad (\text{Verify})$$

where \mathbf{S} is the data covariance matrix, $\mathbf{C}^{-1} = \sigma^{-2}\mathbf{I} - \sigma^{-2}\mathbf{W}\mathbf{M}^{-1}\mathbf{W}^\top$ and $\mathbf{M} = \mathbf{W}^\top\mathbf{W} + \sigma^2\mathbf{I}$

- The MLE solution is given by

$$\begin{aligned}\mathbf{W}_{ML} &= \mathbf{U}_K(\mathbf{L}_K - \sigma_{ML}^2\mathbf{I})^{1/2}\mathbf{R} \\ \sigma_{ML}^2 &= \frac{1}{D-K} \sum_{k=K+1}^D \lambda_k\end{aligned}$$

where \mathbf{U}_K is $D \times K$ matrix of **top K eigvecs** of \mathbf{S} , \mathbf{L}_K : $K \times K$ diagonal matrix of **top K eigvals** $\lambda_1, \dots, \lambda_K$, \mathbf{R} is a $K \times K$ **arbitrary rotation matrix** (equivalent to PCA for $\mathbf{R} = \mathbf{I}$ and $\sigma^2 \rightarrow 0$)

[†] Probabilistic Principal Component Analysis (Tipping and Bishop, 1999)

Parameter Estimation for PPCA via EM

- First of all, why bother doing EM here? Didn't MLE already give us a closed-form solution?

Parameter Estimation for PPCA via EM

- First of all, why bother doing EM here? Didn't MLE already give us a closed-form solution?

$$\begin{aligned}\mathbf{W}_{ML} &= \mathbf{U}_K(\mathbf{L}_K - \sigma_{ML}^2 \mathbf{I})^{1/2} \mathbf{R} \\ \sigma_{ML}^2 &= \frac{1}{D-K} \sum_{k=K+1}^D \lambda_k\end{aligned}$$

Parameter Estimation for PPCA via EM

- First of all, why bother doing EM here? Didn't MLE already give us a closed-form solution?

$$\begin{aligned}\mathbf{W}_{ML} &= \mathbf{U}_K(\mathbf{L}_K - \sigma_{ML}^2 \mathbf{I})^{1/2} \mathbf{R} \\ \sigma_{ML}^2 &= \frac{1}{D-K} \sum_{k=K+1}^D \lambda_k\end{aligned}$$

- Well, several reasons that make EM appealing for PPCA

Parameter Estimation for PPCA via EM

- First of all, why bother doing EM here? Didn't MLE already give us a closed-form solution?

$$\begin{aligned}\mathbf{W}_{ML} &= \mathbf{U}_K(\mathbf{L}_K - \sigma_{ML}^2 \mathbf{I})^{1/2} \mathbf{R} \\ \sigma_{ML}^2 &= \frac{1}{D - K} \sum_{k=K+1}^D \lambda_k\end{aligned}$$

- Well, several reasons that make EM appealing for PPCA
 - MLE for PPCA becomes much easier (and **computationally efficient**) when using EM

Parameter Estimation for PPCA via EM

- First of all, why bother doing EM here? Didn't MLE already give us a closed-form solution?

$$\begin{aligned}\mathbf{W}_{ML} &= \mathbf{U}_K(\mathbf{L}_K - \sigma_{ML}^2 \mathbf{I})^{1/2} \mathbf{R} \\ \sigma_{ML}^2 &= \frac{1}{D-K} \sum_{k=K+1}^D \lambda_k\end{aligned}$$

- Well, several reasons that make EM appealing for PPCA
 - MLE for PPCA becomes much easier (and **computationally efficient**) when using EM
 - EM allows directly **inferring the latent vars.** \mathbf{z}_n 's (MLE only gives parameters Θ)

Parameter Estimation for PPCA via EM

- First of all, why bother doing EM here? Didn't MLE already give us a closed-form solution?

$$\begin{aligned}\mathbf{W}_{ML} &= \mathbf{U}_K(\mathbf{L}_K - \sigma_{ML}^2 \mathbf{I})^{1/2} \mathbf{R} \\ \sigma_{ML}^2 &= \frac{1}{D - K} \sum_{k=K+1}^D \lambda_k\end{aligned}$$

- Well, several reasons that make EM appealing for PPCA
 - MLE for PPCA becomes much easier (and **computationally efficient**) when using EM
 - EM allows directly **inferring the latent vars.** \mathbf{z}_n 's (MLE only gives parameters Θ)
 - No need to construct \mathbf{S} or do **eigen-decomposition** of \mathbf{S} (can be expensive for large D and N)

Parameter Estimation for PPCA via EM

- First of all, why bother doing EM here? Didn't MLE already give us a closed-form solution?

$$\begin{aligned}\mathbf{W}_{ML} &= \mathbf{U}_K(\mathbf{L}_K - \sigma_{ML}^2 \mathbf{I})^{1/2} \mathbf{R} \\ \sigma_{ML}^2 &= \frac{1}{D - K} \sum_{k=K+1}^D \lambda_k\end{aligned}$$

- Well, several reasons that make EM appealing for PPCA
 - MLE for PPCA becomes much easier (and **computationally efficient**) when using EM
 - EM allows directly **inferring the latent vars.** \mathbf{z}_n 's (MLE only gives parameters Θ)
 - No need to construct \mathbf{S} or do **eigen-decomposition** of \mathbf{S} (can be expensive for large D and N)
 - Closed-form MLE **not even possible** for other PPCA variant (e.g., FA where $p(\epsilon) = \mathcal{N}(\mathbf{0}, \Psi)$)

Parameter Estimation for PPCA via EM

- First of all, why bother doing EM here? Didn't MLE already give us a closed-form solution?

$$\begin{aligned}\mathbf{W}_{ML} &= \mathbf{U}_K(\mathbf{L}_K - \sigma_{ML}^2 \mathbf{I})^{1/2} \mathbf{R} \\ \sigma_{ML}^2 &= \frac{1}{D-K} \sum_{k=K+1}^D \lambda_k\end{aligned}$$

- Well, several reasons that make EM appealing for PPCA
 - MLE for PPCA becomes much easier (and **computationally efficient**) when using EM
 - EM allows directly **inferring the latent vars.** \mathbf{z}_n 's (MLE only gives parameters Θ)
 - No need to construct \mathbf{S} or do **eigen-decomposition** of \mathbf{S} (can be expensive for large D and N)
 - Closed-form MLE **not even possible** for other PPCA variant (e.g., FA where $p(\epsilon) = \mathcal{N}(\mathbf{0}, \Psi)$)
 - EM can properly **handle missing data**, e.g., by computing $p(\mathbf{x}_n^{\text{missing}} | \mathbf{x}_n^{\text{obs}})$; often doable easily

Parameter Estimation for PPCA via EM

- First of all, why bother doing EM here? Didn't MLE already give us a closed-form solution?

$$\begin{aligned}\mathbf{W}_{ML} &= \mathbf{U}_K(\mathbf{L}_K - \sigma_{ML}^2 \mathbf{I})^{1/2} \mathbf{R} \\ \sigma_{ML}^2 &= \frac{1}{D - K} \sum_{k=K+1}^D \lambda_k\end{aligned}$$

- Well, several reasons that make EM appealing for PPCA
 - MLE for PPCA becomes much easier (and **computationally efficient**) when using EM
 - EM allows directly **inferring the latent vars.** \mathbf{z}_n 's (MLE only gives parameters Θ)
 - No need to construct \mathbf{S} or do **eigen-decomposition** of \mathbf{S} (can be expensive for large D and N)
 - Closed-form MLE **not even possible** for other PPCA variant (e.g., FA where $p(\epsilon) = \mathcal{N}(\mathbf{0}, \Psi)$)
 - EM can properly **handle missing data**, e.g., by computing $p(\mathbf{x}_n^{\text{missing}} | \mathbf{x}_n^{\text{obs}})$; often doable easily
 - EM can easily be **made online** (enables handling large D and large N)

Parameter Estimation for PPCA via EM

- EM for PPCA iterates between the following two steps

Parameter Estimation for PPCA via EM

- EM for PPCA iterates between the following two steps
 - E Step: Infer the **posterior** $p(\mathbf{z}_n | \mathbf{x}_n)$ given current estimate of $\Theta = (\mathbf{W}, \sigma^2)$

Parameter Estimation for PPCA via EM

- EM for PPCA iterates between the following two steps
 - E Step: Infer the **posterior** $p(\mathbf{z}_n | \mathbf{x}_n)$ given current estimate of $\Theta = (\mathbf{W}, \sigma^2)$ (needed for expectations)

Parameter Estimation for PPCA via EM

- EM for PPCA iterates between the following two steps
 - E Step: Infer the **posterior** $p(\mathbf{z}_n|\mathbf{x}_n)$ given current estimate of $\Theta = (\mathbf{W}, \sigma^2)$ (needed for expectations)

$$p(\mathbf{z}_n|\mathbf{x}_n, \mathbf{W}, \sigma^2) = \mathcal{N}(\mathbf{M}^{-1}\mathbf{W}^\top \mathbf{x}_n, \sigma^2 \mathbf{M}^{-1}) \quad (\text{where } \mathbf{M} = \mathbf{W}^\top \mathbf{W} + \sigma^2 \mathbf{I}_K)$$

Parameter Estimation for PPCA via EM

- EM for PPCA iterates between the following two steps
 - E Step: Infer the **posterior** $p(\mathbf{z}_n|\mathbf{x}_n)$ given current estimate of $\Theta = (\mathbf{W}, \sigma^2)$ (needed for expectations)

$$p(\mathbf{z}_n|\mathbf{x}_n, \mathbf{W}, \sigma^2) = \mathcal{N}(\mathbf{M}^{-1}\mathbf{W}^\top \mathbf{x}_n, \sigma^2 \mathbf{M}^{-1}) \quad (\text{where } \mathbf{M} = \mathbf{W}^\top \mathbf{W} + \sigma^2 \mathbf{I}_K)$$

- M Step: Maximize the **expected complete data log-lik.** (CLL) $\mathbb{E}[\log p(\mathbf{X}, \mathbf{Z}|\Theta)]$ w.r.t. Θ

Parameter Estimation for PPCA via EM

- EM for PPCA iterates between the following two steps
 - E Step: Infer the **posterior** $p(\mathbf{z}_n|\mathbf{x}_n)$ given current estimate of $\Theta = (\mathbf{W}, \sigma^2)$ (needed for expectations)

$$p(\mathbf{z}_n|\mathbf{x}_n, \mathbf{W}, \sigma^2) = \mathcal{N}(\mathbf{M}^{-1}\mathbf{W}^\top \mathbf{x}_n, \sigma^2 \mathbf{M}^{-1}) \quad (\text{where } \mathbf{M} = \mathbf{W}^\top \mathbf{W} + \sigma^2 \mathbf{I}_K)$$

- M Step: Maximize the **expected complete data log-lik.** (CLL) $\mathbb{E}[\log p(\mathbf{X}, \mathbf{Z}|\Theta)]$ w.r.t. Θ
- The CLL (and expected CLL) for PPCA has a simple expression.

Parameter Estimation for PPCA via EM

- EM for PPCA iterates between the following two steps
 - E Step: Infer the **posterior** $p(\mathbf{z}_n|\mathbf{x}_n)$ given current estimate of $\Theta = (\mathbf{W}, \sigma^2)$ (needed for expectations)

$$p(\mathbf{z}_n|\mathbf{x}_n, \mathbf{W}, \sigma^2) = \mathcal{N}(\mathbf{M}^{-1}\mathbf{W}^\top \mathbf{x}_n, \sigma^2 \mathbf{M}^{-1}) \quad (\text{where } \mathbf{M} = \mathbf{W}^\top \mathbf{W} + \sigma^2 \mathbf{I}_K)$$

- M Step: Maximize the **expected complete data log-lik.** (CLL) $\mathbb{E}[\log p(\mathbf{X}, \mathbf{Z}|\Theta)]$ w.r.t. Θ
- The CLL (and expected CLL) for PPCA has a simple expression. The CLL is

Parameter Estimation for PPCA via EM

- EM for PPCA iterates between the following two steps
 - E Step: Infer the **posterior** $p(\mathbf{z}_n|\mathbf{x}_n)$ given current estimate of $\Theta = (\mathbf{W}, \sigma^2)$ (needed for expectations)

$$p(\mathbf{z}_n|\mathbf{x}_n, \mathbf{W}, \sigma^2) = \mathcal{N}(\mathbf{M}^{-1}\mathbf{W}^\top \mathbf{x}_n, \sigma^2 \mathbf{M}^{-1}) \quad (\text{where } \mathbf{M} = \mathbf{W}^\top \mathbf{W} + \sigma^2 \mathbf{I}_K)$$

- M Step: Maximize the **expected complete data log-lik.** (CLL) $\mathbb{E}[\log p(\mathbf{X}, \mathbf{Z}|\Theta)]$ w.r.t. Θ
- The CLL (and expected CLL) for PPCA has a simple expression. The CLL is

$$\log p(\mathbf{X}, \mathbf{Z}|\mathbf{W}, \sigma^2)$$

Parameter Estimation for PPCA via EM

- EM for PPCA iterates between the following two steps
 - E Step: Infer the **posterior** $p(\mathbf{z}_n|\mathbf{x}_n)$ given current estimate of $\Theta = (\mathbf{W}, \sigma^2)$ (needed for expectations)

$$p(\mathbf{z}_n|\mathbf{x}_n, \mathbf{W}, \sigma^2) = \mathcal{N}(\mathbf{M}^{-1}\mathbf{W}^\top \mathbf{x}_n, \sigma^2 \mathbf{M}^{-1}) \quad (\text{where } \mathbf{M} = \mathbf{W}^\top \mathbf{W} + \sigma^2 \mathbf{I}_K)$$

- M Step: Maximize the **expected complete data log-lik.** (CLL) $\mathbb{E}[\log p(\mathbf{X}, \mathbf{Z}|\Theta)]$ w.r.t. Θ
- The CLL (and expected CLL) for PPCA has a simple expression. The CLL is

$$\log p(\mathbf{X}, \mathbf{Z}|\mathbf{W}, \sigma^2) = \log \prod_{n=1}^N p(\mathbf{x}_n, \mathbf{z}_n|\mathbf{W}, \sigma^2)$$

Parameter Estimation for PPCA via EM

- EM for PPCA iterates between the following two steps
 - E Step: Infer the **posterior** $p(\mathbf{z}_n|\mathbf{x}_n)$ given current estimate of $\Theta = (\mathbf{W}, \sigma^2)$ (needed for expectations)

$$p(\mathbf{z}_n|\mathbf{x}_n, \mathbf{W}, \sigma^2) = \mathcal{N}(\mathbf{M}^{-1}\mathbf{W}^\top \mathbf{x}_n, \sigma^2 \mathbf{M}^{-1}) \quad (\text{where } \mathbf{M} = \mathbf{W}^\top \mathbf{W} + \sigma^2 \mathbf{I}_K)$$

- M Step: Maximize the **expected complete data log-lik.** (CLL) $\mathbb{E}[\log p(\mathbf{X}, \mathbf{Z}|\Theta)]$ w.r.t. Θ
- The CLL (and expected CLL) for PPCA has a simple expression. The CLL is

$$\log p(\mathbf{X}, \mathbf{Z}|\mathbf{W}, \sigma^2) = \log \prod_{n=1}^N p(\mathbf{x}_n, \mathbf{z}_n|\mathbf{W}, \sigma^2) = \log \prod_{n=1}^N p(\mathbf{x}_n|\mathbf{z}_n, \mathbf{W}, \sigma^2)p(\mathbf{z}_n)$$

Parameter Estimation for PPCA via EM

- EM for PPCA iterates between the following two steps
 - E Step: Infer the **posterior** $p(\mathbf{z}_n|\mathbf{x}_n)$ given current estimate of $\Theta = (\mathbf{W}, \sigma^2)$ (needed for expectations)

$$p(\mathbf{z}_n|\mathbf{x}_n, \mathbf{W}, \sigma^2) = \mathcal{N}(\mathbf{M}^{-1}\mathbf{W}^\top \mathbf{x}_n, \sigma^2 \mathbf{M}^{-1}) \quad (\text{where } \mathbf{M} = \mathbf{W}^\top \mathbf{W} + \sigma^2 \mathbf{I}_K)$$

- M Step: Maximize the **expected complete data log-lik.** (CLL) $\mathbb{E}[\log p(\mathbf{X}, \mathbf{Z}|\Theta)]$ w.r.t. Θ
- The CLL (and expected CLL) for PPCA has a simple expression. The CLL is

$$\log p(\mathbf{X}, \mathbf{Z}|\mathbf{W}, \sigma^2) = \log \prod_{n=1}^N p(\mathbf{x}_n, \mathbf{z}_n|\mathbf{W}, \sigma^2) = \log \prod_{n=1}^N p(\mathbf{x}_n|\mathbf{z}_n, \mathbf{W}, \sigma^2)p(\mathbf{z}_n) = \sum_{n=1}^N \{\log p(\mathbf{x}_n|\mathbf{z}_n, \mathbf{W}, \sigma^2) + \log p(\mathbf{z}_n)\}$$

Parameter Estimation for PPCA via EM

- EM for PPCA iterates between the following two steps
 - E Step: Infer the **posterior** $p(\mathbf{z}_n|\mathbf{x}_n)$ given current estimate of $\Theta = (\mathbf{W}, \sigma^2)$ (needed for expectations)

$$p(\mathbf{z}_n|\mathbf{x}_n, \mathbf{W}, \sigma^2) = \mathcal{N}(\mathbf{M}^{-1}\mathbf{W}^\top \mathbf{x}_n, \sigma^2 \mathbf{M}^{-1}) \quad (\text{where } \mathbf{M} = \mathbf{W}^\top \mathbf{W} + \sigma^2 \mathbf{I}_K)$$

- M Step: Maximize the **expected complete data log-lik.** (CLL) $\mathbb{E}[\log p(\mathbf{X}, \mathbf{Z}|\Theta)]$ w.r.t. Θ
- The CLL (and expected CLL) for PPCA has a simple expression. The CLL is

$$\log p(\mathbf{X}, \mathbf{Z}|\mathbf{W}, \sigma^2) = \log \prod_{n=1}^N p(\mathbf{x}_n, \mathbf{z}_n|\mathbf{W}, \sigma^2) = \log \prod_{n=1}^N p(\mathbf{x}_n|\mathbf{z}_n, \mathbf{W}, \sigma^2)p(\mathbf{z}_n) = \sum_{n=1}^N \{\log p(\mathbf{x}_n|\mathbf{z}_n, \mathbf{W}, \sigma^2) + \log p(\mathbf{z}_n)\}$$

- Using $p(\mathbf{x}_n|\mathbf{z}_n, \mathbf{W}, \sigma^2) = \frac{1}{(2\pi\sigma^2)^{D/2}} \exp \left[-\frac{(\mathbf{x}_n - \mathbf{W}\mathbf{z}_n)^\top (\mathbf{x}_n - \mathbf{W}\mathbf{z}_n)}{2\sigma^2} \right]$

Parameter Estimation for PPCA via EM

- EM for PPCA iterates between the following two steps
 - E Step: Infer the **posterior** $p(\mathbf{z}_n|\mathbf{x}_n)$ given current estimate of $\Theta = (\mathbf{W}, \sigma^2)$ (needed for expectations)

$$p(\mathbf{z}_n|\mathbf{x}_n, \mathbf{W}, \sigma^2) = \mathcal{N}(\mathbf{M}^{-1}\mathbf{W}^\top \mathbf{x}_n, \sigma^2 \mathbf{M}^{-1}) \quad (\text{where } \mathbf{M} = \mathbf{W}^\top \mathbf{W} + \sigma^2 \mathbf{I}_K)$$

- M Step: Maximize the **expected complete data log-lik.** (CLL) $\mathbb{E}[\log p(\mathbf{X}, \mathbf{Z}|\Theta)]$ w.r.t. Θ
- The CLL (and expected CLL) for PPCA has a simple expression. The CLL is

$$\log p(\mathbf{X}, \mathbf{Z}|\mathbf{W}, \sigma^2) = \log \prod_{n=1}^N p(\mathbf{x}_n, \mathbf{z}_n|\mathbf{W}, \sigma^2) = \log \prod_{n=1}^N p(\mathbf{x}_n|\mathbf{z}_n, \mathbf{W}, \sigma^2)p(\mathbf{z}_n) = \sum_{n=1}^N \{\log p(\mathbf{x}_n|\mathbf{z}_n, \mathbf{W}, \sigma^2) + \log p(\mathbf{z}_n)\}$$

- Using $p(\mathbf{x}_n|\mathbf{z}_n, \mathbf{W}, \sigma^2) = \frac{1}{(2\pi\sigma^2)^{D/2}} \exp \left[-\frac{(\mathbf{x}_n - \mathbf{W}\mathbf{z}_n)^\top (\mathbf{x}_n - \mathbf{W}\mathbf{z}_n)}{2\sigma^2} \right]$ and $p(\mathbf{z}_n) \propto \exp \left[-\frac{\mathbf{z}_n^\top \mathbf{z}_n}{2} \right]$

Parameter Estimation for PPCA via EM

- EM for PPCA iterates between the following two steps
 - E Step: Infer the **posterior** $p(\mathbf{z}_n|\mathbf{x}_n)$ given current estimate of $\Theta = (\mathbf{W}, \sigma^2)$ (needed for expectations)

$$p(\mathbf{z}_n|\mathbf{x}_n, \mathbf{W}, \sigma^2) = \mathcal{N}(\mathbf{M}^{-1}\mathbf{W}^\top \mathbf{x}_n, \sigma^2 \mathbf{M}^{-1}) \quad (\text{where } \mathbf{M} = \mathbf{W}^\top \mathbf{W} + \sigma^2 \mathbf{I}_K)$$

- M Step: Maximize the **expected complete data log-lik.** (CLL) $\mathbb{E}[\log p(\mathbf{X}, \mathbf{Z}|\Theta)]$ w.r.t. Θ
- The CLL (and expected CLL) for PPCA has a simple expression. The CLL is

$$\log p(\mathbf{X}, \mathbf{Z}|\mathbf{W}, \sigma^2) = \log \prod_{n=1}^N p(\mathbf{x}_n, \mathbf{z}_n|\mathbf{W}, \sigma^2) = \log \prod_{n=1}^N p(\mathbf{x}_n|\mathbf{z}_n, \mathbf{W}, \sigma^2)p(\mathbf{z}_n) = \sum_{n=1}^N \{\log p(\mathbf{x}_n|\mathbf{z}_n, \mathbf{W}, \sigma^2) + \log p(\mathbf{z}_n)\}$$

- Using $p(\mathbf{x}_n|\mathbf{z}_n, \mathbf{W}, \sigma^2) = \frac{1}{(2\pi\sigma^2)^{D/2}} \exp\left[-\frac{(\mathbf{x}_n - \mathbf{W}\mathbf{z}_n)^\top (\mathbf{x}_n - \mathbf{W}\mathbf{z}_n)}{2\sigma^2}\right]$ and $p(\mathbf{z}_n) \propto \exp\left[-\frac{\mathbf{z}_n^\top \mathbf{z}_n}{2}\right]$ and simplifying

$$\text{CLL} = -\sum_{n=1}^N \left\{ \frac{D}{2} \log \sigma^2 + \frac{1}{2\sigma^2} \|\mathbf{x}_n\|^2 - \frac{1}{\sigma^2} \mathbf{z}_n^\top \mathbf{W}^\top \mathbf{x}_n + \frac{1}{2\sigma^2} \text{tr}(\mathbf{z}_n \mathbf{z}_n^\top \mathbf{W}^\top \mathbf{W}) + \frac{1}{2} \text{tr}(\mathbf{z}_n \mathbf{z}_n^\top) \right\} \quad (\text{Exercise: Verify})$$

Parameter Estimation for PPCA via EM

- The expected complete data log-likelihood $\mathbb{E}[\log p(\mathbf{X}, \mathbf{Z} | \mathbf{W}, \sigma^2)]$

$$= - \sum_{n=1}^N \left\{ \frac{D}{2} \log \sigma^2 + \frac{1}{2\sigma^2} \|\mathbf{x}_n\|^2 - \frac{1}{\sigma^2} \mathbb{E}[\mathbf{z}_n]^\top \mathbf{W}^\top \mathbf{x}_n + \frac{1}{2\sigma^2} \text{tr}(\mathbb{E}[\mathbf{z}_n \mathbf{z}_n^\top] \mathbf{W}^\top \mathbf{W}) + \frac{1}{2} \text{tr}(\mathbb{E}[\mathbf{z}_n \mathbf{z}_n^\top]) \right\}$$

Parameter Estimation for PPCA via EM

- The expected complete data log-likelihood $\mathbb{E}[\log p(\mathbf{X}, \mathbf{Z}|\mathbf{W}, \sigma^2)]$

$$= - \sum_{n=1}^N \left\{ \frac{D}{2} \log \sigma^2 + \frac{1}{2\sigma^2} \|\mathbf{x}_n\|^2 - \frac{1}{\sigma^2} \mathbb{E}[\mathbf{z}_n]^\top \mathbf{W}^\top \mathbf{x}_n + \frac{1}{2\sigma^2} \text{tr}(\mathbb{E}[\mathbf{z}_n \mathbf{z}_n^\top] \mathbf{W}^\top \mathbf{W}) + \frac{1}{2} \text{tr}(\mathbb{E}[\mathbf{z}_n \mathbf{z}_n^\top]) \right\}$$

- Taking the derivative of $\mathbb{E}[\log p(\mathbf{X}, \mathbf{Z}|\mathbf{W}, \sigma^2)]$ w.r.t. \mathbf{W} and setting to zero

$$\mathbf{W} = \left[\sum_{n=1}^N \mathbf{x}_n \mathbb{E}[\mathbf{z}_n]^\top \right] \left[\sum_{n=1}^N \mathbb{E}[\mathbf{z}_n \mathbf{z}_n^\top] \right]^{-1}$$

(Exercise: verify; can also be done “online”)

Parameter Estimation for PPCA via EM

- The expected complete data log-likelihood $\mathbb{E}[\log p(\mathbf{X}, \mathbf{Z}|\mathbf{W}, \sigma^2)]$

$$= - \sum_{n=1}^N \left\{ \frac{D}{2} \log \sigma^2 + \frac{1}{2\sigma^2} \|\mathbf{x}_n\|^2 - \frac{1}{\sigma^2} \mathbb{E}[\mathbf{z}_n]^\top \mathbf{W}^\top \mathbf{x}_n + \frac{1}{2\sigma^2} \text{tr}(\mathbb{E}[\mathbf{z}_n \mathbf{z}_n^\top] \mathbf{W}^\top \mathbf{W}) + \frac{1}{2} \text{tr}(\mathbb{E}[\mathbf{z}_n \mathbf{z}_n^\top]) \right\}$$

- Taking the derivative of $\mathbb{E}[\log p(\mathbf{X}, \mathbf{Z}|\mathbf{W}, \sigma^2)]$ w.r.t. \mathbf{W} and setting to zero

$$\mathbf{W} = \left[\sum_{n=1}^N \mathbf{x}_n \mathbb{E}[\mathbf{z}_n]^\top \right] \left[\sum_{n=1}^N \mathbb{E}[\mathbf{z}_n \mathbf{z}_n^\top] \right]^{-1} \quad (\text{Exercise: verify; can also be done "online"})$$

- To compute \mathbf{W} , we need two posterior expectations $\mathbb{E}[\mathbf{z}_n]$ and $\mathbb{E}[\mathbf{z}_n \mathbf{z}_n^\top]$

Parameter Estimation for PPCA via EM

- The expected complete data log-likelihood $\mathbb{E}[\log p(\mathbf{X}, \mathbf{Z}|\mathbf{W}, \sigma^2)]$

$$= - \sum_{n=1}^N \left\{ \frac{D}{2} \log \sigma^2 + \frac{1}{2\sigma^2} \|\mathbf{x}_n\|^2 - \frac{1}{\sigma^2} \mathbb{E}[\mathbf{z}_n]^\top \mathbf{W}^\top \mathbf{x}_n + \frac{1}{2\sigma^2} \text{tr}(\mathbb{E}[\mathbf{z}_n \mathbf{z}_n^\top] \mathbf{W}^\top \mathbf{W}) + \frac{1}{2} \text{tr}(\mathbb{E}[\mathbf{z}_n \mathbf{z}_n^\top]) \right\}$$

- Taking the derivative of $\mathbb{E}[\log p(\mathbf{X}, \mathbf{Z}|\mathbf{W}, \sigma^2)]$ w.r.t. \mathbf{W} and setting to zero

$$\mathbf{W} = \left[\sum_{n=1}^N \mathbf{x}_n \mathbb{E}[\mathbf{z}_n]^\top \right] \left[\sum_{n=1}^N \mathbb{E}[\mathbf{z}_n \mathbf{z}_n^\top] \right]^{-1} \quad (\text{Exercise: verify; can also be done "online"})$$

- To compute \mathbf{W} , we need two posterior expectations $\mathbb{E}[\mathbf{z}_n]$ and $\mathbb{E}[\mathbf{z}_n \mathbf{z}_n^\top]$
- These can be easily obtained from the posterior $p(\mathbf{z}_n|\mathbf{x}_n)$ computed in E step

$$p(\mathbf{z}_n|\mathbf{x}_n, \mathbf{W}) = \mathcal{N}(\mathbf{M}^{-1} \mathbf{W}^\top \mathbf{x}_n, \sigma^2 \mathbf{M}^{-1}) \quad \text{where } \mathbf{M} = \mathbf{W}^\top \mathbf{W} + \sigma^2 \mathbf{I}_K$$

Parameter Estimation for PPCA via EM

- The expected complete data log-likelihood $\mathbb{E}[\log p(\mathbf{X}, \mathbf{Z}|\mathbf{W}, \sigma^2)]$

$$= - \sum_{n=1}^N \left\{ \frac{D}{2} \log \sigma^2 + \frac{1}{2\sigma^2} \|\mathbf{x}_n\|^2 - \frac{1}{\sigma^2} \mathbb{E}[\mathbf{z}_n]^\top \mathbf{W}^\top \mathbf{x}_n + \frac{1}{2\sigma^2} \text{tr}(\mathbb{E}[\mathbf{z}_n \mathbf{z}_n^\top] \mathbf{W}^\top \mathbf{W}) + \frac{1}{2} \text{tr}(\mathbb{E}[\mathbf{z}_n \mathbf{z}_n^\top]) \right\}$$

- Taking the derivative of $\mathbb{E}[\log p(\mathbf{X}, \mathbf{Z}|\mathbf{W}, \sigma^2)]$ w.r.t. \mathbf{W} and setting to zero

$$\mathbf{W} = \left[\sum_{n=1}^N \mathbf{x}_n \mathbb{E}[\mathbf{z}_n]^\top \right] \left[\sum_{n=1}^N \mathbb{E}[\mathbf{z}_n \mathbf{z}_n^\top] \right]^{-1} \quad (\text{Exercise: verify; can also be done "online"})$$

- To compute \mathbf{W} , we need two posterior expectations $\mathbb{E}[\mathbf{z}_n]$ and $\mathbb{E}[\mathbf{z}_n \mathbf{z}_n^\top]$
- These can be easily obtained from the posterior $p(\mathbf{z}_n|\mathbf{x}_n)$ computed in E step

$$\begin{aligned} p(\mathbf{z}_n|\mathbf{x}_n, \mathbf{W}) &= \mathcal{N}(\mathbf{M}^{-1} \mathbf{W}^\top \mathbf{x}_n, \sigma^2 \mathbf{M}^{-1}) \quad \text{where } \mathbf{M} = \mathbf{W}^\top \mathbf{W} + \sigma^2 \mathbf{I}_K \\ \mathbb{E}[\mathbf{z}_n] &= \mathbf{M}^{-1} \mathbf{W}^\top \mathbf{x}_n \end{aligned}$$

Parameter Estimation for PPCA via EM

- The expected complete data log-likelihood $\mathbb{E}[\log p(\mathbf{X}, \mathbf{Z}|\mathbf{W}, \sigma^2)]$

$$= - \sum_{n=1}^N \left\{ \frac{D}{2} \log \sigma^2 + \frac{1}{2\sigma^2} \|\mathbf{x}_n\|^2 - \frac{1}{\sigma^2} \mathbb{E}[\mathbf{z}_n]^\top \mathbf{W}^\top \mathbf{x}_n + \frac{1}{2\sigma^2} \text{tr}(\mathbb{E}[\mathbf{z}_n \mathbf{z}_n^\top] \mathbf{W}^\top \mathbf{W}) + \frac{1}{2} \text{tr}(\mathbb{E}[\mathbf{z}_n \mathbf{z}_n^\top]) \right\}$$

- Taking the derivative of $\mathbb{E}[\log p(\mathbf{X}, \mathbf{Z}|\mathbf{W}, \sigma^2)]$ w.r.t. \mathbf{W} and setting to zero

$$\mathbf{W} = \left[\sum_{n=1}^N \mathbf{x}_n \mathbb{E}[\mathbf{z}_n]^\top \right] \left[\sum_{n=1}^N \mathbb{E}[\mathbf{z}_n \mathbf{z}_n^\top] \right]^{-1} \quad (\text{Exercise: verify; can also be done "online"})$$

- To compute \mathbf{W} , we need two posterior expectations $\mathbb{E}[\mathbf{z}_n]$ and $\mathbb{E}[\mathbf{z}_n \mathbf{z}_n^\top]$
- These can be easily obtained from the posterior $p(\mathbf{z}_n|\mathbf{x}_n)$ computed in E step

$$\begin{aligned} p(\mathbf{z}_n|\mathbf{x}_n, \mathbf{W}) &= \mathcal{N}(\mathbf{M}^{-1} \mathbf{W}^\top \mathbf{x}_n, \sigma^2 \mathbf{M}^{-1}) \quad \text{where } \mathbf{M} = \mathbf{W}^\top \mathbf{W} + \sigma^2 \mathbf{I}_K \\ \mathbb{E}[\mathbf{z}_n] &= \mathbf{M}^{-1} \mathbf{W}^\top \mathbf{x}_n \\ \mathbb{E}[\mathbf{z}_n \mathbf{z}_n^\top] &= \mathbb{E}[\mathbf{z}_n] \mathbb{E}[\mathbf{z}_n]^\top + \text{cov}(\mathbf{z}_n) \end{aligned}$$

Parameter Estimation for PPCA via EM

- The expected complete data log-likelihood $\mathbb{E}[\log p(\mathbf{X}, \mathbf{Z}|\mathbf{W}, \sigma^2)]$

$$= - \sum_{n=1}^N \left\{ \frac{D}{2} \log \sigma^2 + \frac{1}{2\sigma^2} \|\mathbf{x}_n\|^2 - \frac{1}{\sigma^2} \mathbb{E}[\mathbf{z}_n]^\top \mathbf{W}^\top \mathbf{x}_n + \frac{1}{2\sigma^2} \text{tr}(\mathbb{E}[\mathbf{z}_n \mathbf{z}_n^\top] \mathbf{W}^\top \mathbf{W}) + \frac{1}{2} \text{tr}(\mathbb{E}[\mathbf{z}_n \mathbf{z}_n^\top]) \right\}$$

- Taking the derivative of $\mathbb{E}[\log p(\mathbf{X}, \mathbf{Z}|\mathbf{W}, \sigma^2)]$ w.r.t. \mathbf{W} and setting to zero

$$\mathbf{W} = \left[\sum_{n=1}^N \mathbf{x}_n \mathbb{E}[\mathbf{z}_n]^\top \right] \left[\sum_{n=1}^N \mathbb{E}[\mathbf{z}_n \mathbf{z}_n^\top] \right]^{-1} \quad (\text{Exercise: verify; can also be done "online"})$$

- To compute \mathbf{W} , we need two posterior expectations $\mathbb{E}[\mathbf{z}_n]$ and $\mathbb{E}[\mathbf{z}_n \mathbf{z}_n^\top]$
- These can be easily obtained from the posterior $p(\mathbf{z}_n|\mathbf{x}_n)$ computed in E step

$$\begin{aligned} p(\mathbf{z}_n|\mathbf{x}_n, \mathbf{W}) &= \mathcal{N}(\mathbf{M}^{-1} \mathbf{W}^\top \mathbf{x}_n, \sigma^2 \mathbf{M}^{-1}) \quad \text{where } \mathbf{M} = \mathbf{W}^\top \mathbf{W} + \sigma^2 \mathbf{I}_K \\ \mathbb{E}[\mathbf{z}_n] &= \mathbf{M}^{-1} \mathbf{W}^\top \mathbf{x}_n \\ \mathbb{E}[\mathbf{z}_n \mathbf{z}_n^\top] &= \mathbb{E}[\mathbf{z}_n] \mathbb{E}[\mathbf{z}_n]^\top + \text{cov}(\mathbf{z}_n) = \mathbb{E}[\mathbf{z}_n] \mathbb{E}[\mathbf{z}_n]^\top + \sigma^2 \mathbf{M}^{-1} \end{aligned}$$

Parameter Estimation for PPCA via EM

- The expected complete data log-likelihood $\mathbb{E}[\log p(\mathbf{X}, \mathbf{Z}|\mathbf{W}, \sigma^2)]$

$$= - \sum_{n=1}^N \left\{ \frac{D}{2} \log \sigma^2 + \frac{1}{2\sigma^2} \|\mathbf{x}_n\|^2 - \frac{1}{\sigma^2} \mathbb{E}[\mathbf{z}_n]^\top \mathbf{W}^\top \mathbf{x}_n + \frac{1}{2\sigma^2} \text{tr}(\mathbb{E}[\mathbf{z}_n \mathbf{z}_n^\top] \mathbf{W}^\top \mathbf{W}) + \frac{1}{2} \text{tr}(\mathbb{E}[\mathbf{z}_n \mathbf{z}_n^\top]) \right\}$$

- Taking the derivative of $\mathbb{E}[\log p(\mathbf{X}, \mathbf{Z}|\mathbf{W}, \sigma^2)]$ w.r.t. \mathbf{W} and setting to zero

$$\mathbf{W} = \left[\sum_{n=1}^N \mathbf{x}_n \mathbb{E}[\mathbf{z}_n]^\top \right] \left[\sum_{n=1}^N \mathbb{E}[\mathbf{z}_n \mathbf{z}_n^\top] \right]^{-1} \quad (\text{Exercise: verify; can also be done "online"})$$

- To compute \mathbf{W} , we need two posterior expectations $\mathbb{E}[\mathbf{z}_n]$ and $\mathbb{E}[\mathbf{z}_n \mathbf{z}_n^\top]$
- These can be easily obtained from the posterior $p(\mathbf{z}_n|\mathbf{x}_n)$ computed in E step

$$\begin{aligned} p(\mathbf{z}_n|\mathbf{x}_n, \mathbf{W}) &= \mathcal{N}(\mathbf{M}^{-1} \mathbf{W}^\top \mathbf{x}_n, \sigma^2 \mathbf{M}^{-1}) \quad \text{where } \mathbf{M} = \mathbf{W}^\top \mathbf{W} + \sigma^2 \mathbf{I}_K \\ \mathbb{E}[\mathbf{z}_n] &= \mathbf{M}^{-1} \mathbf{W}^\top \mathbf{x}_n \\ \mathbb{E}[\mathbf{z}_n \mathbf{z}_n^\top] &= \mathbb{E}[\mathbf{z}_n] \mathbb{E}[\mathbf{z}_n]^\top + \text{cov}(\mathbf{z}_n) = \mathbb{E}[\mathbf{z}_n] \mathbb{E}[\mathbf{z}_n]^\top + \sigma^2 \mathbf{M}^{-1} \end{aligned}$$

- Note: The noise variance σ^2 can also be estimated (take deriv., set to zero..)

Summary: The Full EM Algorithm for PPCA

- Specify K , initialize \mathbf{W} and σ^2 randomly. Also center the data ($\mathbf{x}_n = \mathbf{x}_n - \frac{1}{N} \sum_{n=1}^N \mathbf{x}_n$)

Summary: The Full EM Algorithm for PPCA

- Specify K , initialize \mathbf{W} and σ^2 randomly. Also center the data ($\mathbf{x}_n = \mathbf{x}_n - \frac{1}{N} \sum_{n=1}^N \mathbf{x}_n$)
- **E step:** For each n , compute $p(\mathbf{z}_n|\mathbf{x}_n)$ using current \mathbf{W} and σ^2 . Compute exp. for the M step

Summary: The Full EM Algorithm for PPCA

- Specify K , initialize \mathbf{W} and σ^2 randomly. Also center the data ($\mathbf{x}_n = \mathbf{x}_n - \frac{1}{N} \sum_{n=1}^N \mathbf{x}_n$)
- **E step:** For each n , compute $p(\mathbf{z}_n|\mathbf{x}_n)$ using current \mathbf{W} and σ^2 . Compute exp. for the M step

$$\begin{aligned} p(\mathbf{z}_n|\mathbf{x}_n, \mathbf{W}) &= \mathcal{N}(\mathbf{M}^{-1}\mathbf{W}^\top \mathbf{x}_n, \sigma^2 \mathbf{M}^{-1}) \quad \text{where } \mathbf{M} = \mathbf{W}^\top \mathbf{W} + \sigma^2 \mathbf{I}_K \\ \mathbb{E}[\mathbf{z}_n] &= \mathbf{M}^{-1}\mathbf{W}^\top \mathbf{x}_n \end{aligned}$$

Summary: The Full EM Algorithm for PPCA

- Specify K , initialize \mathbf{W} and σ^2 randomly. Also center the data ($\mathbf{x}_n = \mathbf{x}_n - \frac{1}{N} \sum_{n=1}^N \mathbf{x}_n$)
- E step:** For each n , compute $p(\mathbf{z}_n|\mathbf{x}_n)$ using current \mathbf{W} and σ^2 . Compute exp. for the M step

$$p(\mathbf{z}_n|\mathbf{x}_n, \mathbf{W}) = \mathcal{N}(\mathbf{M}^{-1}\mathbf{W}^\top \mathbf{x}_n, \sigma^2 \mathbf{M}^{-1}) \quad \text{where } \mathbf{M} = \mathbf{W}^\top \mathbf{W} + \sigma^2 \mathbf{I}_K$$

$$\mathbb{E}[\mathbf{z}_n] = \mathbf{M}^{-1}\mathbf{W}^\top \mathbf{x}_n$$

$$\mathbb{E}[\mathbf{z}_n \mathbf{z}_n^\top] = \text{cov}(\mathbf{z}_n) + \mathbb{E}[\mathbf{z}_n]\mathbb{E}[\mathbf{z}_n]^\top = \mathbb{E}[\mathbf{z}_n]\mathbb{E}[\mathbf{z}_n]^\top + \sigma^2 \mathbf{M}^{-1}$$

Summary: The Full EM Algorithm for PPCA

- Specify K , initialize \mathbf{W} and σ^2 randomly. Also center the data ($\mathbf{x}_n = \mathbf{x}_n - \frac{1}{N} \sum_{n=1}^N \mathbf{x}_n$)
- E step:** For each n , compute $p(\mathbf{z}_n|\mathbf{x}_n)$ using current \mathbf{W} and σ^2 . Compute exp. for the M step

$$p(\mathbf{z}_n|\mathbf{x}_n, \mathbf{W}) = \mathcal{N}(\mathbf{M}^{-1}\mathbf{W}^\top \mathbf{x}_n, \sigma^2 \mathbf{M}^{-1}) \quad \text{where } \mathbf{M} = \mathbf{W}^\top \mathbf{W} + \sigma^2 \mathbf{I}_K$$

$$\mathbb{E}[\mathbf{z}_n] = \mathbf{M}^{-1}\mathbf{W}^\top \mathbf{x}_n$$

$$\mathbb{E}[\mathbf{z}_n \mathbf{z}_n^\top] = \text{cov}(\mathbf{z}_n) + \mathbb{E}[\mathbf{z}_n]\mathbb{E}[\mathbf{z}_n]^\top = \mathbb{E}[\mathbf{z}_n]\mathbb{E}[\mathbf{z}_n]^\top + \sigma^2 \mathbf{M}^{-1}$$

- M step:** Re-estimate \mathbf{W} and σ^2

Summary: The Full EM Algorithm for PPCA

- Specify K , initialize \mathbf{W} and σ^2 randomly. Also center the data ($\mathbf{x}_n = \mathbf{x}_n - \frac{1}{N} \sum_{n=1}^N \mathbf{x}_n$)
- E step:** For each n , compute $p(\mathbf{z}_n|\mathbf{x}_n)$ using current \mathbf{W} and σ^2 . Compute exp. for the M step

$$p(\mathbf{z}_n|\mathbf{x}_n, \mathbf{W}) = \mathcal{N}(\mathbf{M}^{-1}\mathbf{W}^\top \mathbf{x}_n, \sigma^2 \mathbf{M}^{-1}) \quad \text{where } \mathbf{M} = \mathbf{W}^\top \mathbf{W} + \sigma^2 \mathbf{I}_K$$

$$\mathbb{E}[\mathbf{z}_n] = \mathbf{M}^{-1}\mathbf{W}^\top \mathbf{x}_n$$

$$\mathbb{E}[\mathbf{z}_n \mathbf{z}_n^\top] = \text{cov}(\mathbf{z}_n) + \mathbb{E}[\mathbf{z}_n]\mathbb{E}[\mathbf{z}_n]^\top = \mathbb{E}[\mathbf{z}_n]\mathbb{E}[\mathbf{z}_n]^\top + \sigma^2 \mathbf{M}^{-1}$$

- M step:** Re-estimate \mathbf{W} and σ^2

$$\mathbf{W}_{new} = \left[\sum_{n=1}^N \mathbf{x}_n \mathbb{E}[\mathbf{z}_n]^\top \right] \left[\sum_{n=1}^N \mathbb{E}[\mathbf{z}_n \mathbf{z}_n^\top] \right]^{-1}$$

Summary: The Full EM Algorithm for PPCA

- Specify K , initialize \mathbf{W} and σ^2 randomly. Also center the data ($\mathbf{x}_n = \mathbf{x}_n - \frac{1}{N} \sum_{n=1}^N \mathbf{x}_n$)
- E step:** For each n , compute $p(\mathbf{z}_n|\mathbf{x}_n)$ using current \mathbf{W} and σ^2 . Compute exp. for the M step

$$p(\mathbf{z}_n|\mathbf{x}_n, \mathbf{W}) = \mathcal{N}(\mathbf{M}^{-1}\mathbf{W}^\top \mathbf{x}_n, \sigma^2 \mathbf{M}^{-1}) \quad \text{where } \mathbf{M} = \mathbf{W}^\top \mathbf{W} + \sigma^2 \mathbf{I}_K$$

$$\mathbb{E}[\mathbf{z}_n] = \mathbf{M}^{-1}\mathbf{W}^\top \mathbf{x}_n$$

$$\mathbb{E}[\mathbf{z}_n \mathbf{z}_n^\top] = \text{cov}(\mathbf{z}_n) + \mathbb{E}[\mathbf{z}_n]\mathbb{E}[\mathbf{z}_n]^\top = \mathbb{E}[\mathbf{z}_n]\mathbb{E}[\mathbf{z}_n]^\top + \sigma^2 \mathbf{M}^{-1}$$

- M step:** Re-estimate \mathbf{W} and σ^2

$$\mathbf{W}_{new} = \left[\sum_{n=1}^N \mathbf{x}_n \mathbb{E}[\mathbf{z}_n]^\top \right] \left[\sum_{n=1}^N \mathbb{E}[\mathbf{z}_n \mathbf{z}_n^\top] \right]^{-1}$$

$$\sigma_{new}^2 = \frac{1}{ND} \sum_{n=1}^N \left\{ \|\mathbf{x}_n\|^2 - 2\mathbb{E}[\mathbf{z}_n]^\top \mathbf{W}_{new}^\top \mathbf{x}_n + \text{tr} \left(\mathbb{E}[\mathbf{z}_n \mathbf{z}_n^\top] \mathbf{W}_{new}^\top \mathbf{W}_{new} \right) \right\}$$

Summary: The Full EM Algorithm for PPCA

- Specify K , initialize \mathbf{W} and σ^2 randomly. Also center the data ($\mathbf{x}_n = \mathbf{x}_n - \frac{1}{N} \sum_{n=1}^N \mathbf{x}_n$)
- E step:** For each n , compute $p(\mathbf{z}_n|\mathbf{x}_n)$ using current \mathbf{W} and σ^2 . Compute exp. for the M step

$$p(\mathbf{z}_n|\mathbf{x}_n, \mathbf{W}) = \mathcal{N}(\mathbf{M}^{-1}\mathbf{W}^\top \mathbf{x}_n, \sigma^2 \mathbf{M}^{-1}) \quad \text{where } \mathbf{M} = \mathbf{W}^\top \mathbf{W} + \sigma^2 \mathbf{I}_K$$

$$\mathbb{E}[\mathbf{z}_n] = \mathbf{M}^{-1}\mathbf{W}^\top \mathbf{x}_n$$

$$\mathbb{E}[\mathbf{z}_n \mathbf{z}_n^\top] = \text{cov}(\mathbf{z}_n) + \mathbb{E}[\mathbf{z}_n]\mathbb{E}[\mathbf{z}_n]^\top = \mathbb{E}[\mathbf{z}_n]\mathbb{E}[\mathbf{z}_n]^\top + \sigma^2 \mathbf{M}^{-1}$$

- M step:** Re-estimate \mathbf{W} and σ^2

$$\mathbf{W}_{new} = \left[\sum_{n=1}^N \mathbf{x}_n \mathbb{E}[\mathbf{z}_n]^\top \right] \left[\sum_{n=1}^N \mathbb{E}[\mathbf{z}_n \mathbf{z}_n^\top] \right]^{-1}$$

$$\sigma_{new}^2 = \frac{1}{ND} \sum_{n=1}^N \left\{ \|\mathbf{x}_n\|^2 - 2\mathbb{E}[\mathbf{z}_n]^\top \mathbf{W}_{new}^\top \mathbf{x}_n + \text{tr} \left(\mathbb{E}[\mathbf{z}_n \mathbf{z}_n^\top] \mathbf{W}_{new}^\top \mathbf{W}_{new} \right) \right\}$$

- Set $\mathbf{W} = \mathbf{W}_{new}$ and $\sigma^2 = \sigma_{new}^2$. If not converged (monitor $p(\mathbf{X}|\Theta)$), go back to E step

Summary: The Full EM Algorithm for PPCA

- Specify K , initialize \mathbf{W} and σ^2 randomly. Also center the data ($\mathbf{x}_n = \mathbf{x}_n - \frac{1}{N} \sum_{n=1}^N \mathbf{x}_n$)
- E step:** For each n , compute $p(\mathbf{z}_n|\mathbf{x}_n)$ using current \mathbf{W} and σ^2 . Compute exp. for the M step

$$p(\mathbf{z}_n|\mathbf{x}_n, \mathbf{W}) = \mathcal{N}(\mathbf{M}^{-1}\mathbf{W}^\top \mathbf{x}_n, \sigma^2 \mathbf{M}^{-1}) \quad \text{where } \mathbf{M} = \mathbf{W}^\top \mathbf{W} + \sigma^2 \mathbf{I}_K$$

$$\mathbb{E}[\mathbf{z}_n] = \mathbf{M}^{-1}\mathbf{W}^\top \mathbf{x}_n$$

$$\mathbb{E}[\mathbf{z}_n \mathbf{z}_n^\top] = \text{cov}(\mathbf{z}_n) + \mathbb{E}[\mathbf{z}_n]\mathbb{E}[\mathbf{z}_n]^\top = \mathbb{E}[\mathbf{z}_n]\mathbb{E}[\mathbf{z}_n]^\top + \sigma^2 \mathbf{M}^{-1}$$

- M step:** Re-estimate \mathbf{W} and σ^2

$$\mathbf{W}_{new} = \left[\sum_{n=1}^N \mathbf{x}_n \mathbb{E}[\mathbf{z}_n]^\top \right] \left[\sum_{n=1}^N \mathbb{E}[\mathbf{z}_n \mathbf{z}_n^\top] \right]^{-1}$$

$$\sigma_{new}^2 = \frac{1}{ND} \sum_{n=1}^N \left\{ \|\mathbf{x}_n\|^2 - 2\mathbb{E}[\mathbf{z}_n]^\top \mathbf{W}_{new}^\top \mathbf{x}_n + \text{tr} \left(\mathbb{E}[\mathbf{z}_n \mathbf{z}_n^\top] \mathbf{W}_{new}^\top \mathbf{W}_{new} \right) \right\}$$

- Set $\mathbf{W} = \mathbf{W}_{new}$ and $\sigma^2 = \sigma_{new}^2$. If not converged (monitor $p(\mathbf{X}|\Theta)$), go back to E step
- Note:** This EM algorithm can easily be made online (Exercise: Think how!)

Summary: The Full EM Algorithm for PPCA

- Specify K , initialize \mathbf{W} and σ^2 randomly. Also center the data ($\mathbf{x}_n = \mathbf{x}_n - \frac{1}{N} \sum_{n=1}^N \mathbf{x}_n$)
- E step:** For each n , compute $p(\mathbf{z}_n|\mathbf{x}_n)$ using current \mathbf{W} and σ^2 . Compute exp. for the M step

$$p(\mathbf{z}_n|\mathbf{x}_n, \mathbf{W}) = \mathcal{N}(\mathbf{M}^{-1}\mathbf{W}^\top \mathbf{x}_n, \sigma^2 \mathbf{M}^{-1}) \quad \text{where } \mathbf{M} = \mathbf{W}^\top \mathbf{W} + \sigma^2 \mathbf{I}_K$$

$$\mathbb{E}[\mathbf{z}_n] = \mathbf{M}^{-1}\mathbf{W}^\top \mathbf{x}_n$$

$$\mathbb{E}[\mathbf{z}_n \mathbf{z}_n^\top] = \text{cov}(\mathbf{z}_n) + \mathbb{E}[\mathbf{z}_n]\mathbb{E}[\mathbf{z}_n]^\top = \mathbb{E}[\mathbf{z}_n]\mathbb{E}[\mathbf{z}_n]^\top + \sigma^2 \mathbf{M}^{-1}$$

- M step:** Re-estimate \mathbf{W} and σ^2

$$\mathbf{W}_{new} = \left[\sum_{n=1}^N \mathbf{x}_n \mathbb{E}[\mathbf{z}_n]^\top \right] \left[\sum_{n=1}^N \mathbb{E}[\mathbf{z}_n \mathbf{z}_n^\top] \right]^{-1}$$

$$\sigma_{new}^2 = \frac{1}{ND} \sum_{n=1}^N \left\{ \|\mathbf{x}_n\|^2 - 2\mathbb{E}[\mathbf{z}_n]^\top \mathbf{W}_{new}^\top \mathbf{x}_n + \text{tr} \left(\mathbb{E}[\mathbf{z}_n \mathbf{z}_n^\top] \mathbf{W}_{new}^\top \mathbf{W}_{new} \right) \right\}$$

- Set $\mathbf{W} = \mathbf{W}_{new}$ and $\sigma^2 = \sigma_{new}^2$. If not converged (monitor $p(\mathbf{X}|\Theta)$), go back to E step
- Note:** This EM algorithm can easily be made online (Exercise: Think how!)
- Note:** For $\sigma^2 = 0$, this EM algorithm can also be used to **efficiently** solve standard PCA

EM for Factor Analysis

- Similar to PPCA except that $\mathbf{x}_n \sim \mathcal{N}(\mathbf{W}\mathbf{z}_n, \mathbf{\Psi})$, where $\mathbf{\Psi}$ is a diagonal matrix

EM for Factor Analysis

- Similar to PPCA except that $\mathbf{x}_n \sim \mathcal{N}(\mathbf{W}\mathbf{z}_n, \mathbf{\Psi})$, where $\mathbf{\Psi}$ is a diagonal matrix
- EM for Factor Analysis is very similar to that for PPCA

EM for Factor Analysis

- Similar to PPCA except that $\mathbf{x}_n \sim \mathcal{N}(\mathbf{W}\mathbf{z}_n, \mathbf{\Psi})$, where $\mathbf{\Psi}$ is a diagonal matrix
- EM for Factor Analysis is very similar to that for PPCA
 - In the E step, for each n , compute the posterior over \mathbf{z}_n and the required expectations

$$p(\mathbf{z}_n | \mathbf{x}_n, \mathbf{W}, \mathbf{\Psi}) = \mathcal{N}(\mathbf{G}^{-1} \mathbf{W}^\top \mathbf{\Psi}^{-1} \mathbf{x}_n, \mathbf{G}) \quad (\text{Exercise: Verify})$$

$$\mathbb{E}[\mathbf{z}_n] = \mathbf{G}^{-1} \mathbf{W}^\top \mathbf{\Psi}^{-1} \mathbf{x}_n$$

$$\mathbb{E}[\mathbf{z}_n \mathbf{z}_n^\top] = \mathbb{E}[\mathbf{z}_n] \mathbb{E}[\mathbf{z}_n]^\top + \mathbf{G}^{-1}$$

where $\mathbf{G} = \mathbf{W}^\top \mathbf{\Psi}^{-1} \mathbf{W} + \mathbf{I}_K$.

EM for Factor Analysis

- Similar to PPCA except that $\mathbf{x}_n \sim \mathcal{N}(\mathbf{W}\mathbf{z}_n, \mathbf{\Psi})$, where $\mathbf{\Psi}$ is a diagonal matrix
- EM for Factor Analysis is very similar to that for PPCA
 - In the E step, for each n , compute the posterior over \mathbf{z}_n and the required expectations

$$p(\mathbf{z}_n | \mathbf{x}_n, \mathbf{W}, \mathbf{\Psi}) = \mathcal{N}(\mathbf{G}^{-1} \mathbf{W}^\top \mathbf{\Psi}^{-1} \mathbf{x}_n, \mathbf{G}) \quad (\text{Exercise: Verify})$$

$$\mathbb{E}[\mathbf{z}_n] = \mathbf{G}^{-1} \mathbf{W}^\top \mathbf{\Psi}^{-1} \mathbf{x}_n$$

$$\mathbb{E}[\mathbf{z}_n \mathbf{z}_n^\top] = \mathbb{E}[\mathbf{z}_n] \mathbb{E}[\mathbf{z}_n]^\top + \mathbf{G}^{-1}$$

where $\mathbf{G} = \mathbf{W}^\top \mathbf{\Psi}^{-1} \mathbf{W} + \mathbf{I}_K$. For $\mathbf{\Psi} = \sigma^2 \mathbf{I}_D$, we get the same equations as in PPCA

EM for Factor Analysis

- Similar to PPCA except that $\mathbf{x}_n \sim \mathcal{N}(\mathbf{W}\mathbf{z}_n, \mathbf{\Psi})$, where $\mathbf{\Psi}$ is a diagonal matrix
- EM for Factor Analysis is very similar to that for PPCA
 - In the E step, for each n , compute the posterior over \mathbf{z}_n and the required expectations

$$p(\mathbf{z}_n | \mathbf{x}_n, \mathbf{W}, \mathbf{\Psi}) = \mathcal{N}(\mathbf{G}^{-1} \mathbf{W}^\top \mathbf{\Psi}^{-1} \mathbf{x}_n, \mathbf{G}) \quad (\text{Exercise: Verify})$$

$$\mathbb{E}[\mathbf{z}_n] = \mathbf{G}^{-1} \mathbf{W}^\top \mathbf{\Psi}^{-1} \mathbf{x}_n$$

$$\mathbb{E}[\mathbf{z}_n \mathbf{z}_n^\top] = \mathbb{E}[\mathbf{z}_n] \mathbb{E}[\mathbf{z}_n]^\top + \mathbf{G}^{-1}$$

where $\mathbf{G} = \mathbf{W}^\top \mathbf{\Psi}^{-1} \mathbf{W} + \mathbf{I}_K$. For $\mathbf{\Psi} = \sigma^2 \mathbf{I}_D$, we get the same equations as in PPCA

- In the M step, updates for \mathbf{W}_{new} have the exact same for as PPCA

EM for Factor Analysis

- Similar to PPCA except that $\mathbf{x}_n \sim \mathcal{N}(\mathbf{W}\mathbf{z}_n, \mathbf{\Psi})$, where $\mathbf{\Psi}$ is a diagonal matrix
- EM for Factor Analysis is very similar to that for PPCA
 - In the E step, for each n , compute the posterior over \mathbf{z}_n and the required expectations

$$\begin{aligned}p(\mathbf{z}_n | \mathbf{x}_n, \mathbf{W}, \mathbf{\Psi}) &= \mathcal{N}(\mathbf{G}^{-1} \mathbf{W}^\top \mathbf{\Psi}^{-1} \mathbf{x}_n, \mathbf{G}) \quad (\text{Exercise: Verify}) \\ \mathbb{E}[\mathbf{z}_n] &= \mathbf{G}^{-1} \mathbf{W}^\top \mathbf{\Psi}^{-1} \mathbf{x}_n \\ \mathbb{E}[\mathbf{z}_n \mathbf{z}_n^\top] &= \mathbb{E}[\mathbf{z}_n] \mathbb{E}[\mathbf{z}_n]^\top + \mathbf{G}^{-1}\end{aligned}$$

where $\mathbf{G} = \mathbf{W}^\top \mathbf{\Psi}^{-1} \mathbf{W} + \mathbf{I}_K$. For $\mathbf{\Psi} = \sigma^2 \mathbf{I}_D$, we get the same equations as in PPCA

- In the M step, updates for \mathbf{W}_{new} have the exact same for as PPCA
- In the M step, updates for $\mathbf{\Psi}$ are

$$\mathbf{\Psi}_{new} = \text{diag} \left\{ \mathbf{S} - \mathbf{W}_{new} \frac{1}{N} \sum_{n=1}^N \mathbb{E}[\mathbf{z}_n] \mathbf{x}_n^\top \right\} \quad (\mathbf{S} \text{ is the cov. matrix of data})$$

EM for Factor Analysis

- Similar to PPCA except that $\mathbf{x}_n \sim \mathcal{N}(\mathbf{W}\mathbf{z}_n, \mathbf{\Psi})$, where $\mathbf{\Psi}$ is a diagonal matrix
- EM for Factor Analysis is very similar to that for PPCA
 - In the E step, for each n , compute the posterior over \mathbf{z}_n and the required expectations

$$p(\mathbf{z}_n | \mathbf{x}_n, \mathbf{W}, \mathbf{\Psi}) = \mathcal{N}(\mathbf{G}^{-1} \mathbf{W}^\top \mathbf{\Psi}^{-1} \mathbf{x}_n, \mathbf{G}) \quad (\text{Exercise: Verify})$$

$$\mathbb{E}[\mathbf{z}_n] = \mathbf{G}^{-1} \mathbf{W}^\top \mathbf{\Psi}^{-1} \mathbf{x}_n$$

$$\mathbb{E}[\mathbf{z}_n \mathbf{z}_n^\top] = \mathbb{E}[\mathbf{z}_n] \mathbb{E}[\mathbf{z}_n]^\top + \mathbf{G}^{-1}$$

where $\mathbf{G} = \mathbf{W}^\top \mathbf{\Psi}^{-1} \mathbf{W} + \mathbf{I}_K$. For $\mathbf{\Psi} = \sigma^2 \mathbf{I}_D$, we get the same equations as in PPCA

- In the M step, updates for \mathbf{W}_{new} have the exact same for as PPCA
- In the M step, updates for $\mathbf{\Psi}$ are

$$\mathbf{\Psi}_{new} = \text{diag} \left\{ \mathbf{S} - \mathbf{W}_{new} \frac{1}{N} \sum_{n=1}^N \mathbb{E}[\mathbf{z}_n] \mathbf{x}_n^\top \right\} \quad (\mathbf{S} \text{ is the cov. matrix of data})$$

- This EM algorithm too can be easily made online!

How to Set “K”?

- Several option to select the “best” K , e.g.,

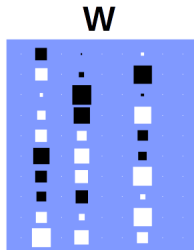
How to Set “K”?

- Several option to select the “best” K , e.g.,
 - Compute the **marginal likelihood** (or its approximation) for each K and choose the best model

How to Set “K”?

- Several options to select the “best” K , e.g.,
 - Compute the **marginal likelihood** (or its approximation) for each K and choose the best model
 - Use **sparsity inducing priors** on \mathbf{W} and/or \mathbf{z}_n (set K to some large value; the unnecessary columns of \mathbf{W} will “turn off” automatically as they will be shrunk to zero during inference)

Using sparsity-inducing Prior
(e.g., **Automatic Relevance
Determination**) on \mathbf{W}

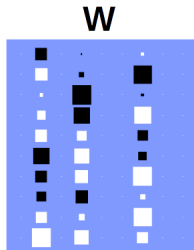


Effect: Only few columns
of \mathbf{W} will have entries with
significant magnitudes

How to Set “K”?

- Several options to select the “best” K , e.g.,
 - Compute the **marginal likelihood** (or its approximation) for each K and choose the best model
 - Use **sparsity inducing priors** on \mathbf{W} and/or \mathbf{z}_n (set K to some large value; the unnecessary columns of \mathbf{W} will “turn off” automatically as they will be shrunk to zero during inference)

Using sparsity-inducing Prior
(e.g., **Automatic Relevance
Determination**) on \mathbf{W}



Effect: Only few columns
of \mathbf{W} will have entries with
significant magnitudes

- **Nonparametric Bayesian methods** (allow K to grow with data)

Some Applications of PPCA/FA

- Compression/dimensionality reduction is a natural application (use \mathbf{z}_n instead of \mathbf{x}_n)

Some Applications of PPCA/FA

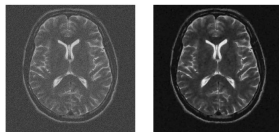
- Compression/dimensionality reduction is a natural application (use \mathbf{z}_n instead of \mathbf{x}_n)
- Also used for learning low-dim. “good” features \mathbf{z}_n from high-dim noisy features \mathbf{x}_n

Some Applications of PPCA/FA

- Compression/dimensionality reduction is a natural application (use \mathbf{z}_n instead of \mathbf{x}_n)
- Also used for learning low-dim. “good” features \mathbf{z}_n from high-dim noisy features \mathbf{x}_n
 - Note that this is different from feature selection (\mathbf{z}_n is a transformed version of \mathbf{x}_n , not a subset)

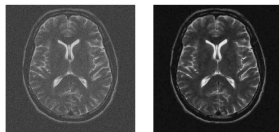
Some Applications of PPCA/FA

- Compression/dimensionality reduction is a natural application (use \mathbf{z}_n instead of \mathbf{x}_n)
- Also used for learning low-dim. “good” features \mathbf{z}_n from high-dim noisy features \mathbf{x}_n
 - Note that this is different from feature selection (\mathbf{z}_n is a transformed version of \mathbf{x}_n , not a subset)
- Learning the noise variance enables “image denoising”: $\mathbf{x}_n = \mathbf{W}\mathbf{z}_n + \epsilon_n$; $\mathbf{W}\mathbf{z}_n$ is the “clean” part



Some Applications of PPCA/FA

- Compression/dimensionality reduction is a natural application (use \mathbf{z}_n instead of \mathbf{x}_n)
- Also used for learning low-dim. “good” features \mathbf{z}_n from high-dim noisy features \mathbf{x}_n
 - Note that this is different from feature selection (\mathbf{z}_n is a transformed version of \mathbf{x}_n , not a subset)
- Learning the noise variance enables “image denoising”: $\mathbf{x}_n = \mathbf{W}\mathbf{z}_n + \epsilon_n$; $\mathbf{W}\mathbf{z}_n$ is the “clean” part



- Ability to fill-in missing data enables “image inpainting” (left: image with 80% missing data, middle: reconstructed, right: original)



Mixture of PPCA and Mixture of FA

- May be appropriate if data also exists in clusters (suppose $M > 1$ clusters)

Mixture of PPCA and Mixture of FA

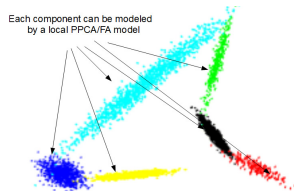
- May be appropriate if data also exists in clusters (suppose $M > 1$ clusters)
- Data in each cluster (say m) can have its own “local” PPCA/FA model defined by $\{\mathbf{W}_m, \sigma_m^2\}$

Mixture of PPCA and Mixture of FA

- May be appropriate if data also exists in clusters (suppose $M > 1$ clusters)
- Data in each cluster (say m) can have its own “local” PPCA/FA model defined by $\{\mathbf{W}_m, \sigma_m^2\}$
- Can use M such PPCA/FA models $\{\mathbf{W}_m, \sigma_m^2\}_{m=1}^M$ (one per cluster) for the entire data

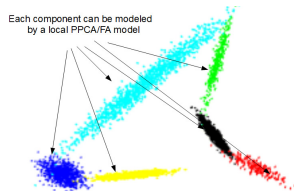
Mixture of PPCA and Mixture of FA

- May be appropriate if data also exists in clusters (suppose $M > 1$ clusters)
- Data in each cluster (say m) can have its own “local” PPCA/FA model defined by $\{\mathbf{W}_m, \sigma_m^2\}$
- Can use M such PPCA/FA models $\{\mathbf{W}_m, \sigma_m^2\}_{m=1}^M$ (one per cluster) for the entire data



Mixture of PPCA and Mixture of FA

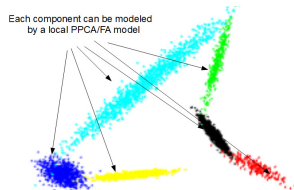
- May be appropriate if data also exists in clusters (suppose $M > 1$ clusters)
- Data in each cluster (say m) can have its own “local” PPCA/FA model defined by $\{\mathbf{W}_m, \sigma_m^2\}$
- Can use M such PPCA/FA models $\{\mathbf{W}_m, \sigma_m^2\}_{m=1}^M$ (one per cluster) for the entire data



- Mixtures of PPCA/FA can be seen as playing several roles

Mixture of PPCA and Mixture of FA

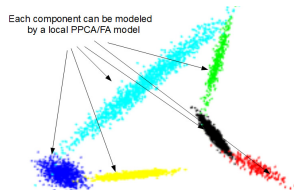
- May be appropriate if data also exists in clusters (suppose $M > 1$ clusters)
- Data in each cluster (say m) can have its own “local” PPCA/FA model defined by $\{\mathbf{W}_m, \sigma_m^2\}$
- Can use M such PPCA/FA models $\{\mathbf{W}_m, \sigma_m^2\}_{m=1}^M$ (one per cluster) for the entire data



- Mixtures of PPCA/FA can be seen as playing several roles
 - Jointly learning clustering and dimensionality reduction

Mixture of PPCA and Mixture of FA

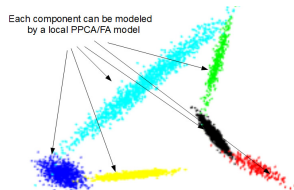
- May be appropriate if data also exists in clusters (suppose $M > 1$ clusters)
- Data in each cluster (say m) can have its own “local” PPCA/FA model defined by $\{\mathbf{W}_m, \sigma_m^2\}$
- Can use M such PPCA/FA models $\{\mathbf{W}_m, \sigma_m^2\}_{m=1}^M$ (one per cluster) for the entire data



- Mixtures of PPCA/FA can be seen as playing several roles
 - Jointly learning clustering and dimensionality reduction
 - Nonlinear dimensionality reduction

Mixture of PPCA and Mixture of FA

- May be appropriate if data also exists in clusters (suppose $M > 1$ clusters)
- Data in each cluster (say m) can have its own “local” PPCA/FA model defined by $\{\mathbf{W}_m, \sigma_m^2\}$
- Can use M such PPCA/FA models $\{\mathbf{W}_m, \sigma_m^2\}_{m=1}^M$ (one per cluster) for the entire data



- Mixtures of PPCA/FA can be seen as playing several roles
 - Jointly learning clustering and dimensionality reduction
 - Nonlinear dimensionality reduction
 - A flexible probability density model: Mixture of **low-rank** Gaussians

Mixture of PPCA and Mixture of FA

- For mixture of PPCA/FA, the generative story for each observation \mathbf{x}_n is as follows

Mixture of PPCA and Mixture of FA

- For mixture of PPCA/FA, the generative story for each observation \mathbf{x}_n is as follows
 - Generate its cluster id as

$$\mathbf{c}_n \sim \text{multinoulli}(\pi_1, \dots, \pi_M)$$

Mixture of PPCA and Mixture of FA

- For mixture of PPCA/FA, the generative story for each observation \mathbf{x}_n is as follows

- Generate its cluster id as

$$\mathbf{c}_n \sim \text{multinoulli}(\pi_1, \dots, \pi_M)$$

- Generate latent variable $\mathbf{z}_n \in \mathbb{R}^K$ as

$$\mathbf{z}_n \sim \mathcal{N}(\mathbf{0}, \mathbf{I}_K)$$

Mixture of PPCA and Mixture of FA

- For mixture of PPCA/FA, the generative story for each observation \mathbf{x}_n is as follows

- Generate its cluster id as

$$\mathbf{c}_n \sim \text{multinoulli}(\pi_1, \dots, \pi_M)$$

- Generate latent variable $\mathbf{z}_n \in \mathbb{R}^K$ as

$$\mathbf{z}_n \sim \mathcal{N}(\mathbf{0}, \mathbf{I}_K)$$

- Generate observation $\mathbf{x}_n \in \mathbb{R}^D$ from the \mathbf{c}_n^{th} PPCA/FA model

$$\mathbf{x}_n \sim \mathcal{N}(\boldsymbol{\mu}_{\mathbf{c}_n} + \mathbf{W}_{\mathbf{c}_n} \mathbf{z}_n, \sigma_{\mathbf{c}_n}^2 \mathbf{I}_D)$$

Mixture of PPCA and Mixture of FA

- For mixture of PPCA/FA, the generative story for each observation \mathbf{x}_n is as follows

- Generate its cluster id as

$$\mathbf{c}_n \sim \text{multinoulli}(\pi_1, \dots, \pi_M)$$

- Generate latent variable $\mathbf{z}_n \in \mathbb{R}^K$ as

$$\mathbf{z}_n \sim \mathcal{N}(\mathbf{0}, \mathbf{I}_K)$$

- Generate observation $\mathbf{x}_n \in \mathbb{R}^D$ from the \mathbf{c}_n^{th} PPCA/FA model

$$\mathbf{x}_n \sim \mathcal{N}(\boldsymbol{\mu}_{\mathbf{c}_n} + \mathbf{W}_{\mathbf{c}_n} \mathbf{z}_n, \sigma_{\mathbf{c}_n}^2 \mathbf{I}_D)$$

- Each PPCA/FA model has its separate mean $\boldsymbol{\mu}_{\mathbf{c}_n}$ (not needed when $M = 1$ if data is centered)

Mixture of PPCA and Mixture of FA

- For mixture of PPCA/FA, the generative story for each observation \mathbf{x}_n is as follows

- Generate its cluster id as

$$\mathbf{c}_n \sim \text{multinoulli}(\pi_1, \dots, \pi_M)$$

- Generate latent variable $\mathbf{z}_n \in \mathbb{R}^K$ as

$$\mathbf{z}_n \sim \mathcal{N}(\mathbf{0}, \mathbf{I}_K)$$

- Generate observation $\mathbf{x}_n \in \mathbb{R}^D$ from the \mathbf{c}_n^{th} PPCA/FA model

$$\mathbf{x}_n \sim \mathcal{N}(\boldsymbol{\mu}_{\mathbf{c}_n} + \mathbf{W}_{\mathbf{c}_n} \mathbf{z}_n, \sigma_{\mathbf{c}_n}^2 \mathbf{I}_D)$$

- Each PPCA/FA model has its separate mean $\boldsymbol{\mu}_{\mathbf{c}_n}$ (not needed when $M = 1$ if data is centered)
- **Exercise:** What will be the **marginal distribution** of \mathbf{x}_n , i.e.. $p(\mathbf{x}_n|\Theta)$?

Mixture of PPCA and Mixture of FA

- For mixture of PPCA/FA, the generative story for each observation \mathbf{x}_n is as follows

- Generate its cluster id as

$$\mathbf{c}_n \sim \text{multinoulli}(\pi_1, \dots, \pi_M)$$

- Generate latent variable $\mathbf{z}_n \in \mathbb{R}^K$ as

$$\mathbf{z}_n \sim \mathcal{N}(\mathbf{0}, \mathbf{I}_K)$$

- Generate observation $\mathbf{x}_n \in \mathbb{R}^D$ from the \mathbf{c}_n^{th} PPCA/FA model

$$\mathbf{x}_n \sim \mathcal{N}(\boldsymbol{\mu}_{\mathbf{c}_n} + \mathbf{W}_{\mathbf{c}_n} \mathbf{z}_n, \sigma_{\mathbf{c}_n}^2 \mathbf{I}_D)$$

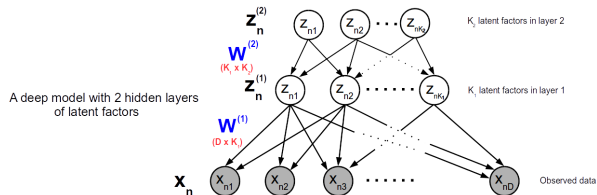
- Each PPCA/FA model has its separate mean $\boldsymbol{\mu}_{\mathbf{c}_n}$ (not needed when $M = 1$ if data is centered)
- **Exercise:** What will be the **marginal distribution** of \mathbf{x}_n , i.e.. $p(\mathbf{x}_n|\Theta)$?
- EM can be used in this model to learn the parameters and latent variables

PPCA/FA Extensions

- Already saw the mixture of PPCA/FA

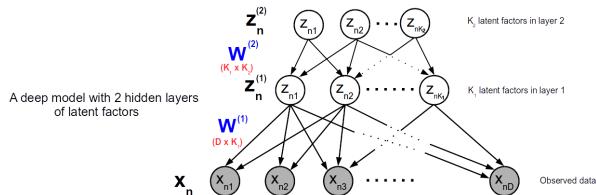
PPCA/FA Extensions

- Already saw the mixture of PPCA/FA
- Can stack PPCA/FA models to construct deep generative models (more on this later)

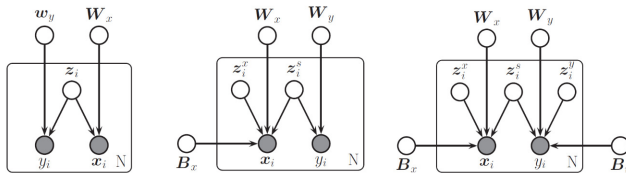


PPCA/FA Extensions

- Already saw the mixture of PPCA/FA
- Can stack PPCA/FA models to construct deep generative models (more on this later)



- Many other extensions for supervised dim-red, multi-modality data such as image+caption, etc.



Some Concluding Thoughts

- Latent variable models for dim-red have many advantages over non-probabilistic counterparts

Some Concluding Thoughts

- Latent variable models for dim-red have many advantages over non-probabilistic counterparts
 - Brings in all the benefits of probabilistic modeling with diverse likelihood/priors

Some Concluding Thoughts

- Latent variable models for dim-red have many advantages over non-probabilistic counterparts
 - Brings in all the benefits of probabilistic modeling with diverse likelihood/priors
 - Can be easily extended to learning more complex patterns (e.g., mixture of PPCA)

Some Concluding Thoughts

- Latent variable models for dim-red have many advantages over non-probabilistic counterparts
 - Brings in all the benefits of probabilistic modeling with diverse likelihood/priors
 - Can be easily extended to learning more complex patterns (e.g., mixture of PPCA)
 - Often more efficient too (e.g., using PPCA with EM can be faster than PCA with eigen-decomp.)

Some Concluding Thoughts

- Latent variable models for dim-red have many advantages over non-probabilistic counterparts
 - Brings in all the benefits of probabilistic modeling with diverse likelihood/priors
 - Can be easily extended to learning more complex patterns (e.g., mixture of PPCA)
 - Often more efficient too (e.g., using PPCA with EM can be faster than PCA with eigen-decomp.)
 - Handle missing data in a principled way

Some Concluding Thoughts

- Latent variable models for dim-red have many advantages over non-probabilistic counterparts
 - Brings in all the benefits of probabilistic modeling with diverse likelihood/priors
 - Can be easily extended to learning more complex patterns (e.g., mixture of PPCA)
 - Often more efficient too (e.g., using PPCA with EM can be faster than PCA with eigen-decomp.)
 - Handle missing data in a principled way
 - Nonlinear maps from z to x possible (e.g., using deep neural nets or Gaussian Processes)

Some Concluding Thoughts

- Latent variable models for dim-red have many advantages over non-probabilistic counterparts
 - Brings in all the benefits of probabilistic modeling with diverse likelihood/priors
 - Can be easily extended to learning more complex patterns (e.g., mixture of PPCA)
 - Often more efficient too (e.g., using PPCA with EM can be faster than PCA with eigen-decomp.)
 - Handle missing data in a principled way
 - Nonlinear maps from \mathbf{z} to \mathbf{x} possible (e.g., using deep neural nets or Gaussian Processes)
- These models can also be used for **estimating the probability density** $p(\mathbf{x})$

$$p(\mathbf{x}) = \int p(\mathbf{x}|\mathbf{z}, \theta)p(\mathbf{z}|\phi)d\mathbf{z}$$

Some Concluding Thoughts

- Latent variable models for dim-red have many advantages over non-probabilistic counterparts
 - Brings in all the benefits of probabilistic modeling with diverse likelihood/priors
 - Can be easily extended to learning more complex patterns (e.g., mixture of PPCA)
 - Often more efficient too (e.g., using PPCA with EM can be faster than PCA with eigen-decomp.)
 - Handle missing data in a principled way
 - Nonlinear maps from \mathbf{z} to \mathbf{x} possible (e.g., using deep neural nets or Gaussian Processes)
- These models can also be used for [estimating the probability density](#) $p(\mathbf{x})$

$$p(\mathbf{x}) = \int p(\mathbf{x}|\mathbf{z}, \theta)p(\mathbf{z}|\phi)d\mathbf{z}$$

- Can also be extended to model sequential data, e.g., [Linear Dynamical Systems](#) (will see later)

Some Concluding Thoughts

- Latent variable models for dim-red have many advantages over non-probabilistic counterparts
 - Brings in all the benefits of probabilistic modeling with diverse likelihood/priors
 - Can be easily extended to learning more complex patterns (e.g., mixture of PPCA)
 - Often more efficient too (e.g., using PPCA with EM can be faster than PCA with eigen-decomp.)
 - Handle missing data in a principled way
 - Nonlinear maps from \mathbf{z} to \mathbf{x} possible (e.g., using deep neural nets or Gaussian Processes)
- These models can also be used for **estimating the probability density** $p(\mathbf{x})$

$$p(\mathbf{x}) = \int p(\mathbf{x}|\mathbf{z}, \theta)p(\mathbf{z}|\phi)d\mathbf{z}$$

- Can also be extended to model sequential data, e.g., **Linear Dynamical Systems** (will see later)
- EM is only one of the ways to do inference in LVMs. Can also do fully Bayesian inference

Some Concluding Thoughts

- Latent variable models for dim-red have many advantages over non-probabilistic counterparts
 - Brings in all the benefits of probabilistic modeling with diverse likelihood/priors
 - Can be easily extended to learning more complex patterns (e.g., mixture of PPCA)
 - Often more efficient too (e.g., using PPCA with EM can be faster than PCA with eigen-decomp.)
 - Handle missing data in a principled way
 - Nonlinear maps from \mathbf{z} to \mathbf{x} possible (e.g., using deep neural nets or Gaussian Processes)
- These models can also be used for **estimating the probability density** $p(\mathbf{x})$

$$p(\mathbf{x}) = \int p(\mathbf{x}|\mathbf{z}, \theta)p(\mathbf{z}|\phi)d\mathbf{z}$$

- Can also be extended to model sequential data, e.g., **Linear Dynamical Systems** (will see later)
- EM is only one of the ways to do inference in LVMs. Can also do fully Bayesian inference
 - More on this when we discuss MCMC and Variational Inference