

Acknowledgements

The slides for this lecture are a modified versions of the offering by **Prof. Sanjeev K Aggarwal**

Lexical Analysis

- Recognize tokens and ignore white spaces, comments

i	f		(x	1		*	x	2	<	1	.	0)	{
---	---	--	---	---	---	--	---	---	---	---	---	---	---	---	---

Generates token stream


if	(x1	*	x2	<	1.0)	{
----	---	----	---	----	---	-----	---	---

- Error reporting
- Model using regular expressions
- Recognize using Finite State Automata

Lexical Analysis

- Sentences consist of string of tokens (a syntactic category) for example, number, identifier, keyword, string
- Sequences of characters in a token is a **lexeme** for example, 100.01, counter, const, “How are you?”
- Rule of description is a pattern for example, letter (letter | digit)*
- Discard whatever does not contribute to parsing like white spaces (blanks, tabs, newlines) and comments
- construct constants: convert numbers to token num and pass number as its attribute, for example, integer 31 becomes <num, 31>
- recognize keyword and identifiers for example counter = counter + increment becomes id = id + id /*check if id is a keyword*/

Interface to other phases



- Push back is required due to lookahead for example $> =$ and $>$
- It is implemented through a buffer
 - Keep input in a buffer
 - Move pointers over the input

Approaches to implementation

- Use assembly language
Most efficient but most difficult to implement
- Use high level languages like C
Efficient but difficult to implement
- Use tools like lex, flex
Easy to implement but not as efficient as the first two cases

Construct a lexical analyzer

- Allow white spaces, numbers and arithmetic operators in an expression
- Return tokens and attributes to the syntax analyzer
- A global variable **tokenval** is set to the value of the number
- Design requires that
 - A finite set of tokens be defined
 - Describe strings belonging to each token

```

#include <stdio.h>
#include <ctype.h>
int lineno = 1;
int tokenval = NONE;
int lex() {
    int t;
    while (1) {
        t = getchar ();
        if (t == ' ' || t == '\t');
            else if (t == '\n') lineno = lineno + 1;
            else if (isdigit (t) ) {
                tokenval = t - '0' ;
                t = getchar ();
                while (isdigit(t)) {
                    tokenval = tokenval * 10 + t - '0' ;
                    t = getchar();
                }
                ungetc(t,stdin);
                return num;
            }
        else { tokenval = NONE; return t; }
    }
}

```

Problems

- Scans text character by character
- Look ahead character determines what kind of token to read and when the current token ends
- First character cannot determine what kind of token we are going to read

Difficulties in design of lexical analyzers

- Is it as simple as it sounds?
- Lexemes in a fixed position. Fix format vs. free format languages
- Handling of blanks
 - in Pascal, blanks separate identifiers
 - in Fortran, blanks are important only in literal strings for example variable **counter** is same as **count er**
 - Another example

DO 10 I = 1.25	DO10I=1.25
DO 10 I = 1,25	DO10I=1,25

- The first line is a variable assignment
`DO10I=1.25`
- second line is beginning of a
`Do loop`
- Reading from left to right one can not distinguish between the two until the ";" or "." is reached
- Fortran white space and fixed format rules came into force due to punch cards and errors in punching

[illegible]

PL/1 Problems

- Keywords are not reserved in PL/1
if then then then = else else else = then
if if then then = then + 1
- PL/1 declarations
Declare(arg₁,arg₂,arg₃,.....,arg_n)
- Cannot tell whether Declare is a keyword or array reference until after ")"
- Requires arbitrary lookahead and very large buffers. Worse, the buffers may have to be reloaded.

Problem continues even today!!

- C++ template syntax: `Foo<Bar>`
- C++ stream syntax: `cin >> var;`
- Nested templates: `Foo<Bar<Bazz>>`
- Can these problems be resolved by lexical analyzers alone?

Building a lexer...

Need to answer the following questions:

- How to specify tokens (patterns)?
- How to recognize tokens (efficiently)?
 - Not just a language acceptance question...
- How to resolve conflicts?
 - Priority
 - Maximal munch

How to specify tokens?

- How to describe tokens

2.e0 20.e-01 2.000

- How to break text into token

if (x==0) a = x << 1;

iff (x==0) a = x < 1;

- How to break input into tokens efficiently
 - Tokens may have similar prefixes
 - Each character should be looked at only once

How to describe tokens?

- Programming language tokens can be described by regular languages
- Regular languages
 - Are easy to understand
 - There is a well understood and useful theory
 - They have efficient implementation
- Regular languages have been discussed in great detail in the "Theory of Computation" course

Notation

- Let Σ be a set of characters. A language over Σ is a set of strings of characters belonging to Σ
- A regular expression r denotes a language $L(r)$
- Rules that define the regular expressions over Σ
 - ϵ is a regular expression that denotes $\{\epsilon\}$ the set containing the empty string
 - If a is a symbol in Σ then a is a regular expression that denotes $\{a\}$

- If r and s are regular expressions denoting the languages $L(r)$ and $L(s)$ then
- $(r)|(s)$ is a regular expression denoting $L(r) \cup L(s)$
- $(r)(s)$ is a regular expression denoting $L(r)L(s)$
- $(r)^*$ is a regular expression denoting $(L(r))^*$
- (r) is a regular expression denoting $L(r)$

- Let $\Sigma = \{a, b\}$
- The regular expression $a|b$ denotes the set $\{a, b\}$
- The regular expression $(a|b)(a|b)$ denotes $\{aa, ab, ba, bb\}$
- The regular expression a^* denotes the set of all strings $\{\epsilon, a, aa, aaa, \dots\}$
- The regular expression $(a|b)^*$ denotes the set of all strings containing ϵ and all strings of a's and b's
- The regular expression $a|a^*b$ denotes the set containing the string a and all strings consisting of zero or more a's followed by a character b

How to specify tokens

- Regular definitions
 - Let r_i be a regular expression and d_i be a distinct name
 - Regular definition is a sequence of definitions of the form
$$\begin{array}{l} d_1 \rightarrow r_1 \\ d_2 \rightarrow r_2 \\ \dots \\ d_n \rightarrow r_n \end{array}$$
 - Where each r_i is a regular expression over $\Sigma \cup \{d_1, d_2, \dots, d_{i-1}\}$

Examples

- My fax number
91-(512)-259-7586
- $\Sigma = \text{digits} \cup \{-, (,)\}$
- Country $\rightarrow \text{digit}^+$ digit^2
- Area $\rightarrow '(' \text{digit}^+ ')'$ digit^3
- Exchange $\rightarrow \text{digit}^+$ digit^3
- Phone $\rightarrow \text{digit}^+$ digit^4
- Number $\rightarrow \text{country '-' area '-' exchange '-' phone}$

Examples ...

- Email address
abc@iitk.ac.in
- $\Sigma = \text{letter} \cup \{ @, . \}$
- Letter $\rightarrow a \mid b \mid \dots \mid z \mid A \mid B \mid \dots \mid Z$
- Name $\rightarrow \text{letter}^+$
- Address $\rightarrow \text{name '@' name '.' name '.' name}$

Examples ...

- Identifier

letter $\rightarrow a \mid b \mid \dots \mid z \mid A \mid B \mid \dots \mid Z$

digit $\rightarrow 0 \mid 1 \mid \dots \mid 9$

identifier $\rightarrow \text{letter}(\text{letter} \mid \text{digit})^*$

- Unsigned number in Pascal

digit $\rightarrow 0 \mid 1 \mid \dots \mid 9$

digits $\rightarrow \text{digit}^+$

fraction $\rightarrow \text{'.' digits} \mid \epsilon$

exponent $\rightarrow (\text{'E' ('+' | '-' | '\epsilon') digits}) \mid \epsilon$

number $\rightarrow \text{digits fraction exponent}$

Regular expressions in specifications

- Regular expressions describe many useful languages
- Regular expressions are only specifications; implementation is still required
- Given a string s and a regular expression R , does $s \in L(R)$?
- Solution to this problem is the basis of the lexical analyzers
- However, just the yes/no answer is not important
- Goal: Partition the input into tokens

1. Write a regular expression for lexemes of each token
 - number \rightarrow digit⁺
 - identifier \rightarrow letter(letter|digit)⁺
2. Construct R matching all lexemes of all tokens
 - $R = R1 + R2 + R3 + \dots$
3. Let input be $x_1 \dots x_n$
 - for $1 \leq i \leq n$ check $x_1 \dots x_i \in L(R)$
4. $x_1 \dots x_i \in L(R) \rightarrow x_1 \dots x_i \in L(R_j)$ for some j
 - smallest such j is token class of $x_1 \dots x_i$
5. Remove $x_1 \dots x_i$ from input; go to (3)

- The algorithm gives priority to tokens listed earlier
 - Treats "if" as keyword and not identifier
- How much input is used? What if
 - $x_1 \dots x_i \in L(R)$
 - $x_1 \dots x_j \in L(R)$
 - Pick up the longest possible string in $L(R)$
 - The principle of "maximal munch"
- Regular expressions provide a concise and useful notation for string patterns
- Good algorithms require a single pass over the input

How to break up text

- Elsex=0

else	x	=	0
------	---	---	---


elsex	=	0
-------	---	---
- Regular expressions alone are not enough
- Normally the longest match wins
- Ties are resolved by prioritizing tokens
- Lexical definitions consist of regular definitions, priority rules and maximal munch principle

Transition Diagrams


- Regular expression are declarative specifications
- Transition diagram is an implementation
- A transition diagram consists of
 - An input alphabet belonging to Σ
 - A set of states S
 - A set of transitions $\text{state}_i \xrightarrow{\text{input}} \text{state}_j$
 - A set of final states F
 - A start state n
- Transition $s1 \xrightarrow{a} s2$ is read:
in state $s1$ on input a go to state $s2$
- If end of input is reached in a final state then accept
- Otherwise, reject

Pictorial notation

- A state




- A final state



- Transition



- Transition from state i to state j on an input a



How to recognize tokens

- Consider

relop $\rightarrow < \mid <= \mid = \mid <> \mid >= \mid >$

id $\rightarrow \text{letter}(\text{letter} \mid \text{digit})^*$

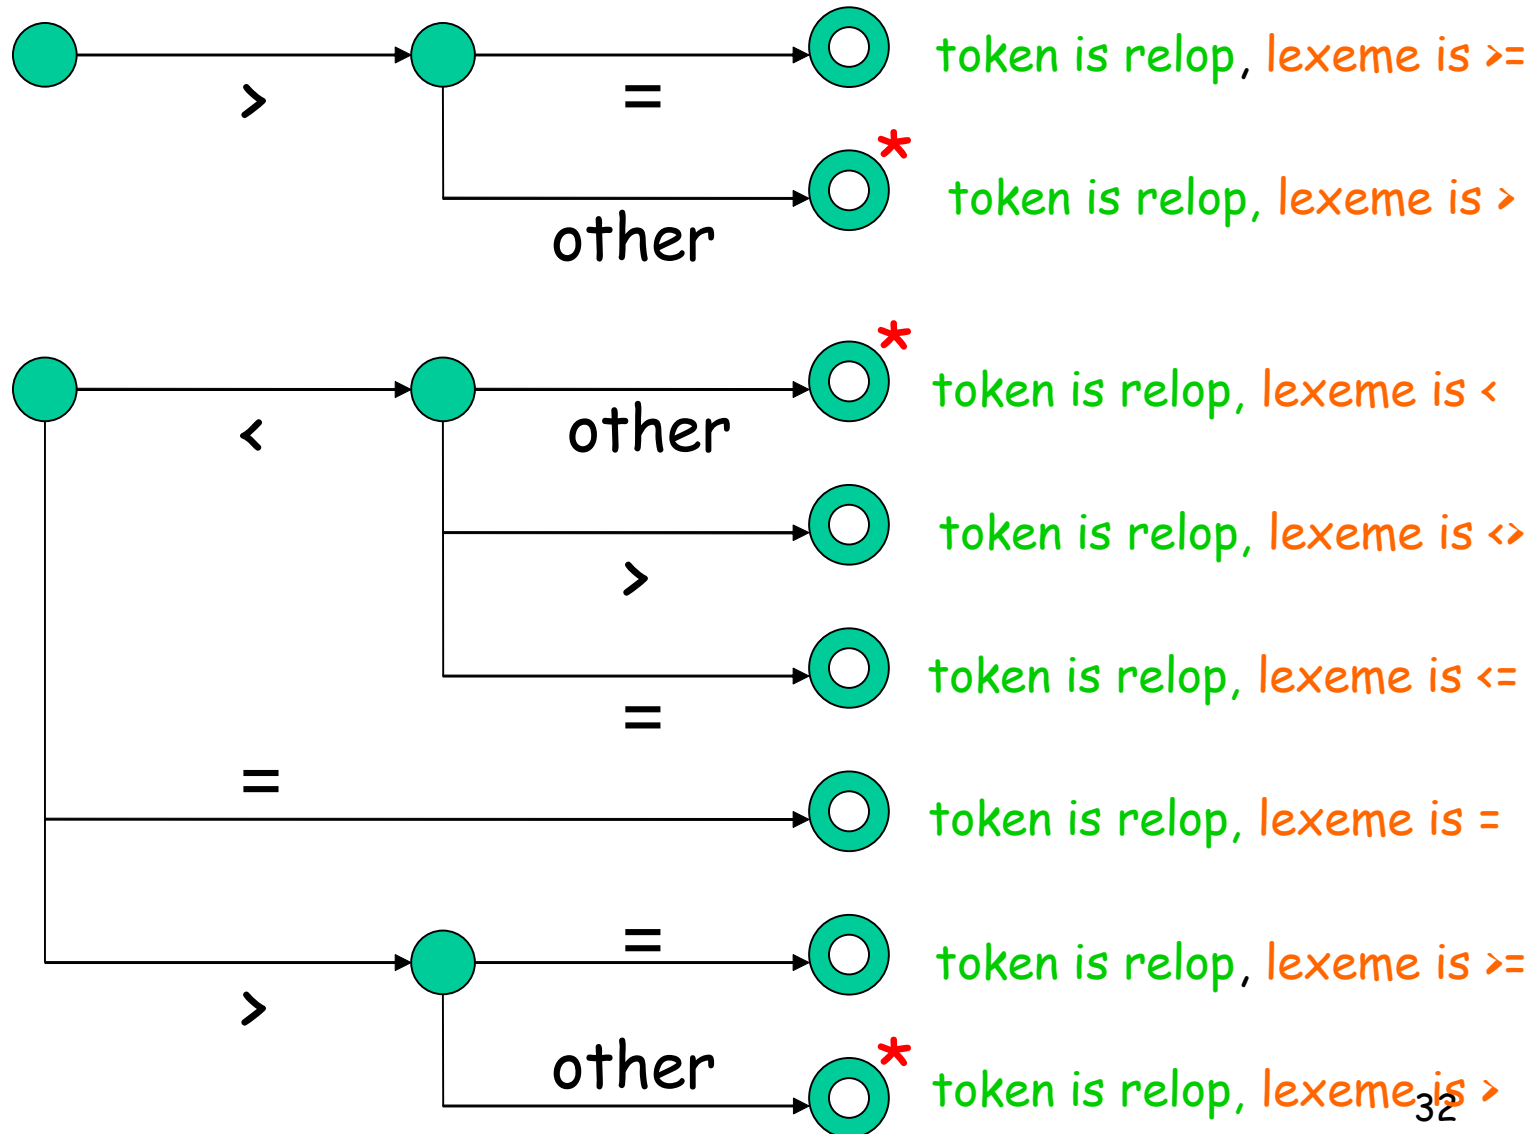
num $\rightarrow \text{digit}^+ (\text{'.' digit}^+)? (\text{E}(\text{'+'} \mid \text{'-'})? \text{digit}^+)?$

delim $\rightarrow \text{blank} \mid \text{tab} \mid \text{newline}$


ws $\rightarrow \text{delim}^+$

- Construct an analyzer that will return $\langle \text{token}, \text{attribute} \rangle$ pairs


Transition diagram for relops





Transition diagram for identifier




Transition diagram for white spaces



Transition diagram for unsigned numbers



Real numbers




Integer number

- The lexeme for a given token must be the longest possible
- Assume input to be 12.34E56
- Starting in the third diagram the accept state will be reached after 12
- Therefore, the matching should always start with the first transition diagram
- If failure occurs in one transition diagram then retract the forward pointer to the start state and activate the next diagram
- If failure occurs in all diagrams then a lexical error has occurred

Implementation of transition diagrams

```
Token nexttoken() {  
    while(1) {  
        switch (state) {  
            .....  
            case 10: c=nextchar();  
                    if(isletter(c)) state=10;  
                    elseif (isdigit(c)) state=10;  
                    else state=11;  
                    break;  
            .....  
        }  
    }  
}
```

Another transition diagram for unsigned numbers

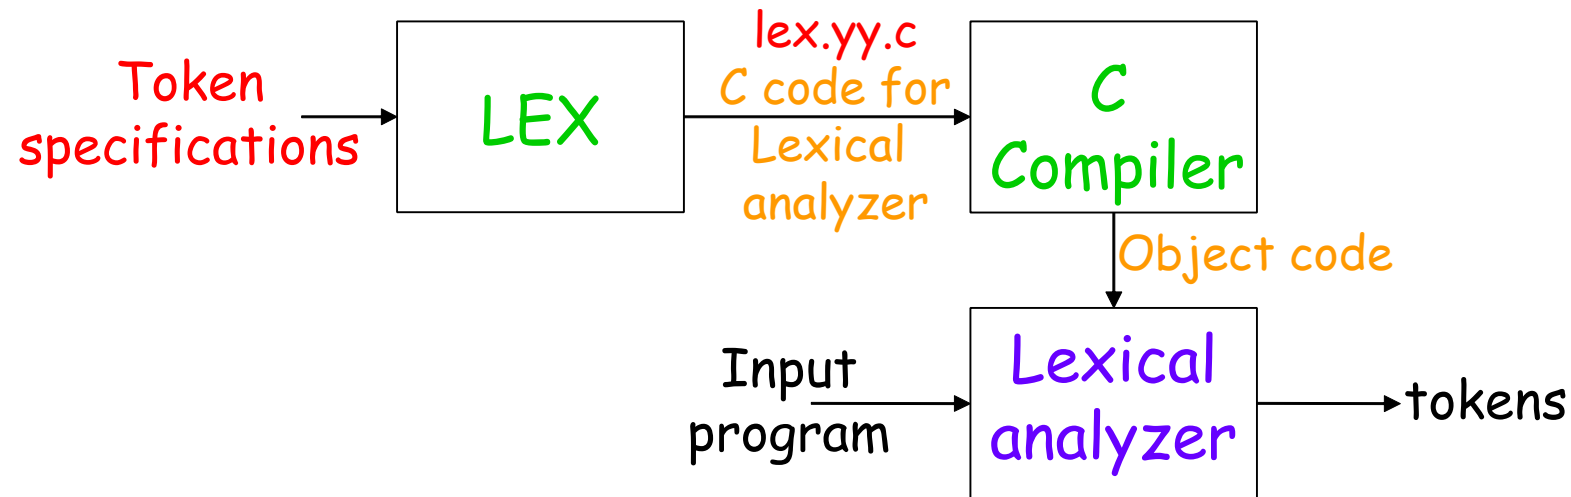


A more complex transition diagram
is difficult to implement and
may give rise to errors during coding,
however,
there are ways to better implementation

Lexical analyzer generator

- Input to the generator
 - List of regular expressions in priority order
 - Associated actions for each of regular expression (generates kind of token and other book keeping information)
- Output of the generator
 - Program that reads input character stream and breaks that into tokens
 - Reports lexical errors (unexpected characters), if any

LEX: A lexical analyzer generator



Refer to LEX User's Manual

How does LEX work?

- Regular expressions describe the languages that can be recognized by finite automata
- Translate each token regular expression into a non deterministic finite automaton (NFA)
- Convert the NFA into an equivalent DFA
- Minimize the DFA to reduce number of states
- Emit code driven by the DFA tables