


Backpropagation derivation



Error: $e(\mathbf{w}) = \frac{1}{2} \sum_{k=1}^{n_o} (t_k - z_k)^2$ assuming square error.

General update: $\Delta \mathbf{w} = -\eta \nabla \mathbf{e}$

Specific update for w_{kj} : $\Delta w_{kj} = -\eta \frac{\partial e}{\partial w_{kj}}$

Feed forward equations:

$$net_k = \sum_{j=1}^{n_h} w_{kj} y_j$$

$z_k = f(net_k)$ f is the activation function.


$$\begin{aligned}\frac{\partial e}{\partial w_{kj}} &= \frac{\partial e}{\partial net_k} \frac{\partial net_k}{\partial w_{kj}} \\ &= -(t_k - z_k) f'(net_k) \frac{\partial net_k}{\partial w_{kj}}\end{aligned}$$

$$-\frac{\partial e}{\partial w_{kj}} = \delta_k y_j, \quad \text{where } \delta_k = (t_k - z_k) f'(net_k)$$

$$\begin{aligned}\Delta w_{kj} &= -\eta \frac{\partial e}{\partial w_{kj}} \\ &= \eta \delta_k y_j \\ &= \eta (t_k - z_k) f'(net_k) y_j\end{aligned}$$

Assumes f is differentiable.

Hidden layer backpropagation



$$\Delta w_{ji} = -\eta \frac{\partial e}{\partial w_{ji}}$$

$$\frac{\partial e}{\partial w_{ji}} = \frac{\partial e}{\partial y_j} \frac{\partial y_j}{\partial net_j} \frac{\partial net_j}{\partial w_{ji}}$$

$$\frac{\partial net_j}{\partial w_{ji}} = x_i \quad \text{since } net_j = \sum_{i=0}^{n_i} w_{ji} x_i$$

$$\frac{\partial y_j}{\partial net_j} = f'(net_j) \quad \text{since } y_j = f(net_j)$$

$$\begin{aligned}
\frac{\partial e}{\partial y_j} &= \frac{\partial}{\partial y_j} \left[\frac{1}{2} \sum_{k=1}^{n_o} (t_k - z_k)^2 \right] \\
&= - \sum_{k=1}^{n_o} (t_k - z_k) \frac{\partial z_k}{\partial y_j} \\
&= - \sum_{k=1}^{n_o} (t_k - z_k) \frac{\partial z_k}{\partial net_k} \frac{\partial net_k}{\partial y_j} \\
&= - \sum_{k=1}^{n_o} (t_k - z_k) f'(net_k) w_{kj} \\
&= - \sum_{k=1}^{n_o} \delta_k w_{kj}
\end{aligned}$$

$$\Delta w_{ji} = \eta f'(net_j) x_i \sum_{k=1}^{n_o} \delta_k w_{kj}$$

$$= \eta \delta_j x_i \quad \text{where } \delta_j = f'(net_j) \sum_{k=1}^{n_o} \delta_k w_{kj}$$

Stochastic gradient descent alg.

Algorithm 0.1: $\text{SGD}(nw, \eta, \text{stopCriterion})$

Init \mathbf{w} randomly

repeat

$(\mathbf{x}_m, \mathbf{t}_m) \leftarrow$ choose randomly from \mathcal{L}
 $\mathcal{L} \leftarrow \mathcal{L} - (\mathbf{x}_m, \mathbf{t}_m)$
Feed forward
comment: y_j s and z_k s are now available.
 $w_{kj} \leftarrow w_{kj} + \eta \delta_k y_j$
 $w_{ji} \leftarrow w_{ji} + \eta \delta_j \mathbf{x}_{mi}$

until stopCriterion

return (\mathbf{w})

Stop criterion: a) error on a validation set b) limit θ on ∇e c) no. of epochs. An epoch is one exposure of every element in \mathcal{L} .

Batch alg.

Algorithm 0.2: $\text{BGD}(nw, \eta, \text{stopCriterion})$

Init \mathbf{w} randomly


repeat

$\left\{ \begin{array}{l} \text{epoch}, \Delta w_{ij}, \Delta w_{kj} \leftarrow 0 \\ \text{comment: Update being batched} \\ \text{repeat} \\ \left\{ \begin{array}{l} (\mathbf{x}_m, \mathbf{t}_m) \leftarrow \text{choose randomly from } \mathcal{L} \\ \mathcal{L} \leftarrow \mathcal{L} - (\mathbf{x}_m, \mathbf{t}_m) \\ \text{Feed forward comment: } y_j\text{s and } z_k\text{s are now available.} \\ \Delta w_{kj} \leftarrow \Delta w_{kj} + \eta \delta_k y_j \\ \Delta w_{ji} \leftarrow \Delta w_{ji} + \eta \delta_j \mathbf{x}_{mi} \end{array} \right. \\ \text{until } \mathcal{L} \text{ is empty} \\ \text{comment: Apply batched update} \\ \begin{array}{l} w_{kj} \leftarrow w_{kj} + \Delta w_{kj} / |\mathcal{L}| \\ w_{ji} \leftarrow w_{ji} + \Delta w_{ji} / |\mathcal{L}| \end{array} \end{array} \right.$

until stopCriterion

return (\mathbf{w})

Behaviour of error



Stochastic vs Batch

- ▶ Stochastic:
 - ▶ Usually faster than batch.
 - ▶ Higher probability of reaching better minimum.
 - ▶ Useful in tracking changes.
- ▶ Batch:
 - ▶ Convergence is well understood.
 - ▶ Many second order techniques to speed up convergence - but computationally expensive.
 - ▶ Theoretical analysis of convergence and dynamics is simpler.

In practice one often uses mini-batches. Often, the mini-batch size increases as epochs increase.

Activation functions

- ▶ Activation function: non-linear, smooth, continuous, saturating. Typical functions are tanh, sigmoid, relu (or approximator).
- ▶ $\tanh(x) = \frac{e^x - e^{-x}}{e^x + e^{-x}}$
- ▶ $\sigma(x) = \frac{1}{1 + e^{-x}}$
- ▶ $\text{relu}(x) = \max(0, x)$, smooth approx. $\text{softplus}(x) = \ln(1 + e^x)$

tanh




Figure: tanh - From: efunda.com

Sigmoid or logistic fn.




Figure: sigmoid - From: wikipedia

Relu and approx.




Figure: A Relu function and its smooth approx. From: wikipedia

Practical issues¹

- ▶ Weights cannot be 0. Initialize to values such that activation function is in the linear region. One recommendation is to choose weights from a distribution with mean 0 and
$$\sigma_w = \frac{1}{\sqrt{\text{fan-in-to-node}}}.$$
- ▶ The values of the input vector are normalized such that mean is 0 and variance is 1.
- ▶ Preprocess to remove correlated attributes (e.g. PCA).
- ▶ Presentation: order examples such that information is maximum.
 - ▶ Successive training examples from different classes.
 - ▶ Order examples such that successive examples produce large error more frequently. Danger: outliers can cause problems.

¹Mueller, et al., Neural networks: tricks of the trade, Springer, 2012

Practical issues - contd.

- ▶ Learning rate η : ideally, it should be chosen such that all weights in the network converge at the same rate.
 - ▶ If possible, give each weight or group of weights its own η .
 - ▶ Learning rate should be proportional to the fan-in.
 - ▶ Weights in earlier layers should typically be larger than later ones.
 - ▶ Learning rates can also be made adaptive if the gradient is remembered.