

Probabilistic/Bayesian Models for Deep Learning


Piyush Rai

Topics in Probabilistic Modeling and Inference (CS698X)

April 12, 2018

Neural Networks


- A simple neural network with one intermediate (also called “hidden”) layer and a single output



- Each intermediate layer computes a **nonlinear transformation** of its previous layer's nodes
- In traditional neural nets, \mathbf{h} is a **Linear transform** (e.g., $\mathbf{W}^{(1)}\mathbf{x}$ in the above picture) **followed by a nonlinearity** (e.g., sigmoid, ReLU, tanh, etc)
- Neural nets are awesome but brittle in many ways
 - Lots of parameters, difficult to train, need lots of data to train
 - Do not provide uncertainty estimates

What We Want..

- Neural networks with additional benefits of probabilistic/Bayesian modeling



- Basically, nonlinear models with estimates of uncertainty in the model/its predictions
- Note: We already have seen something that accomplishes this - [Gaussian Processes](#)
- Probabilistic/Bayesian neural nets are another alternative to this

Neural Networks as Probabilistic Models

- A probabilistic model for neural network for supervised learning


$$y_n \sim \mathcal{N}(\text{NN}(\mathbf{x}_n; \mathbf{W}), \beta^{-1}) \quad (\text{for real-valued responses})$$

$$y_n \sim \text{Bernoulli}(\sigma(\text{NN}(\mathbf{x}_n; \mathbf{W}))) \quad (\text{for binary responses})$$

$$y_n \sim \text{ExpFam}(\text{NN}(\mathbf{x}_n; \mathbf{W})) \quad (\text{for general types of responses modeled by exp-family})$$


where $\text{NN}(\mathbf{x}_n; \mathbf{W})$ is a neural network with features \mathbf{x}_n as inputs and parameters \mathbf{W}

- This enables learning probabilistic nonlinear input-to-output mappings
- We can perform point estimation or fully Bayesian inference for such probabilistic neural networks



Left: Standard NN or NN with point estimation, Right: Bayesian Neural Network

Learning Bayesian Neural Networks




- Even if the prior $p(\mathbf{W})$ on NN weights is Gaussian, the model is not conjugate
- MCMC methods can be used to learn the posterior $p(\mathbf{W}|\mathcal{D})$ but can be slow
 - However, methods such as SGLD[†] allow efficient MCMC inference for such models (recall that SGLD only requires gradient expressions of the log-joint probability $\log p(\mathcal{D}, \mathbf{W})$ of the model)
- Variational inference is another popular alternative to MCMC for such models
 - But ELBO is intractable (due to non-conjugacy); has to be approximated via Monte-Carlo/BBVI
- Note: Hybrid architectures also possible (only last layer modeled in a fully Bayesian way)

[†] "Preconditioned Stochastic Gradient Langevin Dynamics for Deep Neural Networks" (Li et al, 2016)

Constructing Generative Models using Neural Nets

- Probabilistic view enables neural net based **latent variable models** (“deep generative models”)
- Useful for **unsupervised learning**. Nonlinear latent variable to data mapping f modeled by NN




- Example: A probabilistic neural network for **latent variable modeling** (e.g., PPCA)
$$\mathbf{x}_n \sim \mathcal{N}(\text{NN}(\mathbf{z}_n; \mathbf{W}), \sigma^2 \mathbf{I}_D) \quad (\text{for real-valued features})$$
$$\mathbf{x}_n \sim \text{ExpFam}(\text{NN}(\mathbf{z}_n; \mathbf{W})) \quad (\text{for general types of features modeled by exp-family})$$

where $\text{NN}(\mathbf{z}_n; \mathbf{W})$ is a **neural network** with latent variables \mathbf{z}_n as inputs and parameters \mathbf{W}
- The NN enables learning a nonlinear latent variable to data mapping f
- If \mathbf{z}_n has a Gaussian prior, such models are called “**Deep Latent Gaussian Models**” (DLGM)

Inference for Deep Latent Gaussian Models


- Assume θ to be the global parameters of the model (params defining $p(\mathbf{z})$, $p(\mathbf{x}|\mathbf{z})$, etc.)



- The usual approach for inference in such models (as in most Bayesian models) is iterative, e.g.,
- Initialize θ . Then iterate until convergence
 - For $n = 1, \dots, N$
 - Infer $p(\mathbf{z}_n|\mathbf{x}_n)$ using MCMC. If doing VB, update variational parameters ϕ_n of $q(\mathbf{z}_n|\phi_n)$
 - Infer θ (its full posterior using MCMC or VB, or a point estimate)
- This iterative approach can be slow for large N
- Also, inferring \mathbf{z} for **new** data point(s) \mathbf{x} would require using the same iterative procedure

Variational Auto-encoder (VAE)

- Esseentially a DLGM, i.e., the \mathbf{z} to \mathbf{x} mapping $p(\mathbf{x}|\mathbf{z})$ is defined by a neural net. Proposed almost simultaneously by Kingma & Welling (2013), and Rezende *et al* (2014)



- VAE uses VB for inference but has a **fast, non-iterative** way of computing \mathbf{z}_n for a data point \mathbf{x}_n
- Key idea: For each point \mathbf{x}_n , instead of learning a separate $q(\mathbf{z}_n|\phi_n)$ with local params ϕ_n , assume


$$q(\mathbf{z}_n|\phi_n) = q(\mathbf{z}_n|\text{NN}(\mathbf{x}_n; \phi))$$

so, basically, each ϕ_n is computed by a neural net with global parameters ϕ and input \mathbf{x}_n


- Once ϕ is learned, we can get $q(\mathbf{z}_*|\mathbf{x}_*) = q(\mathbf{z}_*|\phi_*)$ for any \mathbf{x}_* by just using $\phi_* = \text{NN}(\mathbf{x}_*; \phi)$
- $p(\mathbf{x}|\mathbf{z})$ is known as **decoder** and $q(\mathbf{z}|\mathbf{x})$ is known as **encoder**

Standard Auto-encoder vs Variational Auto-encoder

- A standard auto-encoder learns to (nonlinearly) compress and uncompress an input




- Model is trained to minimize the **reconstruction error** (difference b/w x and \hat{x})
- However, it can't "generate" a "realistic" input from a random z (the model isn't trained for that)
- VAE allows this by assuming a distribution (e.g., Gaussian) over z and learning to generate x from random z 's drawn from that distribution (so the model is trained to do this!)



- Note: Simple generative models like PPCA or factor analysis also have this ability to generate data from random z but the linear map from z to x limits the type of data that can be generated well

VAE: The Encoder

- Role of encoder: Take \mathbf{x} as input and generate an encoding \mathbf{z}



Encoder


- Unlike standard autoencoders, for each \mathbf{x} , VAE gives us a distribution $q(\mathbf{z}|\mathbf{x})$ over its encoding
- Assume $q(\mathbf{z}|\mathbf{x})$ to be Gaussian whose mean/var are computed by a NN with global params ϕ

$$\mu_z = \text{NN}(\mathbf{x}; \phi) \quad \sigma_z^2 = \text{NN}(\mathbf{x}; \phi)$$

- Since μ_z, σ_z are outputs of neural networks, the \mathbf{x} to \mathbf{z} mapping is nonlinear

VAE: The Decoder

- Role of decoder: Generate \mathbf{x} given \mathbf{z} . Defined by the likelihood model $p_{\theta}(\mathbf{x}|\mathbf{z})$




- Unlike PPCA, the \mathbf{z} to \mathbf{x} mapping is nonlinear (modeled by a neural network)
- Assume $p(\mathbf{x}|\mathbf{z})$ to be Gaussian whose mean/var are computed by a NN with global params θ

$$\mu_{\mathbf{x}} = \text{NN}(\mathbf{z}; \theta) \quad \sigma^2_{\mathbf{x}} = \text{NN}(\mathbf{z}; \theta)$$

- Thus in the VAE, both \mathbf{x} to \mathbf{z} (encoder) and \mathbf{z} to \mathbf{x} (decoder) mappings are nonlinear

Inference for VAE

- VAE uses variational inference (hence the name!) to learn the model parameters θ and ϕ



- Typically a prior $p(z) = \mathcal{N}(\mathbf{0}, \mathbf{I}_K)$ is assumed on z . The ELBO for a single \mathbf{x}_n will be


$$\text{ELBO} = \mathbb{E}_{q_\phi} [\log p(\mathbf{x}_n, \mathbf{z}_n | \theta) - \log q(\mathbf{z}_n | \mathbf{x}_n)] \quad (\text{note: } q_\phi \text{ and } q(\mathbf{z}_n | \mathbf{x}_n) \text{ mean the same})$$

- Variational inference uses the reparametrization trick[†] for computing ELBO derivatives

[†] "Auto-encoding Variational Bayes" (Kingma and Welling, 2013)

Some Other Architectures: Deep Exponential Families


- Standard VAE has only one layer of latent z and a neural net to transform z into x
- Many other deep architectures have multiple layers of latent variables
- Deep Exponential Family (DEF) is one such recently proposed popular model
- in DEF, latent variables in every layer, as well as observations, are from exp. family distributions



- Overall model not conjugate but BBVI (Ranganath et al, 2013) or MCMC methods can be used

Some Other Architectures: Generative Adversarial Networks

- Based on a game between a generator and a discriminator (Goodfellow et al, 2013)



- Generator (a neural net) generates realistic looking “fake” data x from random z
- Discriminator tries to detect fake data from real data
- At game’s equilibrium, the $p_{gen} = p_{data}$ and success rate of discriminator = 50% (i.e., random)
- Originally designed mainly for synthetic data generation tasks but recent work extends GANs for many other problems such as latent variable inference, semi-supervised learning, etc.

Learning the right size of a deep neural network

- How to decide the number of layers and width of each layer?
- Nonparametric Bayesian methods can help here



- A [cascaded Indian Buffet Prior](#) can model the relationships between nodes in adjacent layers
 - The bottom-most layer is the data layer (fixed/known size)
 - Width of each intermediate layer and active connections can be inferred by the IBP prior
- Another option is to use sparsity inducing priors on the connection weights

Neural Nets vs Gaussian Processes

- Both can be used learn nonlinear input to output mappings, e.g.,

$$\begin{aligned}y_n &\sim \mathcal{N}(\text{NN}(\mathbf{x}_n; \mathbf{W}), \beta^{-1}) \\y_n &\sim \mathcal{N}(f(\mathbf{x}_n), \beta^{-1}) \quad \text{where } f \sim \text{GP}\end{aligned}$$

- Both have their pros and cons
- NN pros: Fast to train (e.g., using SGD methods) and also fast at test time
- NN cons: Difficult to train, also not Bayesian (but can be made Bayesian)
- GP pros: Simple formulation, especially for regression settings; natively Bayesian in formulation
- GP cons: Slow to train and also slow at test time
- Nowadays Bayesian NN and GPs are competitive in many applications

Summary

- Probabilistic modeling allows developing very flexible deep learning models
- Much of the recent progress is fuelled by advances in probabilistic modeling and inference
- State-of-the-art results on a variety of tasks such as
 - Representation Learning (latent variables used as a new learned representation of data)
 - Density Estimation (i.e., $p(x)$)
 - Data generation (models like VAE and GAN can generate very realistic looking synthetic data)
 - Semi-supervised learning (models like VAE and GAN can be combined with supervised learning)
- Several other models such as Deep Boltzmann Machines, Neural Autoregressive Density Estimator, etc. that we didn't cover here
- An important distinction between **explicit** and **implicit** generative models
 - Models like PPCA, FA, DLGM, SBN, VAE, etc. have an explicit likelihood model for data
 - A model like GAN only defines $p(x)$ implicitly (no “likelihood” model for data)